Answer Sheet

Q1)

An organization has decided to reserve EC2 compute capacity for three years to get more discounts.

Their application workloads are likely to change during this time period.

What is the EC2 Reserved Instance (RI) type that allows them to change the attributes of the RI whenever they need?

- Elastic RIs

Explanation:-This option is not correct. Elastic RI is not a valid RI type.

- Scheduled RIs

Explanation:-This option is not correct. Scheduled RIs are available to launch within the time windows you reserve. This option allows you to match your capacity reservation to a predictable recurring schedule that only requires a fraction of a day, a week, or a month.

- ✅ Convertible RIs

Explanation:-This option is correct. Convertible RIs provide a discount (up to 54% off On-Demand) and the capability to change the attributes of the RI as long as the exchange results in the creation of Reserved Instances of equal or greater value. These attributes include instance family, instance type, platform, scope, and tenancy.

- Standard RIs

Explanation:-This option is not correct. Standard RIs provide the most significant discount (up to 75% off On-Demand) and are best suited for steady-state usage.

Q2) What does the AWS Snowball provide?

- An encrypted SSL endpoint for backups in the Cloud.

Explanation:-This option is not correct.

- ✅ Secure transfer of large amounts of data into and out of the AWS Cloud.

Explanation:-This option is correct. Snowball is a petabyte-scale data transport solution that uses devices designed to be secure to transfer large amounts of data into and out of the AWS Cloud. Using Snowball addresses common challenges with large-scale data transfers including high network costs, long transfer times, and security concerns. Customers today use Snowball to migrate analytics data, genomics data, video libraries, image repositories, backups, and to archive part of data center shutdowns.

- A backup solution that provides on-premises Cloud storage.

Explanation:-This option is not correct.

- A direct encrypted connection to Amazon S3.

Explanation:-This option is not correct.

Q3)

Derek is running a web application and is noticing that he is paying for way more server capacity than required.

What AWS feature should Derek set up and configure to ensure that his application is automatically adding/removing server capacity to closely match the required demand?

- AWS Inspector

Explanation:-This option is not correct. AWS Inspector is an automated security assessment service to help improve the security and compliance of applications deployed on AWS

- AWS EC2

Explanation:-This option is not correct. AWS EC2 is the service that provides the compute capacity that your applications need.

- ✅ AWS Autoscaling

Explanation:-This option is correct. Auto scaling is the feature that automates the process of adding/removing the server capacity (based on demand). Autoscaling allows you to reduce your costs by automatically turning off resources that aren't in use. On the other hand Autoscaling ensures that your application runs effectively by provisioning more server capacity if required.

- AWS ELB

Explanation:-This option is not correct. AWS ELB is the service that distributes the incoming application traffic to multiple targets that you define.

Q4)

You work as an on-premises MySQL DBA. The work of database configuration, backups, patching, and DR can be time-consuming and repetitive. Your company has decided to migrate to the AWS Cloud.

Which of the following can help save time on the regular database tasks so you can focus on giving users the fast performance and high availability that they need?

- Amazon CloudWatch

Explanation:- This option is not correct. Amazon CloudWatch is a monitoring service that gives you complete visibility of your cloud resources and applications.

- ✅ Amazon RDS

Explanation:-This option is correct. Amazon Relational Database Service (Amazon RDS) makes it easy to set up, operate, and scale a relational database in the cloud. It provides cost-efficient, resizable capacity while automating time-consuming administration tasks such as hardware provisioning, database setup, patching and backups. It frees you to focus on your applications so you can give them the fast performance, high availability, security and compatibility they need.

- Amazon Redshift

Explanation:-This option is not correct. Amazon Redshift is a fast, fully managed data warehouse that makes it simple and cost-effective to analyze all your data using standard SQL and your existing Business Intelligence (BI) tools.

- Amazon DynamoDB

Explanation:-This option is not correct. Amazon DynamoDB is a NoSQL database service.

Q5)

An organization runs many systems and uses many AWS products.

Which of the following services allow them to control how each developer interacts with these products?

✅ AWS Identity and Access Management
Explanation:-This option is correct. AWS Identity and Access Management (IAM) is a web service for securely controlling access to AWS services. With IAM, you can centrally manage users, security credentials such as access keys, and permissions that control which AWS resources users and applications can access.
⚪ Amazon EMR
Explanation:-This option is incorrect.
⚪ Network Access Control Lists.
Explanation:-This option is incorrect.
⚪ Amazon RDS
Explanation:-This option is incorrect.

Q6) Which service allows the customer to retain full administrative privileges of the underlying virtual infrastructure?

⚪ Amazon RDS
Explanation:-This option is incorrect.
⚪ Amazon Redshift.
Explanation:-This option is incorrect.
✅ Amazon EC2
Explanation:-This option is correct. Amazon EC2 provides you with complete control of your virtual instances including root access and the ability to interact with them as you would any machine.
All other services belong to the managed services. AWS Managed Services (AMS) operates AWS on your behalf, providing a secure and compliant AWS Landing Zone and day-to-day infrastructure management.
⚪ Amazon DynamoDB
Explanation:-This option is incorrect.

Q7) Which of the following is one of the benefits of AWS Security?

⚪ None of these
Explanation:-This option is incorrect.
✅ Scales Quickly
Explanation:-This option is correct. Security scales with your AWS Cloud usage. No matter the size of your business, the AWS infrastructure is designed to keep your data safe.
⚪ Starts automatically once you upload your data.
Explanation:-This option is not correct. AWS Security doesn't start automatically, you have to go on and set up how your data will be accessed and decide whether this data will be encrypted or not and so on.
⚪ Free for AWS premium members.
Explanation:-This option is not correct. Not all security features are free. For example security groups are free for all customers however Amazon Inspector is not free.

Q8) Which of the following services allows you to run containerized applications on a cluster of EC2 instances?

⚪ Amazon Elastic Compute Service.
Explanation:-This option is incorrect.
⚪ AWS Docker Manager.
Explanation:-This option is incorrect.
⚪ Amazon Elastic Docker Service.
Explanation:-This option is incorrect.
✅ Amazon Elastic Container Service.
Explanation:-This option is correct. Amazon Elastic Container Service (Amazon ECS) is a highly scalable, high-performance container orchestration service that supports Docker containers and allows you to easily run and scale containerized applications on AWS. Amazon ECS eliminates the need for you to install and operate your own container orchestration software, manage and scale a cluster of virtual machines, or schedule containers on those virtual machines.
Other options presented are bogus.

Q9) What information is required to calculate the Total Cost of Ownership for the AWS Cloud?

⚪ The number of end users you are currently serving
Explanation:-This option is incorrect.
⚪ The number of on-premise applications
Explanation:-This option is not correct. The TCO Calculator doesn't ask for the number of applications migrated to AWS. Also, what really matters is the size of your application and your business.
✅ The number of on-premise virtual machines
Explanation:-This option is correct. The AWS TCO (Total Cost of Ownership) Calculator provides directional guidance on possible realized savings when deploying AWS. This tool is built on an underlying calculation model, that generates a fair assessment of value that a customer may achieve given the data provided by the user which includes the number of servers migrated to AWS, the server type, the number of processors and so on.
⚪ The number of active databases
Explanation:-This option is not correct. The TCO calculator asks for the number of servers that run your databases NOT the number of active databases.

Q10) What services/features are required to maintain a highly available and fault-tolerant architecture in AWS? (Choose two)

✅ Amazon EC2 Auto Scaling

Explanation:-This option is correct. Amazon EC2 Auto Scaling continually monitors the utilization of the instances underlying your application to make sure that your application always has the right amount of compute. In other words Amazon EC2 Auto Scaling automatically scales the instances up during demand spikes (to increase the availability of the application) or scales them down when demand lulls (to minimize costs). In addition to that, Amazon EC2 Auto Scaling can detect when an instance is unhealthy.

⚪ CloudFormation

Explanation:-This option is not correct. CloudFormation provides an organized method to deploy all of your AWS resources.

⚪ AWS NACL

Explanation:-This option is not correct. AWS NACL is used to control traffic at the subnet level.

✅ Elastic Load Balancer

Explanation:-This option is correct. Amazon EC2 Auto Scaling continually monitors the utilization of the instances underlying your application to make sure that your application always has the right amount of compute. In other words Amazon EC2 Auto Scaling automatically scales the instances up during demand spikes (to increase the availability of the application) or scales them down when demand lulls (to minimize costs). In addition to that, Amazon EC2 Auto Scaling can detect when an instance is unhealthy.

Q11)

A company is hosting a web application in the AWS Cloud using a set of EC2 instances.

Which of the following can help to protect them from DDoS attacks? (Choose two)

⚪ Using AWS Batch

Explanation:-This option is not correct. AWS Batch is a compute service that allows you to run hundreds of thousands of batch computing jobs on AWS.

✅ Using Network Access Control Lists

Explanation:-This option is correct. A security group acts as a virtual firewall for your instance to control inbound and outbound traffic.
A network access control list (ACL) acts as a firewall for controlling traffic in and out of one or more subnets. Therefore if they configured properly, they can protect your instances from DDoS attacks.

✅ Using Security Groups

Explanation:-This option is correct. A security group acts as a virtual firewall for your instance to control inbound and outbound traffic.
A network access control list (ACL) acts as a firewall for controlling traffic in and out of one or more subnets. Therefore if they configured properly, they can protect your instances from DDoS attacks.

⚪ Using CloudHSM

Explanation:-This option is not correct. AWS CloudHSM is a cloud-based hardware security module (HSM) that enables you to easily generate and use your own encryption keys on the AWS Cloud.

⚪ Using the Internet gateway

Explanation:-This option is not correct. An internet gateway is a VPC component that allows communication between instances in your VPC and the internet. It therefore imposes no availability risks or bandwidth constraints on your network traffic.

Q12)

Kim is managing a web application running on the AWS Cloud. The application is currently utilizing eight EC2 servers for its computing platform.

Earlier today, two of those web servers crashed; however, none of her customers were affected.

What has Kim done correctly in this scenario?

✅ She has properly built a fault tolerant system.

Explanation:-This option is correct. Fault tolerance is the property that enables a system to continue operating properly in the event of the failure of some (one or more faults within) of its components.

⚪ None of the above.

Explanation:-This option is incorrect.

⚪ She has properly built a scalable system.

Explanation:-This option is not correct. Scalability is the ability of the system to accommodate larger loads just by adding resources either making hardware stronger (scale up) or adding additional nodes (scale out).

⚪ She has properly built an elastic system.

Explanation:-This option is not correct. Elasticity is the ability of a system to scale the resources needed to cope with load dynamically. So that when the load increases you scale by adding more resources and when demand wanes you shrink back and remove unneeded resources.

Q13)

A company created a solution that will help AWS customers improve their architectures on AWS.

Which AWS program may support this company?

⚪ AWS TAM

Explanation:-This option is not correct. A Technical Account Manager (TAM) is your designated technical point of contact who provides advocacy and guidance to help plan and build solutions using best practices and proactively keep your AWS environment operationally healthy. TAM is available only for the Enterprise support plan.

⚪ AWS Professional Services

Explanation:-This option is not correct. AWS Professional Services shares a collection of offerings to help you achieve specific outcomes related to enterprise cloud adoption. AWS Professional Services also trains your team with specialized skills and provides global specialty practices to support your efforts in focused areas of enterprise cloud computing.

✅ APN Consulting Partners

Explanation:-This option is correct. APN Consulting Partners are professional services firms that help customers design, architect, build, migrate, and manage their workloads and applications on AWS. Consulting Partners include System Integrators, Strategic Consultancies, Agencies, Managed Service Providers, and Value-Added Resellers. AWS supports the APN Consulting Partners by providing a wide range of resources and training to support their customers.

⚪ APN Technology Partners

Explanation:-This option is not correct. APN Technology Partners provide software solutions that are either hosted on, or integrated with, the AWS platform. APN Technology Partners include Independent Software Vendors (ISVs), SaaS, PaaS, Developer Tools, Management and Security Vendors.

Q14) AWS provides a storage option known as Amazon Glacier, which is designed for_____ & _____. (Choose two)

- Active database storage

Explanation:-This option is not correct. Active databases require consistent and low-latency storage performance. For example DB instances for Amazon RDS for MySQL, MariaDB, PostgreSQL, Oracle, and Microsoft SQL Server use Amazon Elastic Block Store (Amazon EBS) volumes for database and log storage.

- ✅ Active archive storage

Explanation:-This option is correct. Amazon S3 Glacier provides three retrieval options to fit your use case. Expedited retrievals typically return data in 1-5 minutes, and are great for Active Archive use cases. Standard retrievals typically complete between 3-5 hours work, and work well for less time-sensitive needs like backup data, media editing, or long-term analytics. Bulk retrievals are the lowest-cost retrieval option, returning large amounts of data within 5-12 hours.

- Cached session data

Explanation:-This option is not correct. In computing, a cache is a high-speed data storage layer which stores a subset of data, typically transient in nature, so that future requests for that data are served up faster than is possible by accessing the data's primary storage location. Caching allows you to efficiently reuse previously retrieved or computed data. The data in a cache is generally stored in fast access hardware such as RAM (Random-access memory) and may also be used in correlation with a software.

- ✅ Long-term analytics

Explanation:-This option is correct. Amazon S3 Glacier provides three retrieval options to fit your use case. Expedited retrievals typically return data in 1-5 minutes, and are great for Active Archive use cases. Standard retrievals typically complete between 3-5 hours work, and work well for less time-sensitive needs like backup data, media editing, or long-term analytics. Bulk retrievals are the lowest-cost retrieval option, returning large amounts of data within 5-12 hours.

Q15) Which of the following can be used to automate the management of multiple AWS services through scripts?

- AWS Console

Explanation:-This option is not correct. AWS Console lets you access and manage Amazon Web Services through a web-based user interface.

- ✅ AWS CLI

Explanation:-This option is correct. The AWS Command Line Interface (CLI) is a unified tool to manage your AWS services. With just one tool to download and configure, you can control multiple AWS services from the command line and automate them through scripts.

- AWS Service Catalog

Explanation:-This option is not correct. AWS Service Catalog allows organizations to create and manage catalogs of IT services that are approved for use on AWS.

- AWS OpsWorks

Explanation:-This option is not correct. AWS OpsWorks can be used to automate one service which is EC2. AWS OpsWorks is a configuration management service that provides managed instances of Chef and Puppet. Chef and Puppet are automation platforms that allow you to use code to automate the configurations of your servers. OpsWorks lets you use Chef and Puppet to automate how servers are configured, deployed, and managed across your Amazon EC2 instances or on-premises compute environments.

Q16)

You have developed a microservices-based application.

Which of the following should you use to make sure that each EC2 instance in the system gets the same amount of traffic?

- Network Load Balancer.

Explanation:-This option is not correct. The traffic comes to the instances through HTTP or HTTPS. Network Load Balancer is best suited for load balancing of TCP and TLS traffic.

- Auto Scaling

Explanation:-This option is not correct. AWS Auto Scaling is not for distributing traffic. AWS Auto Scaling monitors your applications and automatically adjusts capacity (up or down) to maintain steady, predictable performance at the lowest possible cost.

- Amazon EC2 Auto Recovery

Explanation:-This option is not correct. Auto Recovery is an Amazon EC2 feature that is designed to increase instance availability. Auto Recovery can be configured to automatically recover EC2 Instances when a system or hardware impairment is detected.

- ✅ Application Load Balancer

Explanation:-This option is correct. Elastic Load Balancing automatically distributes incoming application traffic across multiple targets, such as Amazon EC2 instances, containers, IP addresses, and Lambda functions. Elastic Load Balancing offers three types of load balancers: 1- Application Load Balancer. 2- Network Load Balancer. 3- Classic Load Balancer. Application Load Balancer is best suited for load balancing of HTTP and HTTPS traffic. In our case, the microservices application receives HTTP or HTTPS

Q17) Which of the following AWS services help migrate an existing database to AWS?

- AWS Snowball

Explanation:-This option is not correct. Snowball is a petabyte-scale data transport solution that uses devices designed to securely transfer large amounts of data into and out of the AWS Cloud.

- AWS Lambda

Explanation:-This option is not correct. AWS Lambda is a serverless compute service.

- AWS Storage Gateway

Explanation:-This option is not correct. AWS Storage Gateway is a hybrid cloud storage.

- ✅ AWS DMS

Explanation:-This option is correct. AWS Database Migration Service helps you migrate databases to AWS quickly and securely. The source database remains fully operational during the migration, minimizing downtime to applications that rely on the database. The AWS Database Migration Service can migrate your data to and from most widely used commercial and open-source databases.

Q18)

Your company uses on-demand EC2 Instances dedicated to a project that has just been cancelled.

The company does not want to incur charges for these on-demand Instances.

However, it also does not want to lose the data yet because there is a chance the project may be revived in the next few days.

What should you do to minimize charges for these Instances in the meantime?

✅ Stop the instances as soon as possible
Explanation:-This option is correct. The best way to minimize charges is to stop the instances to avoid any data transfer charges that the instance may incur if left running.
⚪ You cannot minimize charges for this type of instances
Explanation:-This option is not correct. You can minimize charges by stopping the instances when you don't need them.
⚪ Sell the instances on the AWS On-Demand Instance Marketplace as you can buy other on-demand EC2 instances later if needed.
Explanation:-This option is not correct. You cannot sell on-demand instances. With the On-demand option you don't pay any thing upfront and you pay only for what you use. On-Demand instances let you pay for compute capacity by the hour or second (minimum of 60 seconds).
⚪ Terminate the instances as soon as possible
Explanation:-This option is not correct. If you terminate the instances without taking an image (AMI) of them, you will lose their data.

Q19) What best describes penetration testing?

✅ Testing your own network/application for vulnerabilities.
Explanation:-This option is correct. Penetration testing is the practice of testing a network or web application to find security vulnerabilities that an attacker could exploit.
⚪ Testing your IAM users' access to AWS services.
Explanation:-This option is not correct.
⚪ None of these
Explanation:-This option is not correct.
⚪ Testing your applications ability to penetrate other applications.
Explanation:-This option is not correct.

Q20) What is the framework created by AWS Professional Services that helps your organization design a road map to successful cloud adoption?

⚪ AWS KAF
Explanation:-This option is incorrect.
⚪ AWS DAF
Explanation:-This option is incorrect.
✅ AWS CAF
Explanation:-This option is correct. AWS Professional Services created the AWS Cloud Adoption Framework (AWS CAF) to help organizations design and travel an accelerated path to successful cloud adoption. The guidance and best practices provided by the framework help you build a comprehensive approach to cloud computing across your organization, and throughout your IT lifecycle. Using the AWS CAF helps you realize measurable business benefits from cloud adoption faster and with less risk.
⚪ AWS MAF
Explanation:-This option is incorrect.

Q21) Which of the following statements describes the AWS Cloud's agility?

⚪ AWS allows you to host your applications in multiple regions around the world.
Explanation:-This option is not correct. It is true that AWS provides global infrastructure, but this statement doesn't describe AWS' agility.
✅ AWS allows you to provision resources in minutes.
Explanation:-This option is correct. In a cloud computing environment, new IT resources are only a click away, which means that you reduce the time to make those resources available to your developers from weeks to just minutes. This results in a dramatic increase in agility for the organization, since the cost and time it takes to experiment and develop is significantly lower.
⚪ AWS provides virtually unlimited compute and storage capacity.
Explanation:-This option is not correct. It is true that AWS provides virtually unlimited compute and storage capacity, but this statement doesn't describe AWS' agility.
⚪ AWS provides customizable hardware at the lowest possible cost.
Explanation:-This option is not correct. AWS doesn't provide customizable hardware. AWS provides virtualized resources in the cloud.

Q22) What is the DynamoDB replication technology that provides fast, local, read/write performance for globally-deployed applications?

⚪ Point-in-time recovery
Explanation:-This option is not correct. PITR refers to Point-in-time recovery. PITR is used to back up your data with per-second granularity and restore to any single second from the time PITR was enabled up to the prior 35 days. PITR works as additional insurance against accidental loss of data. Also, PITR is not global, you can only enable it within a single region.
⚪ DynamoDB DAX
Explanation:-This option is not correct. DynamoDB DAX is an in-memory cache for DynamoDB that reduces response times from milliseconds to microseconds.
✅ Global Tables
Explanation:-This option is correct. Global Tables builds upon DynamoDB's global footprint to provide you with a fully managed, multi-region, and multi-master database that provides fast, local, read and write performance for massively scaled, global applications. Global Tables replicates your Amazon DynamoDB tables automatically across your choice of AWS regions. Multi-master replication ensures that updates performed in any region are propagated to other regions, and that data in all regions are eventually
⚪ Global PITR
Explanation:-This option is not correct. PITR refers to Point-in-time recovery. PITR is used to back up your data with per-second granularity and

restore to any single second from the time PITR was enabled up to the prior 35 days. PITR works as additional insurance against accidental loss of data. Also, PITR is not global, you can only enable it within a single region.

Q23) What is the most cost effective EC2 Instance purchasing option for users with urgent computing needs for large amounts of additional capacity?

- On-demand instances

Explanation:-This option is incorrect.

- Reserved instances

Explanation:-This option is incorrect.

- ✅ Spot instances

Explanation:-This option is correct. Amazon EC2 Spot instances allow you to request spare Amazon EC2 computing capacity for up to 90% off the On-Demand price. Spot instances are recommended for:
*Applications that have flexible start and end times
**Applications that are only feasible at very low compute prices
***Users with urgent computing needs for large amounts of additional capacity
The other options are incorrect. Spot Instances are the cheapest EC2 option as you can save up to 90%.

- Dedicated instances

Explanation:-This option is incorrect.

Q24) You have multiple accounts in AWS. Which of the following can be used to reduce costs?

- There are no cost benefits when having multiple AWS accounts

Explanation:-This option is incorrect.

- ✅ Consolidated billing

Explanation:-This option is correct. The Consolidated Billing feature is part of the “AWS organization” service. Once you add your multiple accounts to an Organization, AWS charges you based on a consolidated bill for all of the linked accounts. With Consolidated Billing, you can see a combined view of AWS charges incurred by all accounts, as well as get a cost report for each individual account associated with your payer account. For billing purposes, AWS treats all the accounts in the organization as if th

- Costs are automatically reduced for the one who has multiple accounts by AWS.

Explanation:-This option is not correct. Costs are not automatically reduced for multiple accounts. If you don’t create an AWS organization for them, you will not benefit from the consolidated billing.

- Combined billing

Explanation:-This option is not correct. Combined billing is bogus option.

Q25) Which of the following can be used to process a large number of binary files while following the AWS well-architected design principles?

- Use a number of parallel RDS instances

Explanation:-This option not correct. RDS instances are used to store and run databases.

- Use a number of vertically scalable RDS instances

Explanation:-This option not correct. RDS instances are used to store and run databases.

- Use a number of vertically scalable EC2 instances

Explanation:-This option is not correct. Horizontal scaling is recommended over vertical scaling.

- ✅ Use a number of parallel EC2 instances

Explanation:-This option is correct. One of the most important design principles is to “scale horizontally to increase aggregate system availability”. Replace one large resource with multiple small resources to reduce the impact of a single failure on the overall system. For example if you want to convert a large number of binary files to text files or if you want to transcode large number of video files to another format, it is recommended that you use multiple EC2 instances in parallel instead of using one

---

Q26)

You are planning to offload some of your batch processing workloads to AWS.

These jobs can be interrupted and resumed at any time.

Which of the following instance types would be the most cost effective to use?

- Partial Upfront Reserved

Explanation:-This option is not correct. Since these jobs can be interrupted and resumed at any time, then you should use Spot instances as it provides the largest discount compared to any other payment option.

- On-Demand

Explanation:-This option is not correct. The Spot option provides discounts up to 90% off compared to the On-Demand price.

- ✅ Spot

Explanation:-This option is correct. Spot Instances are a cost-effective choice if you can be flexible about when your applications run and if your applications can be interrupted. For example, Spot Instances are well-suited for data analysis, batch jobs, background processing, and optional tasks.

- Full Upfront Reserved

Explanation:-This option is not correct. Since these jobs can be interrupted and resumed at any time, then you should use Spot instances as it provides the largest discount compared to any other payment option.

Q27) Which of the following storage mechanisms can be used to store and reliably deliver messages effectively across distributed systems?

- Amazon EBS Snapshots

Explanation:-This option is not correct. Amazon EBS Snapshots are backups for your EBS volumes.

- Amazon Glacier

Explanation:-This option is not correct. Amazon Glacier is used for backups and data archiving.

Amazon EBS Volumes
Explanation:-This option is not correct. Amazon EBS Volumes can only be used as a mountable disk drive for Amazon EC2 or Amazon RDS.
Amazon SQS
Explanation:-This option is correct. Amazon Simple Queue Service (Amazon SQS) offers a reliable, highly-scalable hosted queue for storing messages as they travel between applications or microservices. It moves data between distributed application components and helps you decouple these components.

Q28) Which of the following reserved instance payment options result in you paying a discounted hourly rate throughout the duration of the term? (Choose two)

All Upfront option.
Explanation:-This option is not correct. When choosing “All Upfront”, a full payment is made at the start of the term, with no other costs or additional hourly charges incurred for the remainder of the term, regardless of hours used.
Percentage Upfront option.
Explanation:-This option is not correct. Percentage upfront is not an available option.
Partial Upfront option.
Explanation:-This option is correct. You can choose between three payment options when you purchase a Standard or Convertible Reserved Instance:

No Upfront option.
Explanation:-This option is correct. You can choose between three payment options when you purchase a Standard or Convertible Reserved Instance:
1- No Upfront:
No upfront payment is required. You are billed a discounted hourly rate for every hour within the term, regardless of whether the Reserved Instance is being used. No Upfront Reserved Instances are based on a contractual obligation to pay monthly for the entire term of the reservation. A successful billing history is required .

Q29) Which statement is correct with regards to service limits? (Choose two)

Each IAM user has the same service limits.
Explanation:-This option is not correct. Service limits are applied at the AWS account level by aggregating usage from all users in the account.
The AWS Simple Email Service is responsible for sending email notifications when usage approaches a service limit.
Explanation:-This option is not correct. Amazon Simple Email Service (Amazon SES) is a cloud-based email sending service designed to help digital marketers and application developers send marketing, notification, and transactional emails.
Additional information:
You can configure the AWS Limit Monitor to send email notification when usage approaches a service limit.
You can contact support to increase the service limits.
Explanation:-This option is correct. Understanding your service limits (and how close you are to them) is an important part of managing your AWS deployments – continuous monitoring allows you to request limit increases or shut down resources before the limit is reached. One of the easiest ways to do this is via AWS Trusted Advisor’s Service Limit Dashboard, which currently covers 39 limits across 10 services.
AWS maintains service limits for each account to help guarantee the availability of AWS resource.
You can use the AWS Trusted Advisor to monitor your service limits.
Explanation:-This option is correct. Understanding your service limits (and how close you are to them) is an important part of managing your AWS deployments – continuous monitoring allows you to request limit increases or shut down resources before the limit is reached. One of the easiest ways to do this is via AWS Trusted Advisor’s Service Limit Dashboard, which currently covers 39 limits across 10 services.
AWS maintains service limits for each account to help guarantee the availability of AWS resource.
There are no service limits on AWS.
Explanation:-This option is not correct. You can find a full list of the AWS service limits on this page:

Q30) Which services does AWS offer for free? (Choose TWO)

Elastic Load Balancing
Amazon EC2
Elastic Beanstalk
Explanation:-AWS Identity and Access Management is a feature of your AWS account offered at no additional charge. You will be charged only for use of other AWS services by your Users.
There is no additional charge for AWS Elastic Beanstalk. You pay for AWS resources (e.g. EC2 instances or S3 buckets) you create to store and run your application. You only pay for what you use, as you use it; there are no minimum fees and no upfront commitments.
AWS IAM
Explanation:-AWS Identity and Access Management is a feature of your AWS account offered at no additional charge. You will be charged only for use of other AWS services by your Users.
There is no additional charge for AWS Elastic Beanstalk. You pay for AWS resources (e.g. EC2 instances or S3 buckets) you create to store and run your application. You only pay for what you use, as you use it; there are no minimum fees and no upfront commitments.
Amazon RDS

Q31) What does Amazon Elastic Beanstalk provide?

A scalable cluster of EC2 instances.
Explanation:-This option is incorrect.
A service by this name doesn't exist.
Explanation:-This option is incorrect.
An application container on top of Amazon Web Services.
Explanation:-This option is correct. AWS Elastic Beanstalk makes it easy for developers to quickly deploy and manage applications in the AWS Cloud. Developers simply upload their application code, and Elastic Beanstalk automatically handles the deployment details of capacity provisioning, load balancing, auto-scaling, and application health monitoring.
A scalable storage appliance on top of Amazon Web Services.
Explanation:-This option is incorrect.

Q32)

Sarah has created a web application, placing its underlining infrastructure in the N. Virginia (US-East-1) region.

After several months, Sarah notices that much of her website's traffic is coming from Japan.

What can she do to reduce latency for her users in Japan?

- None of these

Explanation:-This option is incorrect.

- ✅ Create a CDN using CloudFront, so that content is cached at Edge Locations close to and in Japan

Explanation:-This option is correct. CloudFront is AWS's content delivery network (CDN) service. You can use it to cache web content at edge locations that are closest to your end users. This will reduce latency and improve the overall performance.

- Migrate the application to a Japanese hosting service

Explanation:-This option is not correct. Cloud Computing now can deal with most of the customers' requirements. Whatever your problem is you can find a solution.

- Recreate the website content

Explanation:-This option is not correct. There is no relation between the website content and the traffic that comes to the web application.

Q33)

You are working on two projects that require a completely different network configuration.

Which of the following would allow you to isolate resources and network configurations for each of them?

- Edge Locations

Explanation:-This option is not correct. Edge Locations are used by CloudFront to distribute content to end users with low latency.

- Virtual Public Cloud

Explanation:-This option is not correct. There is nothing called Virtual Public Cloud.

- ✅ Virtual Private Cloud

Explanation:-This option is correct. Amazon Virtual Private Cloud (Amazon VPC) lets you provision a logically isolated section of the AWS Cloud where you can launch AWS resources in a virtual network that you define. You have complete control over your virtual networking environment, including selection of your own IP address range, creation of subnets, and configuration of route tables and network gateways. You can use both IPv4 and IPv6 in your VPC for secure and easy access to resources and applications.

- Security Groups

Explanation:-This option is not correct. Security Groups are used to control traffic.

Q34) Which of the following are important design principles when architecting cloud-based systems? (Choose two)

- ✅ Build loosely-coupled components

Explanation:-This option is correct. Always design your application components to be loosely coupled. This is to ensure that even if one component fails, the entire system will not fail. Also if you design your system with the assumption that everything will fail, then you will ensure that the right measures are taken to build a highly available and fault tolerant system.

- Build Tightly-coupled components

Explanation:-This option is not correct. You should aim to build loosely-coupled components NOT tightly-coupled components. If you build your architectures with tightly-coupled components, then a small interruption in one component may get your entire system down.

- Use as many services as possible

Explanation:-This option is not correct. You should only use the services that you need.

- ✅ Assume everything will fail

Explanation:-This option is correct. Always design your application components to be loosely coupled. This is to ensure that even if one component fails, the entire system will not fail. Also if you design your system with the assumption that everything will fail, then you will ensure that the right measures are taken to build a highly available and fault tolerant system.

Q35)

There are performance issues with your under-development application, being built using microservices architecture.

Which of the following AWS services would help you analyze these issues?

- AWS Inspector

Explanation:-This option is not correct. Amazon Inspector helps you to identify security vulnerabilities as well as deviations from security best practices in applications NOT for analyzing performance issues.

- AWS Config

Explanation:-This option is not correct. AWS Config is a service that enables you to assess, audit, and evaluate the configurations of your AWS resources.

- ✅ X-Ray

Explanation:-This option is correct. AWS X-Ray helps developers analyze and debug distributed applications in production or under development, such as those built using microservice architecture. With X-Ray, you can understand how your application and its underlying services are performing so you can identify and troubleshoot the root cause of performance issues and errors. X-Ray provides an end-to-end view of requests as they travel through your application.

- AWS CodePipeline

Explanation:-This option is not correct. AWS CodePipeline is a fully managed continuous delivery service that helps you automate your release pipelines for fast and reliable application and infrastructure updates.

Q36) Which of the following AWS offerings is a MySQL-compatible database that has the ability to grow in storage size on its own?

- RDS PostgreSQL

Explanation:-This option is not correct. RDS PostgreSQL is used to run PostgreSQL databases NOT MySQL databases.

✅ Aurora

Explanation:-This option is correct. Amazon Aurora is a MySQL and PostgreSQL compatible relational database built for the cloud, that combines the performance and availability of high-end commercial databases with the simplicity and cost-effectiveness of open source databases. Aurora is up to five times faster than standard MySQL databases and three times faster than standard PostgreSQL databases. It provides the security, availability, and reliability of commercial-grade databases at 1/10th the cost.

⚪ RDS Microsoft SQL Server

Explanation:-This option is not correct. RDS Microsoft SQL Server is used to run SQL Server databases NOT MySQL databases.

⚪ DynamoDB

Explanation:-This option is not correct. DynamoDB is NoSQL database service.

Q37) What are your options for protecting the confidentiality of data in transit in Amazon S3? (Choose two)

⚪ RDS Encryption

Explanation:-This option is not correct. RDS Encryption is related to the Amazon RDS service.

✅ Use SSL

Explanation:-This option is correct. Data protection refers to protecting data while in-transit (as it travels to and from Amazon S3) and at rest (while it is stored on disks in Amazon S3 data centers). You can protect data in transit by using SSL or by using client-side encryption.

⚪ Use Server-Side Encryption

Explanation:-This option is not correct. Server-Side Encryption is an option of protecting data at rest in Amazon S3. Server-Side Encryption involves requesting Amazon S3 to encrypt your object before saving it on disks in its data centers and decrypt it when you download the objects.

✅ Use Client-Side Encryption

Explanation:-This option is correct. Data protection refers to protecting data while in-transit (as it travels to and from Amazon S3) and at rest (while it is stored on disks in Amazon S3 data centers). You can protect data in transit by using SSL or by using client-side encryption.

Q38) Which of the following features of Amazon RDS can you use to improve the availability of your database? (Choose two)

⚪ Automatic patching

Explanation:-This option is not correct. The purpose of patching is to resolve functionality issues, improve security or add new features.

✅ Read Replicas

Explanation:-This option is correct. If you are looking to use replication to increase database availability while protecting your latest database updates against unplanned outages, consider running your DB instance as a Multi-AZ deployment. You can use Multi-AZ deployments and Read Replicas in conjunction to enjoy the complementary benefits of each. You can simply specify that a given Multi-AZ deployment is the source DB instance for your Read Replica(s). That way you gain both the data durability and availability

✅ Multi-AZ Deployment

Explanation:-This option is correct. If you are looking to use replication to increase database availability while protecting your latest database updates against unplanned outages, consider running your DB instance as a Multi-AZ deployment. You can use Multi-AZ deployments and Read Replicas in conjunction to enjoy the complementary benefits of each. You can simply specify that a given Multi-AZ deployment is the source DB instance for your Read Replica(s). That way you gain both the data durability and availability.

⚪ VPC Peering

Explanation:-This option is not correct. VPC peering is not a feature of Amazon RDS. A VPC peering connection is a networking connection between two VPCs that enables you to route traffic between them using private IPv4 addresses or IPv6 addresses. A VPC peering connection helps you to facilitate the transfer of data. For example, if you have more than one AWS account, you can peer the VPCs across those accounts to create a file sharing network. You can also use a VPC peering connection to allow other VPCs.

Q39) How much data can you store in S3?

⚪ You can store up to 1 PetaByte of data, then you are required to pay an additional fee.

Explanation:-This option is incorrect.

⚪ Each account is given 50 gigabytes of storage capacity and no more can be used.

Explanation:-This option is incorrect.

✅ Storage capacity is virtually unlimited.

Explanation:-This option is correct. As an S3 user, there is virtually no limit on the amount of data you can store in S3.

⚪ You can store up to 1 PetaByte of data.

Explanation:-This option is incorrect.

Q40)

A company has infrastructure hosted in an on-premises data center. They currently have an operations team that takes care of ID management.

If they decided to move to the AWS cloud, which of the following services would help them performing the same role in AWS?

⚪ AWS X-Ray

Explanation:-This option is not correct. AWS X-Ray is used to analyze the behavior of your application by providing request tracing, exception collection, and profiling capabilities.

⚪ AWS Cloud Trail

Explanation:-This option is not correct. AWS CloudTrail is used to track user activity on your AWS account. The service provides a history of AWS API calls made in your account, including API calls made via the AWS Management Console, the AWS SDKs, the command line tools, and higher-level AWS services.

⚪ AWS KMS

Explanation:-This option is not correct. AWS Key Management Service (AWS KMS) is a managed service that makes it easy for you to create and control the encryption keys used to encrypt your data.

✅ AWS IAM

Explanation:-This option is correct. AWS Identity and Access Management (IAM) is a web service that helps you securely control access to AWS resources. You use IAM to control who is authenticated (signed in) and authorized (has permissions) to use resources.

Q41)

You need to run a number of Amazon EC2 Instances that are physically isolated at the host hardware level from instances that belong to any other

AWS account. How can you meet this requirement in a cost effective way?

✅ Use EC2 Dedicated Instances
Explanation:-This option is correct. Dedicated Instances are Amazon EC2 instances that run in a VPC on hardware that's dedicated to a single customer. Your Dedicated instances are physically isolated at the host hardware level from instances that belong to other AWS accounts.
⚪ Use EC2 Spot Instances
Explanation:-This option is not correct. Spot and Reserved Instances don't provide physical isolation for its instances.
⚪ Use EC2 Dedicated Hosts
Explanation:-This option is not correct. It is not cost effective compared to the Amazon EC2 Dedicated Instances.
⚪ Use EC2 Reserved Instances
Explanation:-This option is not correct. Spot and Reserved Instances don't provide physical isolation for its instances.

Q42) Which of the following can be attached to EC2 Instances to store data?

⚪ Amazon Glacier
Explanation:-This option is not correct. Amazon Glacier is a storage class for S3 that provides an object level storage for backup and archiving.
⚪ Amazon SQS
Explanation:-This option is not correct. Amazon SQS is not a storage volume. It is a messaging queuing service that can be used to send messages between application components. SQS enables you to decouple and scale microservices, distributed systems, and serverless applications.
⚪ Amazon EBS Snapshots
Explanation:-This option is not correct. Amazon EBS Snapshots are copies (backups) of EBS volumes.
✅ Amazon EBS Volumes
Explanation:-This option is correct. An Amazon EBS volume is a durable, block-level storage device that can be attached to a single EC2 instance. You can use EBS volumes as primary storage for data that requires frequent updates, such as the system drive for an instance or storage for database software.

Q43) What is the advantage of the AWS-recommended practice of decoupling applications?

⚪ Allows tracking of any API call made to any AWS service
⚪ Allows updates of any monolithic application quickly and easily
✅ Reduces inter-dependencies so that failures do not impact other components of the application
Explanation:-As application complexity increases, a desirable attribute of an IT system is that it can be broken into smaller, loosely coupled components. This means that IT systems should be designed in a way that reduces interdependencies—a change or a failure in one component should not cascade to other components. On the other hand if the components of an application are tightly coupled and one component fails, the entire application will also fail. Therefore when designing your application, you should always decouple its components.
⚪ Allows treating an application as a single, cohesive unit

Q44) Which of the following services can be used to build video analytics applications?

✅ Amazon Kinesis
Explanation:-This option is correct. You can use Amazon Kinesis to securely stream video from camera-equipped devices in homes, offices, factories, and public places to AWS. You can then use these video streams for video playback, security monitoring, face detection, machine learning, and other analytics.
⚪ Amazon Athena
Explanation:-This option is not correct. Amazon Athena is an interactive query service that makes it easy to analyze data in Amazon S3 instantly using standard SQL commands.
⚪ Amazon QuickSight
Explanation:-This option is not correct. Amazon QuickSight is used to deliver insights quickly to everyone in your organization.
⚪ Amazon S3
Explanation:-This option is not correct. Amazon S3 is a storage service.

Q45) What does the term "Economies of scale" mean?

✅ It means that AWS will continuously lower costs as it grows.
Explanation:-This option is correct. By using cloud computing, you can achieve a lower variable cost than you would get on your own. Because usage from hundreds of thousands of customers is aggregated in the cloud, providers such as AWS can achieve higher economies of scale, which translates into lower pay as-you-go prices.
⚪ It means that you have the ability to pay as you go.
Explanation:-This option is not correct.
⚪ It means that you save more when you consume more.
Explanation:-This option is not correct.
⚪ It means as more time passes using AWS, you pay more for its services.
Explanation:-This option is not correct.

Q46) What is a placement group in Amazon EC2?

⚪ It is a group used to span multiple Availability Zones.
Explanation:-This option is incorrect.
⚪ It is a group of IAM users that are granted to use EC2.
Explanation:-This option is incorrect.
✅ It is a group of EC2 instances within a single Availability Zone.
Explanation:-This option is correct. Placement Groups are logical groupings or clusters of instances within a single Availability Zone. Placement groups are specifically used for launching cluster of compute instance types. Placement groups are recommended for applications that benefit from

low network latency, high network throughput, or both.
- It is a group of network components that helps protect your traffic.

Explanation:-This option is incorrect.

Q47) Which of the following is NOT a factor when estimating the cost of Amazon EC2?

- Elastic Load Balancing.

Explanation:-This option is not correct.
- ✅ Number of security groups

Explanation:-This option is correct. When you begin to estimate the cost of using Amazon EC2, consider the following:
Clock hours of server time: Resources incur charges when they are running—for example, from the time Amazon EC2 instances are launched until they are terminated, or from the time Elastic IPs are allocated until the time they are de-allocated.
Instance type: Amazon EC2 provides a wide selection of instance types optimized to fit different use cases.
- Number of instances

Explanation:-This option is incorrect.
- Clock hours of server time

Explanation:-This option is incorrect.
- Elastic IP Addresses

Explanation:-This option is incorrect.

Q48) Which of the following services can be used to monitor the HTTP and HTTPS requests that are forwarded to Amazon CloudFront?

- ✅ AWS WAF

Explanation:-This option is correct. AWS WAF is a web application firewall that lets you monitor the HTTP and HTTPS requests that are forwarded to Amazon CloudFront or an Application Load Balancer. AWS WAF also lets you control access to your content by defining customizable web security rules.
- NAT Gateways

Explanation:-This option is not correct. NAT Gateways are not for monitoring. They are used in a VPC to enable instances in a private subnet to connect to the internet or other AWS services, but prevent the internet from initiating a connection with those instances.
- AWS CloudWatch

Explanation:-This option is not correct. CloudWatch is used to monitor AWS resources utilization.
- AWS CloudTrail

Explanation:-This option is not correct. CloudTrail is used to monitor the users' API calls to AWS services.

Q49)

A company has a web application that is running on a number of Amazon EC2 instances.

The app is approaching 100% CPU Utilization on one of them.

How can they reduce the load on that instance?

- Recreate the app to handle such huge traffic.

Explanation:-This option is not correct. Large spikes in traffic cannot be handled by editing the application code. You can handle spikes in traffic by balancing load across the instances and create an Auto Scaling system.
- Use a CloudFront distribution.

Explanation:-This option is not correct. CloudFront is used to distribute content to global users with low latency.
- Terminate the instance and recreate new one.

Explanation:-This option is not correct. The problem is that the traffic is not balanced across the running instances. Terminating the instance and recreating a new one will not solve the problem.
- ✅ Create a load balancer, and register the Amazon EC2 instances with it.

Explanation:-This option is correct. Elastic Load Balancing automatically distributes incoming application traffic across multiple targets, such as Amazon EC2 instances, containers, and IP addresses. Once you register the Amazon EC2 instances with a Load Balancer, it will automatically distribute the incoming traffic across those instances. The Load Balancer also continues to perform health checks on the instances and route traffic only to the healthy ones.

Q50)

You manage a blog on AWS that has different stages such as development, testing, and production.

How can you create a custom console in each stage to view and manage your resources easily?

- AWS Service Groups

Explanation:-This option is not correct. There is nothing called "AWS Service Groups".
- AWS Management Console

Explanation:-This option is not correct. AWS Management Console lets you access and manage individual AWS resources through a web-based user interface.
- AWS Tag Editor

Explanation:-This option is not correct. AWS Tag Editor is used to add tags to multiple resources at once.
- ✅ AWS Resource Groups

Explanation:-This option is correct. If you work with multiple resources in multiple stages, you might find it useful to manage all the resources in each stage as a group rather than move from one AWS service to another for each task. Resource Groups help you do just that. By default, the AWS Management Console is organized by AWS service. But with the Resource Groups tool, you can create a custom console that organizes and consolidates information based on your project and the resources that you use.

Q51) Which DynamoDB feature can be used to reduce the latency of requests to a database from milliseconds to microseconds?

Multi-AZ
Explanation:-This options is are not correct. Multi-AZ and Read Replica are Amazon RDS features.
Read Replica
Explanation:-This options is are not correct. Multi-AZ and Read Replica are Amazon RDS features.
DAX
Explanation:-This option is correct. Amazon DynamoDB Accelerator (DAX) is a fully managed, highly available, in-memory cache for DynamoDB that delivers performance improvements from milliseconds to microseconds – even at millions of requests per second. DAX adds in-memory acceleration to your DynamoDB tables without requiring you to manage cache invalidation, data population, or cluster management.
Memcached.
Explanation:-This option is not correct. Memcached is an in-memory key-value store, originally intended for caching. The AWS offering for Memcached is Amazon ElastiCache Not DynamoDB.

Q52) Which of the following is not an AWS reservation model?

S3 Reserved Capacity
Explanation:-This option is correct. There are no reservations in S3. You pay for what you use. While the cloud is well-suited for running variable workloads and rapid deployments, many cloud-based workloads display a more predictable pattern. For these stable applications, organizations can achieve significant cost savings by taking advantage of the available reservation models such as EC2 reserved instances, RDS reserved instances, ElastiCache Reserved Nodes, DynamoDB Reserved Capacity and Redshift Reserve.
Redshift Reserved Nodes
Explanation:-This option is not correct.
DynamoDB Reserved Capacity
Explanation:-This option is not correct.
ElastiCache Reserved Nodes
Explanation:-This option is not correct.

Q53) Which of the following storage classes is most appropriate to be used for dynamic websites with predictable access patterns?

S3 Glacier
Explanation:-This option is not correct. S3 Glacier is for data archiving and would not be appropriate to use with dynamic websites.
S3 Standard
Explanation:-This option is correct. S3 Standard offers high durability, availability, and performance object storage for frequently accessed data. Because it delivers low latency and high throughput, S3 Standard is appropriate for a wide variety of use cases, including cloud applications, dynamic websites, content distribution, mobile and gaming applications, and big data analytics. S3 Storage Classes can be configured at the object level and a single bucket can contain objects stored across S3 Standard.
S3 Standard-IA
Explanation:-This option is not correct. S3 Standard-IA is not for dynamic websites. S3 Standard-IA is for data that is accessed less frequently, but requires rapid access when needed. S3 Standard-IA is ideal for long-term storage, backups, and as a data store for disaster recovery files.
S3 Intelligent-Tiering
Explanation:-This option is not correct. S3 Intelligent-Tiering is the ideal storage class for long-lived data with access patterns that are unknown or unpredictable. It is designed to optimize costs by automatically moving data to the most cost-effective access tier (Standard and Standard-IA), without performance impact or operational overhead.

Q54) How are S3 storage classes rated?

Scalability & Accessibility.
Explanation:-This option is not correct. All S3 storage classes provide access to your S3 buckets. All S3 storage classes provide the same level of storage scalability. S3 will scale to store any amount of data from anywhere.
Availability & Durability.
Explanation:-This option is correct. Each S3 storage class is rated on its availability and durability.
Availability & Scalability.
Explanation:-This option is not correct. All S3 storage classes provide the same level of storage scalability. S3 will scale to store any amount of data from anywhere.
Durability & Accessibility.
Explanation:-This option is not correct. All s3 storage classes provide access to your S3 buckets.

Q55) What is the document that provides a formal statement of one or more permissions?

Policy
Explanation:-This option is correct. The policy is a JSON document that consists of:
>> Actions: what actions you will allow. Each AWS service has its own set of actions.
>> Resources: which resources you allow the action on.
>> Effect: what will be the effect when the user requests access—either allow or deny.
>> Conditions – which conditions must be present for the policy to take effect. For example, you might allow access only to the specific S3 buckets
Resource
Explanation:-This option is not correct. In AWS, a resource is an entity that you can work with. Examples include an Amazon EC2 instance, an AWS CloudFormation stack, and an Amazon S3 bucket.
Role
Explanation:-This option is not correct. An IAM role is an IAM identity that you can create in your account that has specific permissions. When you assume a role, it provides you with temporary security credentials for your role session. You can use roles to delegate access to users, applications, or services that don't normally have access to your AWS resources. For example, you might want to grant users in your AWS account access to resources they don't usually have, or grant users in one AWS account access.
Permission
Explanation:-This option is not correct. Permissions are granted to IAM entities (users, groups, and roles) to determine whether they are authorized to perform an action or not.

Q56) You want to transfer 200 Terabytes of data from on-premises locations to the AWS Cloud, which of the following can do the job in a cost effective way?

- AWS DMS

Explanation:-This option is not correct. AWS Database Migration Service (DMS) is used to migrate databases to AWS.

- AWS SnowMobile.

Explanation:-This option is not correct. SnowMobile is not a cost effective solution here. AWS Snowmobile is an Exabyte-scale data transfer service used to move extremely large amounts of data to AWS. You can transfer up to 100PB per Snowmobile, a 45-foot long ruggedized shipping container, pulled by a semi-trailer truck.

- ✅ AWS Snowball

Explanation:-This option is correct. Snowball is a petabyte-scale data transport solution that uses secure appliances to transfer large amounts of data into and out of the AWS cloud. Using Snowball addresses common challenges with large-scale data transfers including high network costs, long transfer times, and security concerns. Transferring data with Snowball is simple, fast, secure, and can cost as little as one-fifth the cost of using high-speed Internet.

- AWS Import/Export

Explanation:-This option is not correct. AWS Import/Export has been replaced by the Snowball service.

Q57)

Your web application currently faces performance issues and suffers from long delays.

Which of the following could help you in this situation?

- Amazon Aurora

Explanation:-This option is not correct. Amazon Aurora is a database service.

- AWS Shield

Explanation:-This option is not correct. AWS Shield is a managed Distributed Denial of Service (DDoS) protection service that safeguards applications running on AWS.

- ✅ AWS X-Ray

Explanation:-This option is correct. AWS X-Ray helps you identify performance bottlenecks. X-Ray's service maps let you see relationships between services and resources in your application in real time. You can easily detect where high latencies are occurring, visualize node and edge latency distribution for services, and then drill down into the specific services and paths impacting application performance.

- AWS OpsWorks

Explanation:-This option is not correct. AWS OpsWorks is used to automate operations with Chef and Puppet.

Q58)

You are sure that your application deployed in AWS needs frequent updates for the next 6 months.

Which of the following services would allow you to make these updates easily, retaining the ability to change the AWS resources powering the application any time?

- AWS Simple Storage Service.

Explanation:-This options is not correct. S3 and EFS are storage services.

- AWS Elastic File System.

Explanation:-This options is not correct. S3 and EFS are storage services.

- ✅ AWS Elastic Beanstalk.

Explanation:-This option is correct. AWS Elastic Beanstalk is considered a Platform as a Service (PaaS). it is an easy-to-use service for deploying, scaling and updating web applications and services developed with Java, .NET, PHP, Node.js, Python, Ruby, Go, and Docker on familiar servers such as Apache, Nginx, Passenger, and IIS. You can simply upload your code and Elastic Beanstalk automatically handles the deployment, from capacity provisioning, load balancing, auto-scaling to application health monitoring.

- AWS CodeCommit.

Explanation:-This option is not correct. AWS CodeCommit is a source code control service that hosts secure Git-based repositories. You can use CodeCommit to securely store anything from source code to binaries, and it works seamlessly with your existing Git tools.

Q59) Why are Serverless Architectures more economical than Server-based Architectures?

- When you reserve serverless capacity, you will get large discounts compared to server reservation.

Explanation:-This option is not correct. There are no reservations when using the Serverless Architectures.

- With Serverless Architectures you have the ability to scale automatically up or down as demand changes.

Explanation:-This option is not correct. With Serverless Architectures, you don't have to worry about scaling the compute capacity. AWS handles that for you but that has no bearing on saving money while you are not changing servers.

- ✅ With the Server-based Architectures, servers continue to run all the time but with the serverless architectures the code runs only when needed.

Explanation:-This option is correct. Serverless architectures can **reduce costs** because you don't have to manage or pay for underutilized servers, or provision redundant infrastructure to implement high availability. For example, you can upload your code to the AWS Lambda compute service, and the service can run the code on your behalf using AWS infrastructure. With AWS Lambda, you are charged for every 100ms your code executes and the number of times your code is triggered.

- Serverless Architectures use new powerful computing devices.

Explanation:-This option is not correct. AWS uses the same devices for both server-based and serverless architectures.

Q60) You are working as a site reliability engineer (SRE), which of the following services helps monitor your applications?

- Amazon CloudHSM

Explanation:-This option is not correct. AWS CloudHSM is a cloud-based hardware security module (HSM) that enables you to easily generate and use your own encryption keys on the AWS Cloud.

- ✅ Amazon CloudWatch

Explanation:-This option is correct. Amazon CloudWatch is a monitoring service for AWS cloud resources and the applications you run on AWS. You can use Amazon CloudWatch to collect and track metrics, collect and monitor log files, set alarms, and automatically react to changes in your AWS resources.

Amazon Elastic MapReduce.

Explanation:-This option is not correct. Amazon Elastic MapReduce (Amazon EMR) provides a managed Hadoop framework that makes it easy, fast, and cost-effective to process vast amounts of data across dynamically scalable Amazon EC2 instances.

Amazon CloudSearch

Explanation:-This option is not correct. Amazon CloudSearch is used to set up, manage, and scale a search solution for your website or application.

Q61) What can you access by visiting the URL: http://status.aws.amazon.com/?

AWS Service Health Dashboard

Explanation:-This option is correct. The AWS Service Health Dashboard publishes AWS' most up-to-the-minute information on service availability. The dashboard provides access to current status and historical data about each and every Amazon Web Service. Just copy the URL to your browser and see the result.

AWS Cloud Monitor

Explanation:-This option is not correct.

Status of your Amazon RDS DB

Explanation:-This option is not correct.

Amazon Cloud Watch

Explanation:-This option is not correct.

Q62) Which of the following services allows you to store your application assets, like images and libraries, along with your code?

None of the above.

Explanation:-This option is not correct.

AWS X-Ray

Explanation:-This option is not correct. AWS X-Ray is a service that collects data about requests that your application serves, and provides tools you can use to view, filter, and gain insights into that data to identify issues and opportunities for optimization.

AWS CodeCommit

Explanation:-This option is correct. AWS CodeCommit is a fully managed source control service that makes it easy for companies to host secure and highly scalable private Git repositories. AWS CodeCommit eliminates the need to operate your own source control system or worry about scaling its infrastructure. You can use AWS CodeCommit to securely store anything from source code to binaries, and it works seamlessly with your existing Git tools.

AWS CodePipeline

Explanation:-This option is not correct. AWS CodePipeline is a fully managed continuous delivery service that helps you automate your release pipelines for fast and reliable application and infrastructure updates.
