

Q1)

For selling a personalized gadget that needs an average of 3-4 days to produce Jack's start-up wants to implement an order fulfilment process with some orders taking up to 6 months he expects 10 orders per day on his first day.

After 6 months 1000 orders per day and 10,000 orders after 12 months.

Orders coming in are checked for consistency men dispatched to your manufacturing plant for production quality control packaging shipment and payment processing Employees may force the process to repeat a step, if the product does not meet the quality standards at any stage of the process.

Order status is notified via email to the customers about and also regarding any critical issues with their orders such as payment failure Jack's case architecture includes AWS Elastic Beanstalk for his website with an RDS MySQL instance for customer data and orders.

While making sure that the emails are delivered reliably how can he implement the order fulfilment process?

- ☐ Use an SQS queue to manage all process tasks Use an Auto Scaling group of EC2 Instances that poll the tasks and execute them. Use SES to send emails to customers.
- ☒ Use SWF with an Auto Scaling group of activity workers and a decider instance in another Auto Scaling group with min/max=1 use SES to send emails to customers.
- ☐ Use SWF with an Auto Scaling group of activity workers and a decider instance in another Auto Scaling group with min/max=1 Use the decider instance to send emails to customers.
- ☐ Add a business process management application to your Elastic Beanstalk app servers and re-use the ROS database for tracking order status use one of the Elastic Beanstalk instances to send emails to customers.

Q2)

Several FTP servers are run by a web design company currently that their 250 customers use to upload and download large graphic files.

To make it more scalable they wish to move this system to AWS, but at the same time they wish to maintain customer privacy and Keep costs to a minimum.

What AWS architecture would you be recommended by you?

- ☐ Create a single S3 bucket with Requester Pays turned on and ask their customers to use an S3 client instead of an FTP client Create a bucket for each customer with a Bucket Policy that permits access only to that one customer.
- ☐ Create an auto-scaling group of FTP servers with a scaling policy to automatically scale in when minimum network traffic on the auto-scaling group is below a given threshold. Load a central list of ftp users from S3 as part of the user Data startup script on each Instance
- ☐ Create a single S3 bucket with Reduced Redundancy Storage turned on and ask their customers to use an S3 client instead of an FTP client Create a bucket for each customer with a Bucket Policy that permits access only to that one customer.
- ☒ ASK their customers to use an S3 client instead of an FTP client. Create a single S3 bucket Create an IAM user for each customer Put the IAM Users in a Group that has an IAM policy that permits access to sub-directories within the bucket via use of the 'username' Policy variable

Q3)

Select the security aspects in AWS which are the customer's responsibility.

- 1.) Security Group and ACL (Access Control List) settings
- 2.) Encryption of EBS (Elastic Block Storage) volumes
- 3.) Patch management on the EC2 instance's operating system
- 4.) Life-cycle management of IAM credentials
- 5.) Controlling physical access to compute resources

- ☒ Only 1, 2, 3, and 4
- ☐ Only 2, 4 and 5
- ☐ Only 3, and 5
- ☐ Only 1, 3, 4 and 5

Q4)

A large amount of aerial image data to S3 is uploaded to Allen's firm. In the past, in his on premises environment, batch process this data and used RabbitMQ he used a dedicated group of servers to, an open source messaging system, to get job information to the servers.

The data would go to tape and be shipped offsite once processed the.

His manager told him to stay with the current design, and leverage AWS archival storage and messaging services to minimize cost.

State the correct for the above.

- ☐ Setup Auto-Scaled workers triggered by queue depth that use spot instances to process messages in SQS. Once data is processed, change the storage class of the S3 objects to Glacier.
- ☐ Setup Auto-Scaled workers triggered by queue depth that use spot instances to process messages in SQS. Once data is processed, change the storage class of the S3 objects to Reduced Redundancy Storage
- ☐ Use SQS for passing job messages, use CloudWatch alarms to terminate EC2 worker instances when they become idle Once data is processed, change the storage class of the S3 objects to Reduced Redundancy Storage

Q5)

In a company that operates a mobile sleep tracking application Joy is an IT architect.

Every 5 minutes to his backend when activated at night, the mobile app is sending collected data points of 1 kilobyte.

The backend takes care of authenticating the user and writing the data points into an Amazon DynamoDB table. On a per user basis every morning, Joy scans the table to aggregate and extract last night's data, and store the results in Amazon S3.

The availability of new Data is notified to the users via Amazon SNS mobile push notifications, which is parsed and visualized by the mobile app.

He has around 100k users currently who are mostly based out of North America.

He has been tasked to optimize the architecture of the backend system to lower costs. Select the recommended option of the following.

- 1.) Have the mobile app access Amazon DynamoDB directly instead of JSON files stored on Amazon S3.**
- 2.) Write data directly into an Amazon Redshift cluster replacing both Amazon DynamoDB and Amazon S3.**
- 3.) Introduce an Amazon SQS queue to buffer writes to the Amazon DynamoDB table and reduce provisioned write throughput.**
- 4.) Introduce Amazon ElastiCache to cache reads from the Amazon DynamoDB table and reduce provisioned read throughput. Create a new Amazon DynamoDB table each day and drop the one for the previous day after its data is on Amazon S3.**

- ☒ Only 1 and 4
☐ Only 3
☐ Only 2 and 3
☐ Only 2, 3 and 4

Q6)

Using AWS an application is going to be developed .

To store important documents the application needs a storage layer.

To fulfill this requirement select the incorrect option of the following.

- ☐ Amazon EBS
☐ Amazon S3
☐ Amazon EFS
☒ Amazon Storage Gateway VTL

Explanation: -It's used to take the data backups to the cloud. For more information, please refer to the below AWS doc link:

<https://aws.amazon.com/storagegateway/vtl/> NOTE: The question is asking about which of the below options is "incorrect" for storing of "important" documents in the cloud, and Option D is correct. The question is not asking about data archival, rather storing. So, Option D is not suited for our requirement.

Q7)

Ryan's company is planning on hosting an e-commerce application on the AWS Cloud.

There is a requirement for sessions to be always maintained for users.

For storing session data which of the following can be used?

- 1.) CloudWatch**
- 2.) DynamoDB**
- 3.) Elastic Load Balancing**
- 4.) ElastiCache**

- ☐ Only 1 and 2
☐ Only 3 and 4

Explanation: -DynamoDB and ElastiCache are perfect options for storing session data. AWS Documentation mentions the following on these services: Amazon DynamoDB is a fast and flexible NoSQL database service for all applications that need consistent, single-digit millisecond latency at any scale. It is a fully managed cloud database and supports both document and key-value store models. Its flexible data model, reliable performance, and automatic scaling of throughput capacity makes it a great fit for mobile

- ☐ Only 2 and 3
☒ Only 2 and 4

Q8)

A company named: 'XPN' needs to have its object-based data stored on AWS. The initial size of data would be around 500 GB, with overall growth expected to go into 80TB over the next couple of months. The solution must also be durable.

Suggest an ideal storage option to use for such a requirement, of the following.

- ☐ Amazon Aurora
☒ Amazon S3

Explanation: -Amazon S3 is object storage built to store and retrieve any amount of data from anywhere – web sites and mobile apps, corporate applications, and data from IoT sensors or devices. It is designed to deliver 99.999999999% durability, and stores data for millions of applications used by market leaders in every industry. S3 provides comprehensive security and compliance capabilities that meet even the most stringent

regulatory requirements. It gives customers flexibility in the way they manage data

- DynamoDB
- Amazon Redshift

Q9)

A company 'XNP' has decided to host a MongoDB database on an EC2 Instance. There is an expectancy of a large number of reads and writes on the database.

To implement for the database suggest the ideal EBS storage types of the following.

- EBS General Purpose SSD
- EBS Throughput Optimized HDD
- ✔ EBS Provisioned IOPS SSD
- EBS Cold HDD

Q10)

'PNX' is a Company who is planning to use AWS Simple Storage Service for hosting their project documents. At the end of the project, the documents need to be moved to archival storage.

Out of the following suggest the implementation steps that would ensure the documents are managed accordingly.

- Creating an IAM policy for the S3 bucket
- ✔ Configuring lifecycle configuration rules on the S3 bucket

Explanation:-The AWS Documentation mentions the following on lifecycle policies. Lifecycle configuration enables you to specify the lifecycle management of objects in a bucket. The configuration is a set of one or more rules, where each rule defines an action for Amazon S3 to apply to a group of objects. These actions can be classified as follows: Transition actions – In which you define when objects transition to another storage class. For example, you may choose to transition objects to the STANDARD_IA (

- Adding a bucket policy on the S3 bucket
 - Enabling CORS on the S3 bucket
-