Answer Sheet

Q1)

A user has created a queue named "myqueue" with SQS. There are four messages published to queue which are not received by the consumer yet.

If the user tries to delete the queue, what will happen?

- A user can never delete a queue manually. AWS deletes it after 30 days of inactivity on queue
- It will initiate the delete but wait for four days before deleting until all messages are deleted automatically.
- It will ask user to delete the messages first
- ✅ It will delete the queue

Explanation:-SQS allows the user to move data between distributed components of applications so they can perform different tasks without losing messages or requiring each component to be always available. The user can delete a queue at any time, whether it is empty or not. It is important to note that queues retain messages for a set period of time. By default, a queue retains messages for four days.

Q2)

A user has hosted a website on AWS and uses ELB to load balance the multiple instances.

The user application does not have any cookie management.

How can the user bind the session of the requestor with a particular instance?

- Use session synchronization with ELB
- ✅ Let ELB generate a cookie for a specified duration

Explanation:-The key to manage the sticky session is determining how long the load balancer should route the user's request to the same application instance. If the application has its own session cookie, then the user can set the Elastic Load Balancing to create the session cookie to follow the duration specified by the application's session cookie. If the user's application does not have its own session cookie, then he can set the Elastic Load Balancing to create a session cookie by specifying his own stickiness duration.

- Bind the IP address with a sticky cookie
- Create a cookie at the application level to set at ELB

Q3) To analyze performance metrics of Amazon DynamoDB by using _____.

- CloudWatch's own command-line interface
- Programmatically using the CloudWatch API
- ✅ All the other answers

Explanation:-Amazon DynamoDB and Amazon CloudWatch are integrated, so you can gather and analyze performance metrics. You can monitor these metrics using the CloudWatch console, CloudWatch's own command-line interface, or programmatically using the CloudWatch API.

- The CloudWatch console

Q4)

A user is creating multiple IAM users.

What advice should be given to him to enhance the security?

- Grant less privileges for user, but higher privileges for the group
- Grant all higher privileges to the group
- ✅ Grant least privileges to the individual user

Explanation:-It is a recommended rule that the root user should grant the least privileges to the IAM user or the group. The higher the privileges, the more problems it can create.

- Grant more privileges to the user, but least privileges to the group

---

Q5)

A user plans to use RDS as a managed DB platform.

Which of the below mentioned features is not supported by RDS?

- Automated backup
- ✅ Automated scaling to manage a higher load

Explanation:-AWS RDS provides a managed DB platform, which offers features, such as automated backup, patch management, automated failure detection and recovery. The scaling is not automated and the user needs to plan it with a few clicks.

- Automated software patching
- Automated failure detection and recovery

Q6)

Cloud Academy has three AWS accounts. They have created separate IAM users within each account.

Cloud Academy wants a single IAM console URL such as https://cloudacademy.signin.aws.amazon.com/console/ for all account users.

How can this be achieved?

- Create the S3 bucket with an alias name and use the redirect rule to forward requests to various accounts

- Merge all the accounts with consolidated billing
- Create the same account alias with each account ID
- ✅ It is not possible to have the same IAM account login URL for separate AWS accounts

Explanation:-If a user wants the URL of the AWS IAM sign-in page to have a company name instead of the AWS account ID, he can create an alias for his AWS account ID. The alias should be unique.

Q7) Which one of the following statements is NOT an advantage of DyanamoDB being built on Solid State Drives:

- low request pricing
- serve high-scale request workloads
- low-latency response times
- ✅ high I/O performance of WebApp on EC2 instance

Explanation:-In DynamoDB, SSDs help achieve design goals of predictable low-latency response times for storing and accessing data at any scale. The high I/O performance of SSDs also enables to serve high-scale request workloads cost efficiently, and to pass this efficiency along in low request pricing.

Q8)

A user is running a webserver on EC2. The user wants to receive the SMS when the EC2 instance utilization is above the threshold limit.

Which AWS services should the user configure in this case?

- AWS CloudWatch + AWS SES.
- ✅ AWS CloudWatch + AWS SNS.

Explanation:-Amazon SNS makes it simple and cost-effective to push to mobile devices, such as iPhone, iPad, Android, Kindle Fire, and internet connected smart devices, as well as pushing to other distributed services. In this case, the user can configure that Cloudwatch sends an alarm on when the threshold is crossed to SNS which will trigger an SMS.

- AWS EC2 + AWS Cloudwatch.
- AWS CloudWatch + AWS SQS.

Q9)

A user has enabled automated backup for an RDS instance.

What is the longest duration for which the user can retain the automated backup?

- 25 days
- ✅ 35 days

Explanation:-Amazon RDS provides two different methods for backing up and restoring the Amazon DB instances: automated backups and DB snapshots. Automated backups automatically back up the DB instance during a specific, user-definable backup window, and keep the backups for a limited, user-specified period of time. The maximum period can be 35 days.

- 15 days
- 45 days

Q10)

A user has configured an automated backup between 5 AM – 5:30 AM for the MySQL RDS DB.

Will the performance of RDS get frozen momentarily during a backup?

- Yes, always
- Yes, only if the instance size is smaller than large size
- ✅ Yes, provided it is a single zone implementation

Explanation:-Amazon RDS provides two different methods for backing up and restoring the Amazon DB instances. A brief I/O freeze, typically lasting a few seconds, occurs during both automated backups and DB snapshot operations on Single-AZ DB instances.

- No

Q11)

A user wants to configure AutoScaling which scales up when the CPU utilization is above 70% and scales down when the CPU utilization is below 30%.

How can the user configure AutoScaling for the above mentioned condition?

- Use AutoScaling by manually modifying the desired capacity during a condition
- ✅ Use dynamic AutoScaling with a policy

Explanation:-The user can configure the AutoScaling group to automatically scale up and then scale down based on the specified conditions. To configure this, the user must setup policies which will get triggered by the CloudWatch alarms.

- Use AutoScaling with a schedule
- Configure ELB to notify AutoScaling on load increase or decrease

---

Q12) In regard to DynamoDB, what is the Global secondary index?

- An index that has the same hash key as the table, but a different range key
- An index that has the same hash key and range key as the table
- ✅ An index with a hash and range key that can be different from those on the table.

Explanation:-Global secondary index — an index with a hash and range key that can be different from those on the table.

- An index that has the same range key as the table, but a different hash key

---

Q13) Amazon SNS can also deliver notifications by SMS text message or email, to_____.

- Amazon SES
- ✅ Amazon SQS

Explanation:-In Amazon SNS, besides pushing cloud notifications directly to mobile devices, Amazon SNS can also deliver notifications by SMS text message or email, to Amazon Simple Queue Service (SQS) queues, or to any HTTP endpoint.

- Amazon FPS
- Amazon Glacier

Q14)

Cloud Academy (with AWS account ID 111122223333) has created 50 IAM users for its organization's employees.

Cloud Academy wants to make the AWS console login URL for all IAM users like: https://cloudacademy.signin.aws.amazon.com/console/.

How can this be configured?

- It is not possible to have a personalized IAM login URL
- ✅ Create an IAM AWS account alias with the name cloudacademy

Explanation:-If a user wants the URL of the AWS IAM sign-in page to have a company name instead of the AWS account ID, he can create an alias for his AWS account ID.

- Create an IAM hosted zone Identity for the domain cloudacademy
- The user needs to use Route 53 to map the cloudacademy domain and IAM URL

Q15)

CloudAcademy (with AWS account ID 111122223333) has created 50 IAM users for its organization's employees.

CloudAcademy wants to make the AWS console login URL for all IAM users as: https://CloudAcademy.signin.aws.amazon.com/console/.

How can this be configured?

- Create a bucket with the name CloudAcademy and map it with the IAM alias
- For the AWS account, create an alias CloudAcademy for the IAM login
- ✅ It is not possible to have capital letters as a part of the alias name

Explanation:-If a user wants the URL of the AWS IAM sign-in page to have the company name instead of the AWS account ID, he can create an alias for his AWS account ID. The alias must be unique across all Amazon Webservices products and contain only digits, lowercase letters, and hyphens.

- The user needs to use Route 53 to map the CloudAcademy domain and IAM URL

Q16) In regard to AWS CloudFormation, the template is _____, text-based file that describes all the AWS resources you need to deploy to run your application.

- A XML-format
- ✅ A JSON-format

Explanation:-The template is a JSON-format, text-based file that describes all the AWS resources you need to deploy to run your application.

- A HTML-format
- A CSV-format

Q17)

A CloudFormation template successfully runs in the region in which it was developed.

Upon running the CloudFormation template in another region, it creates some resources then fails when it builds EC2 instances.

What could be the problem?

- CloudFormation templates are unique to the region; a new template must be created for the other region.
- ✅ IDs are unique to the region, and the template likely calls AMI ID's, Security Group ID's or others statically. Replace static ID's with Parameters, Mappings, and Conditions where appropriate.

Explanation:-All ID's are unique to each region, account, and VPC. It is best practice to not embed such ID's inside a CloudFormation template. Instead define parameters, mappings and conditions to create a dynamic template that could be run across VPC's, Regions or even accounts. "CloudFormation templates are unique to the region; a new template must be created for the other region" is incorrect because CF templates, when properly written, can run anywhere. "IAM Permissions are preventing template execution" is incorrect because the template is running, just failing on certain parts. "A role is not defined in the template for the EC2 instance; a role must be defined" is incorrect because a role is not a required parameter when creating an EC2 instance.

- IAM Permissions are preventing template execution.
- A role is not defined in the template for the EC2 instance; a role must be defined.

Q18)

You would like to use the developer tools such as CodeCommit, CodeDeploy, CodePipeline, and CodeBuild but are having a hard time coordinating the workflow and integration with it all in a CI/CD pipeline as well as external tools.

What would be a good way to solve this?

- Find a partner to help
- ✅ Use AWS CodeStar

Explanation:-AWS CodeStar is a tool which will bring all these together under one umbrella. These services can all be used independently of each other but are common to be integrated in an overall workflow. Though CodePipeline is going to manage the actual CI/CD pipeline logic, it does not necessarily integrate or bring the other tools together under one management capability. AWS CodeStar will do this.”Find a partner to help” could be a solution, it’s not the first choice since AWS CodeStar is designed for this and is free to use. CodeStar will also integrate with third-party tools such as Confluence JIRA.

- Use AWS Cloud9
- AWS CodePipeline is meant for CI/CD workflow management, thus it is the solution to tie these components together

---

Q19)

An IAM user is part of a group that allows an action. The user assumes a role which does not allow the action.

Is the user denied or allowed?

- ✅ Denied

Explanation:-When a role is assumed, an entity “becomes” that role. While assuming the role, that is the set of permissions which are evaluated; every day user permissions are irrelevant at the time the role is assumed. Only a single role can be assumed at a given time, and thus the policies assigned to that role are the only set of permissions evaluated at the time of action evaluation.

- Allowed
- CC
- DD

Q20) Which product in AWS is used to create automated workflows for CI/CD pipelines?

- AWS CodeStar
- ✅ AWS CodePipeline

Explanation:-AWS CodePipeline is for creating CI/CD pipelines, and automated steps/actions along the way. CodePipeline can be used to trigger various actions to happen throughout steps in the process. CodePipeline can run steps in parallel, series, and even work with manual steps such as a manual manager approval step. AWS CodeStar is incorrect because AWS CodeStar is more of a project workspace for an overall view of the CI/CD pipeline but it does not automate the workflow. AWS CodeDeploy is incorrect because CodeDeploy is a product for deploying code, which may be a step in the process but not the process itself. AWS Step Functions is incorrect, because although AWS Step Functions is a workflow engine, it is designed to be used with tools such as Lambda. This tool is not designed with CI/CD pipelines in mind.

- AWS CodeDeploy
- AWS Step Functions

Q21)

An application will be performing 150 writes per second to a DynamoDB table. Each write averages to just under 3 KB.

How many WCUs (Write Capacity Units) will be utilized?

- 150
- 1000
- 300
- ✅ 450

Explanation:-Write Capacity Units (WCUs) are measured per second, in 1 KB increments per write. If a single write is more than 1 KB, then more WCU’s will be required per write. As per this question, approximately a 3 KB write is taking place, and a single write consumes three WCUs. Thus 150 writes * 3 = 450 WCUs.

Q22) In Amazon SWF, the limit for open activity tasks is up to _____ per workflow execution.

- ✅ 1000

Explanation:-In Amazon SWF, the maximum open activity tasks is up to 1,000 per workflow execution.

- 10000
- 100
- 5000

Q23)

A user has attached one RDS security group with 5 RDS instances. The user has changed the ingress rule for the security group.

What will be the initial status of the ingress rule?

- Approving
- Implementing
- It is not possible to assign a single group to multiple DB instances
- ✅ Authorizing

Explanation:-When the user makes any changes to the RDS security group the rule status will be authorizing for some time until the changes are applied to all instances that the group is connected with. Once the changes are propagated the rule status will change to authorized.

Q24) Regarding Amazon SWF, the coordination logic in a workflow is contained in a software program called a _____.

- Cordinator
- Worker

- Handler
- ✅ Decider

Explanation:-In Amazon SWF, the coordination logic in a workflow is contained in a software program called a decider. The decider schedules activity tasks, provides input data to the activity workers, processes events that arrive while the workflow is in progress, and ultimately ends (or closes) the workflow when the objective has been completed.

Q25)

A user has created an RDS instance with MySQL.

The user is using the HeidiSQL client to connect with the RDS DB.

The client is unable to connect to DB from his home machine.

What is a possible reason for the failure?

- The user has to open port 80 in the RDS security group to connect with RDS DNS
- ✅ The security group is not configured to allow a request from the user's IP on port 3306

Explanation:-If the user needs to connect to RDS then he has to open port 3306 in the RDS security group for his IP address.

- The user has to open port 22 in the RDS security group to connect with RDS DNS
- You can never connect to RDS from your desktop

Q26)

An organization has created 10 IAM users. The organization wants those users to work independently and access AWS.

Which of the below mentioned options is not a possible solution?

- Enable MFA for each IAM user and assign them the virtual MFA device to access the console
- Create the X.509 certificate for each user and provide them access to AWS CLI
- Provide each user with the IAM login and password for the AWS console
- ✅ Create the access key and secret access key for each user and provide access to AWS using the console

Explanation:-If an organization has created the IAM users, the users can access AWS services either with an IAM specific login/password or console. The organization can generate the IAM X.509 certificates to access AWS with CLI. The organization can also enable MFA for each IAM user, which allows an added security for each IAM user. If the organization has created the access key and secret key than the user cannot access the console using those keys. Access key and secret access key are useful for CLI or Webservices.

Q27) How long can you keep your Amazon SQS messages in Amazon SQS queues?

- From 120 secs up to 4 weeks
- ✅ From 60 secs up to 2 weeks

Explanation:-The SQS message retention period is configurable and can be set anywhere from 1 minute to 2 weeks. The default is 4 days and once the message retention limit is reached your messages will be automatically deleted. The option for longer message retention provides greater flexibility to allow for longer intervals between message production and consumption.

- From 30 secs up to 1 week
- From 10 secs up to 7 days

Q28) When AutoScaling is launching a new instance based on condition, which of the below mentioned policies will it follow?

- Launch an instance which has the highest load distribution
- ✅ Launch an instance in the AZ with the fewest instances

Explanation:-AutoScaling attempts to distribute instances evenly between the Availability Zones that are enabled for the user's AutoScaling group. Auto Scaling does this by attempting to launch new instances in the Availability Zone with the fewest instances.

- Based on the criteria defined with cross zone Load balancing
- Launch an instance in the AZ which has the highest instances

---

Q29) In regards to VPC, select the correct statement:

- You can associate multiple subnets with the same Route Table, but you can't associate a subnet with only one Route Table.
- ✅ You can associate multiple subnets with the same Route Table.

Explanation:-Every subnet in your VPC must be associated with exactly one Route Table. However, multiple subnets can be associated with the same Route Table.

- None of these.
- You can't associate multiple subnets with the same Route Table.

Q30)

A user has launched a MySQL RDS. The user wants to plan for the DR and automate the snapshot.

Which of the below mentioned functionality offers this option with RDS?

- Automated synchronization
- Snapshot
- Copy snapshot
- ✅ Automated backup

Explanation:-Amazon RDS provides two different methods for backing up and restoring the Amazon DB instances: automated backups and DB snapshots. Automated backups automatically back up the DB instance during a specific, user-definable backup window, and keep the backups for a limited, user-specified period of time.

Q31) Can one instance be registered with two ELBs in the same region?

- ✅ Yes, always

Explanation:-Yes, it is possible to have one instance part of two separate ELBs, though both ELBs have different configurations. ELBs are never launched in specific zones.

- No
- Yes, provided both ELBs have the same health check configuration
- Yes, provided both ELBs are in the same AZ

Q32) Regarding AWS Elastic Beanstalk, if the environment health is gray, _____.

- the environment is active and ready to use.
- the environment is ready but it should be activated by user.
- the environment is terminated and application should be reinstalled.
- ✅ the environment is still in the process of being launched.

Explanation:-While AWS Elastic Beanstalk creates your AWS resources and launches your application, the environment will be in a Launching (gray) state. Status messages about launch events are displayed in the environment's dashboard. If the environment health is gray, the environment is still in the process of being launched.

Q33)

Your AWS account is now growing to 200 users and you would like to provide each of these users a personal space in the S3 bucket 'my_company_space' with the prefix /home/, where they have read/write access.

How can you do this efficiently?

- Create one customer managed policy per user and attach them to the relevant users
- Create inline policies for each user as they are on-boarded
- ✅ Create one customer managed policy with dynamic variables and attach it to a group of all users
- Create an S3 bucket policy and change it as users are added and removed

---

Q34)

You would like to deploy a Lambda function globally so that requests are filtered at the AWS edge locations.

Which Lambda deployment mode do you need?

- Deploy Lambda in a Global VPC
- Use a Global DynamoDB table as a Lambda source
- ✅ Use a Lambda@Edge
- Deploy Lambda in S3

Q35)

Your organization has setup a full CICD pipeline leveraging CodePipeline and the deployment is done on Elastic Beanstalk. This pipeline has worked for over a year now but you are approaching the limits of Elastic Beanstalk in terms of how many versions can be stored in the service.

How can you remove older versions that are not used by Elastic Beanstalk so that new versions can be created for your applications?

- ✅ Use a Lifecycle Policy
- Define a Lambda function
- Setup an .ebextensions files
- Use Worker Environments

Q36) In Amazon SWF, _____ represents a logical unit of work that is performed by a component of your application.

- a worker
- ✅ a task

Explanation:-In Amazon SWF, a task represents a logical unit of work that is performed by a component of your application. When using Amazon SWF, you implement workers to perform tasks.

- a method
- a schedule

Q37)

A user has created a MySQL RDS instance with PIOPS.

Which of the below mentioned statements will help user understand the advantage of PIOPS?

- It uses a standard EBS volume with optimized configuration the stacks
- ✅ It uses optimized EBS volumes and optimized configuration stacks

Explanation:-RDS DB instance storage comes in two types: standard and provisioned IOPS. Standard storage is allocated on the Amazon EBS volumes and connected to the user's DB instance. Provisioned IOPS uses optimized EBS volumes and an optimized configuration stack. It provides additional, dedicated capacity for the EBS I/O.

- The user can achieve additional dedicated capacity for the EBS I/O with an enhanced RDS option
- It provides a dedicated network bandwidth between EBS and RDS

Q38)

An organization has 10000 employees. The organization wants to give restricted AWS access to each employee.

How can the organization achieve this?

- Use STS and create the users' run time
- Create an IAM user for each employee and make them a part of the group
- It is not recommended to support 10000 users with IAM
- ✅ Use Identity federation with SSO

Explanation:-Identity federation enables users from an existing directory to access resources within your AWS account, making it easier to manage your users by maintaining their identities in a single place. In this case, the federated user is the only solution since AWS does not allow creating more than 5000 IAM users.

Q39)

A user has enabled serverside encryption with S3. The user downloads the encrypted object from S3.

How can the user decrypt it?

- S3 does not support server side encryption
- ✅ S3 manages encryption and decryption automatically

Explanation:-If the user is using the server-side encryption feature, Amazon S3 encrypts the object data before saving it on disks in its data centres and decrypts it when the user downloads the objects. Thus, the user is free from the tasks of managing encryption, encryption keys, and related tools.

- The user needs to decrypt the object using their own private key
- S3 provides a server side key to decrypt the object

Q40)

A user has created a blank EBS volume in the US-East-1 region.

The user is unable to attach the volume to a running instance in the same region.

What could be the possible reason for this?

- The instance must be in a running state. It is required to stop the instance to attach volume
- ✅ The AZ for the instance and volume are different

Explanation:-An EBS volume provides persistent data storage. The user can attach a volume to any instance provided they are both in the same AZ. Even if they are in the same region but in a different AZ, it will not be able to attach the volume to that instance.

- The instance is from an instance store backed AMI
- The instance has enabled the volume attach protection

Q41) n AWS Elastic Beanstalk, applications can have _____ and each application version _____.

- unique versions, can have the same name of before
- ✅ many versions, is unique

Explanation:-In AWS Elastic Beanstalk, applications can have many versions and each application version is unique. In a running environment, you can deploy any application version you already uploaded to the application or you can upload and immediately deploy a new application version. You might upload multiple application versions to test differences between one version of your web application and another.

- many versions, can refer to many applications
- unique versions, can refer to few applications

Q42) A user is planning to use the AWS RDS with MySQL. Which of the below mentioned services the user is not going to pay?

- Data storage
- ✅ RDS Cloudwatch metrics

Explanation:-RDS charges the user on a pay as you go basis. It charges the user based on the instance type, number of hours that the instance is running, data transfer, storage cost as well for the I/O requests. The monitoring is free of cost.

- Data transfer
- I/O requests per month

Q43) Regarding Amazon SNS, when you want to write a policy, each policy is a _____.

- TXT document
- ✅ JSON document

Explanation:-In Amazon SNS, when you want to write a policy, each policy is a JSON document.

- XML document
- AWS Command

Q44) AWS Elastic Beanstalk will change the health status of a web server environment tier to gray color when:

- Your application hasn't responded to the application health check URL within the last five minutes.
- Your application hasn't responded to the application health check URL within the last one hour.
- ✅ Your application's health status is unknown because status is reported when the application is not in the ready state.

Explanation:-AWS Elastic Beanstalk will change the health status of a web server environment tier to gray color when your application's health status is unknown (because status is reported when the application is not in the ready state).

- AWS Elastic Beanstalk detects other problems with the environment that are known to make the application unavailable

---

Q45)

Cloudacademy (with AWS account ID 111122223333) has created 50 IAM users for its organization's employees.

What will be the AWS console URL for these associates?

- https://signin.aws.amazon.com/111122223333/console/
- https:// signin.aws.amazon.com/console/
- https://signin.aws.amazon.com/console/111122223333/
- ✅ https:// 111122223333.signin.aws.amazon.com/console/

Explanation:-When an organization is using AWS IAM for creating various users and manage their access rights, the IAM user can not use the login URL http://aws.amazon.com/console to access AWS management console. The console login URL for the IAM user will have AWS account ID of that organization to identify the IAM user belongs to particular account. The AWS console login URL for the IAM user will be https:// .signin.aws.amazon.com/console/. In this case it will be https:// 111122223333.signin.aws.amazon.com/console/

Q46) In regard to AWS CloudFormation, what is a stack?

- The set of AWS resources that are created and managed as a template
- The set of AWS templates that are created and managed as a single unit
- ✅ The set of AWS resources that are created and managed as a single unit

Explanation:-A stack is the set of AWS resources that are created and managed as a single unit when AWS CloudFormation initiates a template.

- The set of AWS templates that are created and managed as a template

---

Q47)

An organization has 10 departments. The organization wants to track the AWS usage of each department.

Which of the below mentioned options meets the requirement?

- Setup IAM users for each department and track their usage
- Setup IAM groups for each department and track their usage
- Create separate accounts for each department and track them separately
- ✅ Create separate accounts for each department, but use consolidated billing for payment and tracking

Explanation:-The cost of an IAM user or groups can never be tracked separately for the purpose of billing. The best solution in this case is to create a separate account for each department and use consolidated billing.

Q48) AWS Elastic Load Balancer supports SSL termination.

- True. For specific availability zones only.
- ✅ True. For all regions

Explanation:-You can configure your load balancer in ELB (Elastic Load Balancing) to use a SSL certificate in order to improve your system security.The load balancer uses the certificate to terminate and then decrypt requests before sending them to the back-end instances. Elastic Load Balancing uses AWS Identity and Access Management (IAM) to upload your certificate to your load balancer.

- True. For specific regions only

Q49) A user has configured a bucket S3 to host a static website. What difference will there be when static website hosting is enabled?

- ✅ It will provide the region specific website endpoint

Explanation:-To host a static website, the user needs to configure an Amazon S3 bucket for website hosting and then upload the website contents to the bucket. The website is then available at the region-specific website endpoint of the bucket.

- It will create a new version of the bucket
- It will help the user identify this bucket as the website root to map with the domain
- It will not make any difference, but will help the user to configure the error page

---

Q50) To scale up the AWS resources using manual AutoScaling, which of the below mentioned parameters should the user change?

- as-set-desired-capacity --desired-capacity
- Preferred capacity
- ✅ Desired capacity

Explanation:-The Manual Scaling as part of Auto Scaling allows the user to change the capacity of Auto Scaling group. The user can add / remove EC2 instances on the fly. To execute manual scaling, the user should modify the desired capacity. AutoScaling will adjust instances as per the requirements. If the user is trying to CLI, he can use command

- Current capacity
- Maximum capacity

Q51)

A user has launched one EC2 instance in the US West region.

The user wants to access the RDS instance launched in the US East region from that EC2 instance.

How can the user configure the access for that EC2 instance?

- Open the security group of the US West region in the RDS security group's ingress rule

- It is not possible to access RDS of the US East region from the US West region
- Create an IAM role which has access to RDS and launch an instance in the US West region with it
- ✅ Configure the IP range of the US West region instance as the ingress security rule of RDS

Explanation:-The user cannot authorize an Amazon EC2 security group if it is in a different AWS Region than the RDS DB instance. The user can authorize an IP range or specify an Amazon EC2 security group in the same region that refers to an IP address in another region.

Q52)

An organization has 500 employees. The organization wants to set up AWS access for each department.

Which of the below mentioned options is a possible solution?

- Create IAM users and provide individual permission to each
- ✅ Create IAM groups based on the permission and assign IAM users to the groups

Explanation:-An IAM group is a collection of IAM users. Groups let the user specify permissions for a collection of users, which can make it easier to manage the permissions for those users.

- Create IAM roles based on the permission and assign users to each role
- It is not possible to manage more than 100 IAM users with AWS

Q53)

A user has an S3 object in the US Standard region with the content "colour=red".

The user updates the object with the content as "colour="white".

If the user tries to read the value 1 minute after it was uploaded, what will S3 return?

- It will return an error saying that the object was not found
- It will return "colour=red"
- ✅ It may return either "colour=red" or "colour=white" i.e. any of the value

Explanation:-AWS S3 follows the eventual consistent model in the US Standard Region. Once the object is updated it may return the new value or the old value based on whether all the content is replicated across multiple servers until it becomes consistent (eventual).

- It will return "colour=white"

Q54)

A user has developed an application which is required to send the data to a NoSQL database.

The user wants to decouple the data sending such that the application keeps processing and sending data but does not wait for an acknowledgement of DB.

Which of the below mentioned applications helps in this scenario?

- AWS Simple Workflow
- AWS Simple Query Service
- AWS Simple Notification Service
- ✅ AWS Simple Queue Service

Explanation:-Amazon Simple Queue Service (SQS) is a fast, reliable, scalable, and fully managed message queuing service. SQS provides a simple and cost-effective way to decouple the components of an application. In this case, the user can use AWS SQS to send messages which are received from an application and sent to DB. The application can continue processing data without waiting for any acknowledgement from DB. The user can use SQS to transmit any volume of data without losing messages or requiring other services to always be available.

Q55) Which of the below mentioned options is not a best practice to securely manage the AWS access credentials?

- Create individual IAM users
- Enable MFA for privileged users
- Keep rotating your secure access credentials at regular intervals
- ✅ Create strong access key and secret access key and attach to the root account

Explanation:-It is a recommended approach to avoid using the access and secret access keys of the root account. Thus, do not download or delete it. Instead make the IAM user as powerful as the root account and use its credentials. The user cannot generate their own access and secret access keys as they are always generated by AWS.

Q56)

A user is planning to host data with RDS.

Which of the below mentioned databases is not supported by RDS?

- MS SQL
- ✅ SQLDB

Explanation:-Amazon Relational Database Service (Amazon RDS) is a web service that makes it easier to set up, operate, and scale a relational database in the cloud. AWS RDS supports popular DBs, such as MySQL, PostgreSQL, MS SQL and Oracle. This means that the code, applications, and tools user is already using with existing databases can be used with Amazon RDS too. In short, it is a managed Relation Database offering from AWS which manages backups, software patching, automatic failure detection, and recovery of Database.

- Oracle
- PostgreSQL

Q57) With regard to DynamoDB, increases in throughput level of a table will typically _____.

Take around a few seconds

Take anywhere from a few minutes to a few hours.

Explanation:-In general in DynamoDB, decreases in throughput will take anywhere from a few seconds to a few minutes, while increases in throughput will typically take anywhere from a few minutes to a few hours.

Be immediate.

Take 24 hours.

Q58) Amazon SNS has no upfront costs and you can pay as you go. It costs _____mobile push notifications.

$0.10 to send one million

$1.00 to send one million

Explanation:-In Amazon SNS, It costs $1.00 to send one million mobile push notifications ($0.50 per million publishes, plus $0.50 per million mobile push notification deliveries).

$0.10 to send thousand

$0.01 to send thousand

Q59) What is the maximum size for messages stored in SQS?

64KB

128KB

1024KB

256KB

Explanation:-By default, SQS queues allow you to send the largest supported payload size, currently 256KB. You can choose to specify a limit on how many bytes can be sent per payload, using the MaximumMessageSize attribute of the SetQueueAttributes method.

Q60) What is the data model of DynamoDB?

"Database", a collection of Tables; "Tables", with Keys and one or more Attribute; and "Attribute", with Name and Value.

"Table", a collection of Items; "Items", with Keys and one or more Attribute; and "Attribute", with Name and Value.

Explanation:-The data model of DynamoDB is:
“Table”, a collection of Items;
“Items”, with Keys and one or more Attribute;
“Attribute”, with Name and Value.

"Items", with Keys and one or more Attribute; and "Attribute", with Name and Value.

Since DynamoDB is schema-less, there is no data model.

Q61) When a user is detaching an EBS volume from a running instance and attaching it to a new instance, which of the below mentioned options should be followed to avoid file system damage?

Force Detach the volume to ensure that all the data stays intact

Unmount the volume first

Explanation:-When a user is trying to detach an EBS volume, the user can either terminate the instance or explicitly remove the volume. It is a recommended practice to unmount the volume first to avoid any file system damage.

Stop all the I/O of the volume before processing

Take a snapshot of the volume before detaching

Q62) Regarding Amazon SWF, _____ is a program that receives activity tasks, performs them, and provides results back.

A role handler

An activity worker

Explanation:-In Amazon SWF, an activity worker is a program that receives activity tasks, performs them, and provides results back. Note that the task itself might actually be performed by a person, in which case the person would use the activity worker software for the receipt and disposition of the task. An example might be a statistical analyst, who receives sets of data, analyzes them, and then sends back the analysis.

A coordinator

A task worker

Q63)

A user is trying to create a policy for an IAM user from the AWS console.

Which of the below mentioned options is not available to the user while configuring policy?

Use custom policy to create policy

Use policy simulator to create policy

Explanation:-When a user is trying to create a policy from the AWS console, it will have options such as create policy from templates or use a policy generator. The user can also define a custom policy or chose the option to have no permission. The policy simulator is not available in the console.

Use policy generator to create policy

Assign No permission

Q64)

A user is configuring the HTTPS protocol on a front end ELB and the SSL protocol for the back-end listener in ELB.

What will ELB do?

It will allow you to create the configuration, but ELB will not work as expected

It will not allow you to create this configuration

Explanation:-If a user is configuring HTTPS on the front end and TCP on the back end, ELB will not allow saving these listeners and will respond with the message.

“Load Balancer protocol is an application layer protocol, but instance protocol is not. Both the Load Balancer protocol and the instance protocol should be at the same layer. Please fix.”

- It will allow you to create the configuration, but the instance will not pass the health check
- Receives requests on HTTPS and sends it to the back end instance on SSL

Q65)

A user is trying to share a video file with all his friends.

Which of the below mentioned AWS services will be cheapest and easy to use?

AWS RRS

Explanation:-AWS RRS provides the same functionality as AWS S3, but at a cheaper rate. It is ideally suited for non mission critical applications. It provides less durability than S3, but is a cheaper option.

- AWS S3
- AWS Glacier
- AWS EC2
