

Q1)

A user is running a MySQL RDS instance. The user will not use the DB for the next 3 months.

How can the user save costs?

- ☐ Change the instance size to micro
- ☐ Stop the RDS instance
- ☒ Create a snapshot of RDS to launch in the future and terminate the instance now

Explanation:-The RDS instances unlike the AWS EBS backed instances cannot be stopped or paused. The user needs to take the final snapshot, terminate the instance and launch a new instance in the future from that snapshot.

- ☐ Pause the RDS activities from CLI until it is required in the future

Q2)

A user wants to access RDS from an EC2 instance using IP addresses. Both RDS and EC2 are in the same region, but different AZs.

Which of the below mentioned options help configure that the instance is accessed faster?

- ☐ Configure the Public IP of the instance in RDS security group
- ☐ Security group of EC2 allowed in the RDS security group
- ☐ Configuring the elastic IP of the instance in RDS security group
- ☒ Configure the Private IP of the Instance in RDS security group

Explanation:-If the user is going to specify an IP range in RDS security group, AWS recommends using the private IP address of the Amazon EC2 instance. This provides a more direct network route from the Amazon EC2 instance to the Amazon RDS DB instance, and does not incur network charges for the data sent outside of the Amazon network.

Q3) In regard to AWS CloudFormation, to pass values to your template at runtime you should use _____

- ☐ Mapping
- ☒ Parameters

Explanation:-Optional parameters are listed in the Parameters section. Parameters enable you to pass values to your template at runtime, and can be dereferenced in the Resources and Outputs sections of the template.

- ☐ JSON file
- ☐ Conditions

Q4) In regard to DynamoDB, can I delete local secondary indexes?

- ☐ Yes, if it is a local secondary indexes
- ☒ No

Explanation:-In DynamoDB, an index cannot be modified once it is created.

- ☐ Yes, if it is a Global secondary indexes
- ☐ Yes, if it is a primary hash key index

Q5) In regard to CloudFormation, CloudFormer is itself _____.

- ☐ A JSON file
- ☒ An AWS CloudFormation stack

Explanation:-In AWS CloudFormation, CloudFormer is itself an AWS CloudFormation stack. You run CloudFormer by launching the stack from your AWS environment. It runs on a t1.micro Amazon EC2 instance and requires no other resources.

- ☐ An AWS CloudFormation template
- ☐ An Amazon S3 bucket

Q6) What is the maximum number of tags that a user can assign to an EC2 instance?

- ☐ 5
- ☒ 10

Explanation:-To help manage EC2 instances as well as their usage in a better way, the user can tag the instances. The tags are metadata assigned by the user which consists of a key and a value. One resource can have a maximum of 10 tags.

- ☐ 25
- ☐ 50

Q7) AWS Elastic Beanstalk stores your application files and optionally server log files in_____.

- ☐ Amazon Glacier
- ☒ Amazon S3

Explanation:-AWS Elastic Beanstalk stores your application files and optionally server log files in Amazon S3. If you are using the AWS Management Console, Git, the AWS Toolkit for Visual Studio, or AWS Toolkit for Eclipse, an Amazon S3 bucket will be created in your account for you and the files you upload will be automatically copied from your local client to Amazon S3. Optionally, you may configure Elastic Beanstalk to copy your server log files every hour to Amazon S3. You do this by editing the environment configuration settings.

- ☐ Amazon Storage Gateway
- ☐ Amazon EC2

Q8) Which of the following items are required to allow an application deployed on an EC2 instance to write data to a DynamoDB table? Assume that no security Keys are allowed to be stored on the EC2 instance. (Choose two.)

- ☒ Create an IAM Role that allows write access to the DynamoDB table.
 - ☒ Launch an EC2 Instance with the IAM Role included in the launch configuration.
 - ☐ Add an IAM User to a running EC2 Instance.
 - ☐ Add an IAM Role to a running EC2 instance.
 - ☐ Create an IAM User that allows write access to the DynamoDB table.
 - ☐ Launch an EC2 Instance with the IAM User included in the launch configuration.
-

Q9) What is the maximum number of S3 Buckets available per AWS account?

- ☐ 500 per account
 - ☒ 100 per account
 - ☐ there is no limit
 - ☐ 100 per region
 - ☐ 100 per IAM user
-

Q10)

You attempt to store an object in the US-STANDARD region in Amazon S3, and receive a confirmation that it has been successfully stored.

You then immediately make another API call and attempt to read this object. S3 tells you that the object does not exist.

What could explain this behavior?

- ☐ You exceeded the bucket object limit, and once this limit is raised the object will be visible.
 - ☐ US-STANDARD imposes a 1 second delay before new objects are readable.
 - ☐ Objects in Amazon S3 do not become visible until they are replicated to a second region.
 - ☒ US-STANDARD uses eventual consistency and it can take time for an object to be readable in a bucket
-

Q11)

After launching an instance that you intend to serve as a NAT (Network Address Translation) device in a public subnet you modify your route tables to have the NAT device be the target of internet bound traffic of your private subnet.

When you try and make an outbound connection to the Internet from an instance in the private subnet, you are not successful.

Which of the following steps could resolve the issue?

- ☒ Disabling the Source/Destination Check attribute on the NAT instance
 - ☐ Attaching a second Elastic Network Interface (ENI) to the instance in the private subnet, and placing it in the public subnet
 - ☐ Attaching a second Elastic Network interface (ENI) to the NAT instance, and placing it in the private subnet
 - ☐ Attaching an Elastic IP address to the instance in the private subnet
-

Q12) What item operation allows the retrieval of multiple items from a DynamoDB table in a single API call?

- ☐ GetMultipleItems
 - ☒ BatchGetItem
 - ☐ GetItem
 - ☐ GetItemRange
-

Q13) When uploading an object, what request header can be explicitly specified in a request to Amazon S3 to encrypt object data when saved on the server side?

- ☐ x-amz-security-token
 - ☐ Content-MD5
 - ☐ x-amz-storage-class
 - ☒ x-amz-server-side-encryption
-

Q14) Which of the following services are key/value stores? (Choose three.)

- ☐ Simple Workflow Service
 - ☒ DynamoDB
 - ☐ Simple Notification Service
 - ☒ Amazon ElastiCache
 - ☒ Simple Storage Service
-

Q15) If a message is retrieved from a queue in Amazon SQS, how long is the message inaccessible to other users by default?

- ☒ 30 seconds
 - ☐ forever
 - ☐ 1 day
 - ☐ 1 hour
 - ☐ 0 seconds
-

Q16)

You have an environment that consists of a public subnet using Amazon VPC and 3 instances that are running in this subnet.

These three instances can successfully communicate with other hosts on the Internet.

You launch a fourth instance in the same subnet, using the same AMI and security group configuration you used for the others, but find that this instance cannot be accessed from the Internet.

What should you do to enable internet access?

- ☐ Configure a publically routable IP Address In the host OS of the fourth instance.
- ☐ Modify the routing table for the public subnet
- ☐ Deploy a NAT instance into the public subnet.
- ☒ Assign an Elastic IP address to the fourth instance.

Q17) Which of the following are correct statements with policy evaluation logic in AWS Identity and Access Management? (Choose two.)

- ☐ An explicit deny does not override an explicit allow
- ☒ An explicit allow overrides default deny.
- ☐ An explicit allow overrides an explicit deny
- ☒ By default, all requests are denied
- ☐ By default, all request are allowed

Q18) What happens, by default, when one of the resources in a CloudFormation stack cannot be created?

- ☐ CloudFormation templates are parsed in advance so stack creation is guaranteed to succeed.
- ☐ The stack creation continues, and the final results indicate which steps failed.
- ☒ Previously-created resources are deleted and the stack creation terminates.
- ☐ Previously-created resources are kept but the stack creation terminates.

Q19) Which features can be used to restrict access to data in S3? (Choose two.)

- ☒ Set an S3 ACL on the bucket or the object.
- ☐ Enable IAM Identity Federation.
- ☒ Set an S3 Bucket policy.
- ☐ Use S3 Virtual Hosting
- ☐ Create a CloudFront distribution for the bucket

Q20) Which code snippet below returns the URL of a load balanced web site created in CloudFormation with an AWS::ElasticLoadBalancing::LoadBalancer resource name "ElasticLoad Balancer"?

- ☐ "Fn::Join" : [" ", ["http://", {"Ref" : "ElasticLoadBalancerDNSName"}]]
- ☐ "Fn::Join" : [" ", ["http://", {"Ref" : "ElasticLoadBalancerUrl"}]]
- ☐ "Fn::Join" : [" ", ["http://", {"Fn::GetAtr" : ["ElasticLoadBalancer", "Url"]}]]
- ☒ "Fn::Join" : [" ", ["http://", {"Fn::GetAtr" : ["ElasticLoadBalancer", "DNSName"]}]]

Q21) Which of the following platforms are supported by Elastic Beanstalk? (Choose two.)

- ☐ IBM WebsphereOracle JBoss
- ☒ .NET
- ☒ Apache Tomcat
- ☐ E.Jetty

Q22)

Company D is running their corporate website on Amazon S3 accessed from http/www.companyd.com.

Their marketing team has published new web fonts to a separate S3 bucket accessed by the S3 endpoint https://s3-us-west-1.amazonaws.com/cdfonts.

While testing the new web fonts, Company D recognized the web fonts are being blocked by the browser.

What should Company D do to prevent the web fonts from being blocked by the browser?

- ☐ Add the Content-MD5 header to the request for webfonts in the cdfonts bucket from the website
- ☐ Create a policy on the cdfonts bucket to enable access to everyone
- ☐ Enable versioning on the cdfonts bucket for each web font
- ☒ Configure the cdfonts bucket to allow cross-origin requests by creating a CORS configuration

Q23) When using a large Scan operation in DynamoDB, what technique can be used to minimize the impact of a scan on a table's provisioned throughput?

- ☐ Define a range index on the table
- ☐ Use parallel scans
- ☒ Set a smaller page size for the scan
- ☐ Prewarm the table by updating all items

Q24) In AWS, which security aspects are the customer's responsibility? (Choose four.)

- ☐ Controlling physical access to compute resources
- ☒ Encryption of EBS (Elastic Block Storage) volumes
- ☒ Security Group and ACL (Access Control List) settings
- ☐ Decommissioning storage devices
- ☒ Life-cycle management of IAM credentials
- ☒ Patch management on the EC2 instance's operating system

Q25) Which EC2 API call would you use to retrieve a list of Amazon Machine Images (AMIs)?

- ☐ You cannot retrieve a list of AMIs as there are over 10,000 AMIs
- ☐ GetAMIs
- ☒ DescribeImages
- ☐ DescnbeInstances
- ☐ DescribeAMIs

Q26) EC2 instances are launched from Amazon Machine images (AMIs). A given public AMI can:

- ☐ only be used to launch EC2 instances in the same AWS availability zone as the AMI is stored
- ☒ only be used to launch EC2 instances in the same AWS region as the AMI is stored.
- ☐ only be used to launch EC2 instances in the same country as the AMI is stored.
- ☐ be used to launch EC2 Instances in any AWS region.

Q27) How can software determine the public and private IP addresses of the Amazon EC2 instance that it is running on?

- ☐ Query the local instance userdata.
- ☐ Use ipconfig or ifconfig command.
- ☐ Query the appropriate Amazon CloudWatch metric.
- ☒ Query the local instance metadata.

Q28) What is the development practice used when a failure of the other end occurs, and instead of an endless loop, we stop trying and create a "ping" until the service is back up to resume operations?

- ☐ Bulkheads
- ☒ Circuit breaker

Explanation:-A circuit breaker pattern will allow us to stop trying to make the call until the other end is healthy. A circuit breaker will be "tripped" and a sideline process will simply check for the other end to resume operations. Once the other end is back up and ready to accept requests, the breaker is reset and operations may resume. This ensures we do not have endless loops that are not able to handle outages. Graceful outage management is key to high performance systems. In addition, in AWS, endless loops could have a charge to them if we are being charged for each call.

- ☐ Exponential backoff
- ☐ Try and Except

Q29) What is the maximum size of an S3 bucket?

- ☐ 100TB
- ☐ 5TB
- ☒ There is no maximum

Explanation:-AWS S3 has no maximum storage limit. Any account can store unlimited data. Since we are talking about total bucket storage, there is no maximum. A single object however cannot exceed 5 TB.

- ☐ 1PB

Q30) Which of these services can return inconsistent results immediately after a write?

- ☐ EFS
- ☐ Redshift
- ☒ S3

Explanation:-Only S3 provides eventual consistency. Review the CAP theorem for further explanation on eventual consistency. All other options are strongly consistent.

- ☐ RDS

Q31)

You are supposed to write an application utilizing the AWS SDK that writes log events to S3.

You choose to utilize S3 because of its high resiliency, ability to scale, and its cost effectiveness.

You get word from your manager that the costs of S3 have skyrocketed.

You realize S3 is not priced well for frequent, small writes; however, the business requirements need the data to persist in S3.

What may be a way to optimize for cost and still meet the requirements?

- ☐ Send the events to S3 but utilize the S3-IA Storage Class
- ☒ Send the events to AWS Kinesis Firehose to buffer the writes prior to sending to S3

Explanation:-AWS Kinesis Firehose is a way to handle a large, frequent number of writes, and buffer them prior to a write to S3. Though S3 is ideal for long-term retention, it is not designed to handle frequent, small writes such as streaming data. PUT operations into S3 will be many times more

than GET or LIST (in some cases PUT can cost 10x that of a GET). Though S3 is cost efficient to store the data, we have to buffer the writes to get good cost optimization. "Send the events to S3 but utilize the S3-IA Storage Class" is incorrect because it could actually increase costs as the S3-IA storage class is for infrequent operations and PUTs and GETs will be more than the standard class." Send the events to AWS Redshift to buffer the writes prior to sending to S3" is incorrect because Redshift is a data warehouse and not the right tool for the job. "Nothing can be done in the application; the problem is that S3 is just more expensive than we thought" is incorrect, because the problem is caused by using the wrong tool for data ingestion; in this case, S3 is not the right tool for the job to ingest the data, but it is the right tool to store it persistently.

- ☐ Send the events to AWS Redshift to buffer the writes prior to sending to S3
- ☐ Nothing can be done in the application; the problem is that S3 is just more expensive than we thought

Q32)

You develop a user data script to run on a newly created EC2 instance in a VPC.

The script downloads some files (objects) from S3 to the local EC2 instance.

You find the user data script is running, and has no scripting errors.

The logs on the EC2 instance show this is an access/reachability problem.

What are the possible problems? (Choose 3)

- ☐ VPC Peering to S3 is not configured properly
- ☒ The VPC does not have a gateway to reach S3 (VPC Endpoint Gateway or Internet Gateway)

Explanation:-"VPC Peering to S3 is not configured properly" is incorrect because there is no ability to peer to S3. Peering is used to peer between VPC's but not to connect to AWS services. "The EC2 instance SSH keys are invalid" is incorrect because SSH keys are irrelevant in this situation. No SSH access is required to reach S3, nor to run the user data script. If the EC2 instance cannot communicate with S3, this script will not work. If the VPC blocks the access to S3 (, or if the EC2 instance cannot communicate outbound, the communications will be blocked. And lastly, if the URL is broken, of course, the objects or buckets will not be found.

- ☐ The EC2 instance SSH keys are invalid
- ☒ The Security Group is not allowing outbound communications to S3

Explanation:-"VPC Peering to S3 is not configured properly" is incorrect because there is no ability to peer to S3. Peering is used to peer between VPC's but not to connect to AWS services. "The EC2 instance SSH keys are invalid" is incorrect because SSH keys are irrelevant in this situation. No SSH access is required to reach S3, nor to run the user data script. If the EC2 instance cannot communicate with S3, this script will not work. If the VPC blocks the access to S3 (, or if the EC2 instance cannot communicate outbound, the communications will be blocked. And lastly, if the URL is broken, of course, the objects or buckets will not be found.

- ☒ The S3 URL is inaccurate

Explanation:-"VPC Peering to S3 is not configured properly" is incorrect because there is no ability to peer to S3. Peering is used to peer between VPC's but not to connect to AWS services. "The EC2 instance SSH keys are invalid" is incorrect because SSH keys are irrelevant in this situation. No SSH access is required to reach S3, nor to run the user data script. If the EC2 instance cannot communicate with S3, this script will not work. If the VPC blocks the access to S3 (, or if the EC2 instance cannot communicate outbound, the communications will be blocked. And lastly, if the URL is broken, of course, the objects or buckets will not be found.

Q33) Which of the following are valid arguments for an SNS Publish request? (Choose three.)

- ☒ Message
- ☐ Format
- ☐ Destination
- ☒ Subject
- ☒ TopicArn
- ☐ Language

Q34) How is provisioned throughput affected by the chosen consistency model when reading data from a DynamoDB table?

- ☐ Strongly consistent reads use less throughput than eventually consistent reads
- ☒ Strongly consistent reads use more throughput than eventually consistent reads
- ☐ Strongly consistent reads use the same amount of throughput as eventually consistent reads
- ☐ Strongly consistent reads use variable throughput depending on read activity

Q35) Which of the following are valid SNS delivery transports? (Choose two.)

- ☐ DynamoDB
- ☒ SMS
- ☐ UDP
- ☒ HTTP
- ☐ Named Pipes

Q36) A startup s photo-sharing site is deployed in a VPC. An ELB distributes web traffic across two subnets. ELB session stickiness is configured to use the AWS-generated session cookie, with a session TTL of 5 minutes. The webserver Auto Scaling Group is configured as: min-size=4, max-size=4. The startups preparing for a public launch, by running load-testing software installed on a single EC2 instance running in us-west-2 After 60 minutes of load-testing, the webserver logs show: Which recommendations can help ensure load-testing HTTP requests are evenly distributed across the four webserver? (Choose two.)

- ☐ Configure ELB session stickiness to use the app-specific session cookie.
 - ☐ Configure ELB and Auto Scaling to distribute across us-west-2a and us-west-2c.
 - ☒ Use a 3rd-party load-testing service which offers globally-distributed test clients.
 - ☒ Re-configure the load-testing software to re-resolve DNS for each web request.
 - ☐ Launch and run the load-tester EC2 instance from us-east-1 instead.
-

Q37) Which of the following statements about SWF are true? (Choose three.)

- ☒ SWF uses deciders and workers to complete tasks
- ☐ SWF triggers SNS notifications on task assignment
- ☒ SWF workflow executions can last up to a year
- ☐ SWF requires an S3 bucket for workflow storage
- ☒ SWF tasks are assigned once and never duplicated
- ☐ SWF requires at least 1 EC2 instance per domain

Q38) Which of the following is chosen as the default region when making an API call with an AWS SDK?

- ☐ us-central-1
- ☐ eu-west-1
- ☒ us-east-1
- ☐ us-west-2
- ☐ ap-northeast-1

Q39) How can you secure data at rest on an EBS volume?

- ☐ Create an IAM policy that restricts read and write access to the volume.
- ☐ Encrypt the volume using the S3 server-side encryption service.
- ☒ Use an encrypted file system on top of the EBS volume.
- ☐ Write the data randomly instead of sequentially.
- ☐ Attach the volume to an instance using EC2's SSL interface.

Q40) What is one key difference between an Amazon EBS-backed and an instance-store backed instance?

- ☐ Instance-store backed instances can be stopped and restarted.
- ☐ Auto scaling requires using Amazon EBS-backed instances.
- ☒ Amazon EBS-backed instances can be stopped and restarted
- ☐ Virtual Private Cloud requires EBS backed instances

Q41) An application stores payroll information nightly in DynamoDB for a large number of employees across hundreds of offices. Item attributes consist of individual name, office identifier, and cumulative daily hours. Managers run reports for ranges of names working in their office. One query is. "Return all Items in this office for names starting with A through E". Which table configuration will result in the lowest impact on provisioned throughput for this query?

- ☐ Configure a hash index on the name attribute and no range index
- ☒ Configure the table to have a range index on the name attribute, and a hash index on the office identifier
- ☐ Configure the table to have a hash index on the name attribute, and a range index on the office identifier
- ☐ Configure a hash index on the office Identifier attribute and no range index

Q42) Which of the following is an example of a good DynamoDB hash key schema for provisioned throughput efficiency?

- ☐ Device ID, where one is by far more popular than all the others.
- ☐ Status Code where most status codes are the same
- ☒ User ID, where the application has many different users.
- ☐ Game Type, where there are three possible game types

Q43)

Company A has an S3 bucket containing premier content that they intend to make available to only paid subscribers of their website.

The S3 bucket currently has default permissions of all objects being private to prevent inadvertent exposure of the premier content to non-paying website visitors.

How can Company A provide only paid subscribers the ability to download a premier content file in the S3 bucket?

- ☐ Enable server side encryption on the S3 bucket for data protection against the non-paying website visitors
- ☐ Add a bucket policy that requires Multi-Factor Authentication for requests to access the S3 bucket objects
- ☒ Generate a pre-signed object URL for the premier content file when a paid subscriber requests a download
- ☐ Apply a bucket policy that grants anonymous users to download the content from the S3 bucket

Q44) When a Simple Queue Service message triggers a task that takes 5 minutes to complete, which process below will result in successful processing of the message and remove it from the queue while minimizing the chances of duplicate processing?

- ☐ Retrieve the message with increased DelaySeconds, delete the message from the queue, process the message
- ☐ Retrieve the message with increased DelaySeconds, process the message, delete the message from the queue
- ☐ Retrieve the message with an increased visibility timeout, delete the message from the queue, process the message
- ☒ Retrieve the message with an increased visibility timeout, process the message, delete the message from the queue

Q45) Which DynamoDB limits can be raised by contacting AWS support? (Choose two.)

- ☐ The number of local secondary indexes per account
- ☒ The number of tables per account

- ☐ The maximum storage used per account
 - ☐ The number of hash keys per account
 - ☒ The number of provisioned throughput units per account
-

Q46) In DynamoDB, what type of HTTP response codes indicate that a problem was found with the client request sent to the service ?

- ☐ 306 HTTP response code
 - ☐ 200 HTTP response code
 - ☐ A 5xx HTTP response code
 - ☒ 4xx HTTP response code
-

Q47) Which of the following programming languages have an officially supported AWS SDK?

- ☐ Java
 - ☐ Pascal
 - ☒ PHP
 - ☐ Perl
 - ☐ SQL
-

Q48)

Games-R-Us is launching a new game app for mobile devices. Users will log into the game using their existing Facebook account and the game will record player data and scoring information directly to a DynamoDB table.

What is the most secure approach for signing requests to the DynamoDB API?

- ☐ Establish cross account access between the mobile app and the DynamoDB table to sign the requests
 - ☒ Request temporary security credentials using web identity federation to sign the requests
 - ☐ Distribute the AWS root account access credentials with the mobile app to sign the requests
 - ☐ Create an IAM user with access credentials that are distributed with the mobile app to sign the requests
-

Q49) What AWS products and features can be deployed by Elastic Beanstalk? (Choose three.)

- ☒ RDS Instances
 - ☒ Elastic Load Balancers
 - ☐ Route 53 hosted zones
 - ☒ Auto scaling groups
 - ☐ Elastic IP addresses
 - ☐ SQS Queues
-

Q50)

Your application is trying to upload a 6 GB file to Simple Storage Service and receive a “Your proposed upload exceeds the maximum allowed object size.” error message.

What is a possible solution for this?

- ☐ Contact support to increase your object size limit
 - ☐ Use the large object upload API for this object
 - ☒ Use the multi-part upload API for this object
 - ☐ None, Simple Storage Service objects are limited to 5 GB
 - ☐ Upload to a different region
-

Q51) Which of the following services are included at no additional cost with the use of the AWS platform? (Choose two.)

- ☒ CloudFormation
 - ☐ Elastic Load Balancing
 - ☒ Auto Scaling
 - ☐ Elastic Compute Cloud
 - ☐ Simple Storage Service
 - ☐ Simple Workflow Service
-

Q52)

You are writing to a DynamoDB table and receive the following exception: “ProvisionedThroughputExceededException”. though according to your Cloudwatch metrics for the table, you are not exceeding your provisioned throughput.

What could be an explanation for this?

- ☐ You're exceeding your capacity on a particular Sort Key
 - ☒ You're exceeding your capacity on a particular Hash Key
 - ☐ You're exceeding your capacity on a particular Range Key
 - ☐ You haven't provisioned enough DynamoDB storage instances
 - ☐ You haven't configured DynamoDB Auto Scaling triggers
-

Q53)

You are providing AWS consulting services for a company developing a new mobile application that will be leveraging Amazon SNS Mobile Push for push notifications.

In order to send direct notification messages to individual devices each device registration identifier or token needs to be registered with SNS; however the developers are not sure of the best way to do this.

You advise them to:

- ☐ Implement a token vending service to handle the registration.
- ☐ Let the push notification service (e.g. Amazon Device Messaging) handle the registration.
- ☐ Bulk upload the device tokens contained in a CSV file via the AWS Management Console.
- ☒ Call the CreatePlatformEndPoint API function to register multiple device tokens.

Q54) Which statements about DynamoDB are true? (Choose two.)

- ☐ DynamoDB restricts item access during reads
- ☒ DynamoDB uses conditional writes for consistency
- ☒ DynamoDB uses optimistic concurrency control
- ☐ DynamoDB uses a pessimistic locking model
- ☐ DynamoDB restricts item access during writes

Q55)

You run an ad-supported photo sharing website using S3 to serve photos to visitors of your site.

At some point you find out that other sites have been linking to the photos on your site, causing loss to your business.

What is an effective method to mitigate this?

- ☐ Use CloudFront distributions for static content.
- ☒ Remove public read access and use signed URLs with expiry dates.
- ☐ Store photos on an EBS volume of the web server
- ☐ Block the IPs of the offending websites in Security Groups.

Q56)

You have written an application that uses the Elastic Load Balancing service to spread traffic to several web servers.

Your users complain that they are sometimes forced to login again in the middle of using your application, after they have already logged in.

This is not behavior you have designed.

What is a possible solution to prevent this happening?

- ☐ Use EBS to save session state
- ☐ Use instance storage to save session state.
- ☐ Use instance memory to save session state.
- ☒ Use ElastiCache to save session state.
- ☐ Use Glacier to save session slate.

Q57)

You are inserting 1000 new items every second in a DynamoDB table.

Once an hour these items are analyzed and then are no longer needed.

You need to minimize provisioned throughput, storage, and API calls.

Given these requirements, what is the most efficient way to manage these Items after the analysis?

- ☒ Delete the table and create a new table per hour
- ☐ Delete items individually over a 24 hour period
- ☐ Retain the items in a single table
- ☐ Create a new table per hour

Q58) An Amazon S3 bucket, "myawsbucket" is configured with website hosting in Tokyo region, what is the regionspecific website endpoint?

- ☐ myawsbucket.amazonaws.com
- ☒ myawsbucket.s3-website-ap-northeast-1.amazonaws.com
- ☐ www.myawsbucket.ap-northeast-1.amazonaws.com
- ☐ myawsbucket.tokyo.amazonaws.com

Q59)

Company B provides an online image recognition service and utilizes SQS to decouple system components for scalability The SQS consumers poll the imaging queue as often as possible to keep end-to-end throughput as high as possible.

However, Company B is realizing that polling in tight loops is burning CPU cycles and increasing costs with empty responses.

How can Company B reduce the number of empty responses?

- ☐ Set the imaging queue MessageRetentionPeriod attribute to 20 seconds
- ☒ Set the Imaging queue ReceiveMessageWaitTimeSeconds attribute to 20 seconds
- ☐ Set the imaging queue visibility Timeout attribute to 20 seconds
- ☐ Set the DelaySeconds parameter of a message to 20 seconds

Q60)

A corporate web application is deployed within an Amazon VPC, and is connected to the corporate data center via IPSec VPN.

The application must authenticate against the on-premise LDAP server.

Once authenticated, logged-in users can only access an S3 keyspace specific to the user.

Which two approaches can satisfy the objectives? (Choose two.)

- ☐ Develop an identity broker which authenticates against IAM Security Token Service to assume an IAM Role to get temporary AWS security credentials. The application calls the identity broker to get AWS temporary security credentials with access to the appropriate S3 bucket.
- ☒ Develop an identity broker which authenticates against LDAP, and then calls IAM Security Token Service to get IAM federated user credentials. The application calls the identity broker to get IAM federated user credentials with access to the appropriate S3 bucket.
- ☐ The application authenticates against IAM Security Token Service using the LDAP credentials. The application uses those temporary AWS security credentials to access the appropriate S3 bucket.
- ☒ The application authenticates against LDAP, and retrieves the name of an IAM role associated with the user. The application then calls the IAM Security Token Service to assume that IAM Role. The application can use the temporary credentials to access the appropriate S3 bucket.
- ☐ The application authenticates against LDAP. The application then calls the IAM Security Service to login to IAM using the LDAP credentials. The application can use the IAM temporary credentials to access the appropriate S3 bucket.

Q61) Which of the following statements about SQS is Correct?

- ☐ Messages will be delivered one or more times and messages will be delivered in First in, First out order
- ☐ Messages will be delivered exactly once and message delivery order is indeterminate
- ☐ Messages will be delivered exactly once and messages will be delivered in First in, First out order
- ☒ Messages will be delivered one or more times and message delivery order is indeterminate

Q62) If an application is storing hourly log files from thousands of instances from a high traffic web site, which naming scheme would give optimal performance on S3?

- ☐ YYYY-MM-DD-HH-log_instanceID
- ☒ HH-DD-MM-YYYY-log_instanceID
- ☐ instanceID_log-YYYY-MM-DD-HH
- ☐ instanceID_log-HH-DD-MM-YYYY
- ☐ Sequential

Q63) What type of block cipher does Amazon S3 offer for server side encryption?

- ☐ Blowfish
- ☒ Advanced Encryption Standard
- ☐ Triple DES
- ☐ RC5

Q64) In regard to DynamoDB, for which one of the following parameters does Amazon not charge you?

- ☒ I/O usage within the same Region

Explanation:-In DynamoDB, you will be charged for the storage and the throughput you use rather than for the I/O which has been used.

- ☐ Storage cost
- ☐ Cost per provisioned write units
- ☐ Cost per provisioned read units