

Q1) Which of the following services can be used in conjunction with Cloud watch Logs. Choose the 3 most viable services from the options given below

- ☒ Amazon S3
- ☐ Amazon Kinesis
- ☒ Amazon Lambda
- ☐ Amazon SNS

Q2)

You have a requirement to automate the creation of EBS Snapshots.

Which of the following can be used to achieve this in the best way possible.

- ☒ Use Cloud watch Events to trigger the snapshots of EBS Volumes
- ☐ Use the AWS Code Deploy service to create a snapshot of the AWS Volumes
- ☐ Create a power shell script which uses the AWS CLI to get the volumes and then run the script as a cronjob
- ☐ Use the AWS Config service to create a snapshot of the AWS Volumes

Q3)

Your company has the requirement to set up instances running as part of an Auto scaling Group. Part of the requirement is to use Lifecycle hooks to setup custom based software and do the necessary configuration on the Instances.

The time required for this setup might take an hour, or might finish before the hour is up.

How should you setup lifecycle hooks for the Auto scaling Group. Choose 2 ideal actions you would include as part of the lifecycle hook.

- ☐ Configure the lifecycle hook to record heartbeats. If the hour is up, choose to terminate the current instance and start a new one
- ☒ Configure the lifecycle hook to record heartbeats. If the hour is up, restart the timeout period. ?
- ☐ If the software installation and configuration is complete, then restart the time period.
- ☒ If the software installation and configuration is complete, then send a signal to complete the launch of the instance.

Q4)

You have a large number of web servers in an Auto Scaling group behind a load balancer. On an hourly basis, you want to filter and process the logs to collect data on unique visitors, and then put that data in a durable data store in order to run reports. Web servers in the Auto Scaling group are constantly launching and terminating based on your scaling policies, but you do not want to lose any of the log data from these servers during a stop/termination initiated by a user or by Auto Scaling.

What two approaches will meet these requirements? Choose two answers from the options given below?

- ☐ Install an AWS Data Pipeline Logs Agent on every web server during the bootstrap process. Create a log group object in AWS Data Pipeline, and define Metric Filters to move processed log data directly from the web servers to Amazon Redshift and run reports every hour. .
- ☐ On the web servers, create a scheduled task that executes a script to rotate and transmit the logs to Amazon Glacier. Ensure that the operating system shutdown procedure triggers a logs transmission when the Amazon EC2 instance is stopped/terminated. Use Amazon Data Pipeline to process the data in Amazon Glacier and run reports every hour.
- ☒ Install an Amazon CloudWatch Logs Agent on every web server during the bootstrap process. Create a CloudWatch log group and define Metric Filters to create custom metrics that track unique visitors from the streaming web server logs. Create a scheduled task on an Amazon EC2 instance that runs every hour to generate a new report based on the CloudWatch custom metrics.
- ☒ On the web servers, create a scheduled task that executes a script to rotate and transmit the logs to an Amazon S3 bucket. Ensure that the operating system shutdown procedure triggers a logs transmission when the Amazon EC2 instance is stopped/terminated. Use AWS Data Pipeline to move log data from the Amazon S3 bucket to Amazon Redshift in order to process and run reports every hour.

Q5)

You have an Auto scaling Group configured to launch EC2 Instances for your application. But you notice that the Auto scaling Group is not launching instances in the right proportion. In fact instances are being launched too fast.

What can you do to mitigate this issue? Choose 2 answers from the options given below Please select:

- ☒ Set a custom metric which monitors a key application functionality for the scale-in and scale-out process.
- ☒ Adjust the cool down period set for the Auto scaling Group
- ☐ Adjust the Memory threshold set for the Auto scaling scale-in and scale-out process.
- ☐ Adjust the CPU threshold set for the Auto scaling scale-in and scale-out process.

Q6)

Your security officer has told you that you need to tighten up the logging of all events that occur on your AWS account. He wants to be able to access all events that occur on the account across all regions quickly and in the simplest way possible. He also wants to make sure he is the only person that has access to these events in the most secure way possible.

Which of the following would be the best solution to assure his requirements are met? Choose the correct answer from the options below ?

- ☐ Use to log all events to a separate S3 bucket in each region as Cloud Trail cannot write to a bucket in a different region. Use MFA and bucket

policies on all the different buckets.

☐ Use Cloud Trail to log all events to an Amazon Glacier Vault. Make sure the vault access policy only grants access to the security officer's IP address.

☒ Use Cloud Trail to log all events to one S3 bucket. Make this S3 bucket only accessible by your security office with a bucket policy that restricts access to his user only and also add MFA to the policy for a further level of security.

☐ Use Cloud Trail to send all API calls to Cloud Watch and send an email to the security officer every time an API call is made. Make sure the emails are encrypted.

Q7)

Your company has an on-premise Active Directory setup in place. The company has extended their footprint on AWS, but still want to have the ability to use their on-premise Active Directory for authentication.

Which of the following AWS services can be used to ensure that AWS resources such as AWS Workspaces can continue to use the existing credentials stored in the on-premise Active Directory.

☐ Use the AWS Simple AD service

☐ Use the Classic Link feature on AWS

☒ Use the Active Directory connector service on AWS

☐ Use the Active Directory service on AWS

Q8)

Your application is having a very high traffic, so you have enabled auto scaling in multi availability zone to suffice the needs of your application but you observe that one of the availability zone is not receiving any traffic.

What can be wrong here? Please select:

☐ instances need to manually added to availability zone

☐ Auto scaling can be enabled for multi AZ only in north Virginia region

☒ Availability zone is not added to Elastic load balancer

☐ Auto scaling only works for single availability zone

Q9)

Your development team is using access keys to develop an application that has access to S3 and Dynamo DB.

A new security policy has outlined that the credentials should not be older than 2 months and should be rotated.

How can you achieve this

☐ Delete the IAM Role associated with the keys after every 2 months. Then recreate the IAM Role again.

☐ Use the application to rotate the keys in every 2 months via the SDK

☒ Use a script which will query the date the keys are created. If older than 2 months, delete them and recreate new keys.

☐ Delete the user associated with the keys after every 2 months. Then recreate the user again.

Q10)

You are in charge of designing CloudFormation templates for your company. One of the key requirements is to ensure that if a CloudFormation stack is deleted, a snapshot of the relational database which is part of the stack, is created.

How can you achieve this in the best possible way.

☐ Create a snapshot of the relational database beforehand so that when the CloudFormation stack is deleted, the snapshot of the database will be present.

☐ Create a new CloudFormation template to create a snapshot of the relational database.

☐ Use the Update policy of the CloudFormation template to ensure a snapshot is created of the relational database.

☒ Use the Deletion policy of the CloudFormation template to ensure a snapshot is created of the relational database.

Q11)

You're building a mobile application game. The application needs permissions for each user to communicate and store data in DynamoDB tables.

What is the best method for granting each mobile device that installs your application to access DynamoDB tables for storage when required? Choose the correct answer from the options below

☐ Create an IAM group that only gives access to your application and to the DynamoDB tables. Then, when writing to DynamoDB, simply include the unique device ID to associate the data with that specific user.

☐ During the install and game configuration process, have each user create an IAM credential and assign the IAM user to a group with proper permissions to communicate with DynamoDB.

☐ Create an Active Directory server and an AD user for each mobile application user. When the user signs in to the AD sign-on, allow the AD server to federate using SAML 2.0 to IAM and assign a role to the AD user which is the assumed role with SAML.

☒ Create an IAM role with the proper permission policy to communicate with the DynamoDB table. Use web identity federation, which assumes the IAM role using AssumeRoleWithWebIdentity, when the user signs in, granting temporary security credentials using STS.

Q12)

You need to deploy a multi-container Docker environment on to Elastic Beanstalk.

Which of the following files can be used to deploy a set of Docker containers to Elastic Beanstalk? Please select:

- ☒ Dockerrun.aws.json
 - ☐ Docker file
 - ☐ Docker run
 - ☐ Docker Multi file
-

Q13)

You currently have an application deployed via Elastic Beanstalk. You are now deploying a new application and have ensured that Elastic beanstalk has detached the current instances and deployed and reattached new instances. But the new Instances are still not receiving any sort of traffic.

Why is this the case.

- ☒ It takes time for the ELB to register the instances, hence there is a small time frame before your instances can start receiving traffic
 - ☐ The instances are of the wrong AMI. hence they are not being detected by the ELB.
 - ☐ You need to create a new Elastic Beanstalk application, because you cannot detach and then reattach instances to an ELB within an Elastic Beanstalk application
 - ☐ The Instances needed to be reattached before the new application version was deployed
-

Q14)

As part of your continuous deployment process, your application undergoes an I/O load performance test before it is deployed to production using new AM Is. The application uses one Amazon EBS PIOPS volume per instance and requires consistent I/O performance.

Which of the following must be carried out to ensure that I/O load performance tests yield the correct results in a repeatable manner?

- ☐ Ensure that snapshots of the Amazon EBS volumes are created as a backup.
 - ☒ Ensure that the Amazon EBS volumes have been pre-warmed by reading all the blocks before the test.
 - ☐ Ensure that the I/O block sizes for the test are randomly selected.
 - ☐ Ensure that the Amazon EBS volume is encrypted.
-

Q15)

You are creating a cloud formation templates which takes in a database password as a parameter.

How can you ensure that the password is not visible when anybody tries to describes the stack Please select:

- ☐ Set the hidden attribute for the Cloud formation resource
 - ☐ Use the hidden property for the parameter value
 - ☒ Use the No Echo property for the parameter value .
 - ☐ Use the password attribute for the resource
-

Q16)

You run a multi-tier architecture on AWS with web server instances running Nginx. Your users are getting errors when they use the web application.

How can diagnose the errors quickly and efficiently Please select:

- ☐ Install the Cloud Watch Logs agent and send Nginx access log data to Cloud Watch. From there, pipe the log data through to a third party logging and graphing tool.
 - ☒ Install the Cloud Watch Logs agent and send Nginx access log data to Cloud Watch. Then, filter the log streams for searching the relevant errors.
 - ☐ Send all the errors to AWS Config for processing.
 - ☐ Send all the errors to AWS Lambda for processing.
-

Q17) Which of the following services along with Cloud formation helps in building a Continuous Delivery release practice

- ☐ AWS Cloud trail
 - ☐ AWS Config
 - ☐ AWS Lambda
 - ☒ AWS Code Pipeline
-

Q18)

You have a large multi-tier architecture that serves public facing web traffic through a load balancer and is backed by a web tier that is contained within an Amazon EC2 Auto Scaling group. During a peak in traffic, you discover that your web tier is adding new Instances disproportionately compared to the amount of Incoming traffic and the Auto Scaling policy that was set up.

What should you do in order to stop the Auto Scaling group from scaling incorrectly in response to incoming traffic?

- ☐ Using a third-party configuration management tool and the AWS SDK suspend all Scheduled Actions of the Auto Scaling group until after the traffic peak and then resume all scheduled Actions.
- ☒ Using a custom Cloud Watch metric insert the elapsed time since the instance launch to the time the instance responds to an Elastic Load Balancing health check, and periodically adjust the Pause Time and the Cool Down property on the Auto Scaling group to be over the value of the metric.
- ☐ Using Cloud Watch and the instance Boot Time metric, increase the Pause Time and Cool Down property on the Auto Scaling group to be over

the value of the metric.

- Using a custom Cloud Watch metric insert the elapsed time since the instance launch to the time the instance responds to an Elastic Load Balancing health check, and periodically adjust the Pause Time of the Update Policy and reduce the Scaling Adjustment property by 50%

Q19) Which of the following is not a supported platform for the Elastic beanstalk service Please select:

- ✔ Angular JS
- PHP
- .Net
- Java

Q20)

You have configured the following AWS services in your organization - Auto Scaling group, Elastic Load Balancer, and EC2 instances.

You have been requested to terminate an instance from the Auto scaling Group when the Cpu utilization is less than 30%.

How can you achieve this. Please select:

- ✔ Create a Cloud Watch alarm to send a notification to the Auto Scaling group when the aggregated CPU utilization is less than 30% and configure the Auto Scaling policy to remove one instance.
- Create a Cloud Watch alarm to send a notification to SQS. SQS can then remove one instance from the Auto scaling Group.
- Create a Cloud Watch alarm to send a notification to the ELB. The ELB can then remove one instance from the Auto scaling Group.
- Create a Cloud Watch alarm to send a notification to the admin team. The admin team can then manually terminate an instance from the Auto scaling Group.

Q21)

Your EMR cluster uses twelve m4.large instances and runs 24 hours per day, but it is only used for processing and reporting during business hours.

Which options can you use reduce the costs? Choose 2 answers from the options given below Please select:

- ✔ Use Spot instances for tasks nodes when needed
- Use Auto scaling to scale in and scale out the cluster whenever required.
- ✔ Migrate the data from HDFS to S3 using S3DisCp and turn off the cluster when not in use
- Use Reserved Instances for the task nodes

Q22) Which of the following tools for EC2 can be used to administer instances without the need to SSH or RDP into the instance?

- AWS Config
- ✔ Run Command
- EC2Config
- AWS Code Pipeline

Q23)

You have are managing an application that contains Go as the front end, Mongo DB for document management and is hosted on a relevant Web server. You prebake AM I?s with the latest version of the Web server, then user the User Data section to setup the application.

You now have a change to the underlying Operating system version and need to deploy that accordingly.

How can this be done in the easiest way possible.

- Create an Ops work stack with the new AMi and then deploy the application accordingly.
- Create a Cloud formation stack with the new AM1 and then deploy the application accordingly.
- ✔ Create a new pre-baked AM1 with the new OS and use the User Data section to deploy the application.
- Create a new EBS Volume with the relevant Os patches and attach it to the EC2 instance.

Q24)

You are working with a customer who is using Chef Configuration management in their data center.

Which service is designed to let the customer leverage existing Chef recipes in AWS?

- AWS Elastic Beanstalk
- AWS Cloud Formation
- ✔ AWS Ops Works
- Amazon Simple Workflow Service

Q25)

You have defined a Linux based instance stack in Ops work. You now want to attach a database to the Ops work stack.

Which of the below is an important step to ensure that the application on the Linux instances can communicate with the database?

- ☒ Add the appropriate driver packages to ensure the application can work with the database
 - ☐ Add another stack with the database layer and attach it to the application stack.
 - ☐ Configure database tags for the Ops work application layer
 - ☐ Configure SSL so that the instance can communicate with the database
-

Q26)

You have a web application composed of an Auto Scaling group of web servers behind a load balancer, and create a new AMI for each application version for deployment. You have a new version to release, and you want to use the A/B deployment technique to migrate users over in a controlled manner while the size of the fleet remains constant over a period of 12 hours, to ensure that the new version is performing well.

What option should you choose to enable this technique while being able to roll back easily?

- ☒ Create a load balancer. Create an Auto Scaling launch configuration with the new AMI to use the new launch configuration and to register instances with the new load balancer. Use Amazon Route53 weighted Round Robin to vary the proportion of requests sent to the load balancers.
 - ☐ Create an Auto scaling launch configuration with the new AMI. Configure Auto Scaling to vary the proportion of instances launched from the two launch configurations.
 - ☐ Create an Auto scaling launch configuration with the new AMI. Configure the Auto Scaling group with the / new launch configuration. Use the Auto Scaling rolling updates feature to migrate to the new version.
 - ☐ Launch new Instances using the new AMI and attach them to the Auto Scaling group. Configure Elastic Load Balancing to vary the proportion of requests sent to instances running the two application versions.
-

Q27) Which of the following tools is available to send log data from EC2 Instances. Please select:

- ☒ Cloud Watch Logs Agent
 - ☐ Cloud Watch Agent
 - ☐ Logs console.
 - ☐ Logs Stream
-

Q28)

A company wants to create standard templates for deployment of their Infrastructure.

Which AWS service can be used in this regard? Please choose one option. Please select:

- ☐ Amazon Simple Workflow Service
 - ☒ AWS Cloud Formation
 - ☐ AWS Elastic Beanstalk
 - ☐ AWS Ops Works
-

Q29)

A company is building a two-tier web application to serve dynamic transaction-based content. The data tier is leveraging an Online Transactional Processing (OLTP) database.

What services should you leverage to enable an elastic and scalable web tier?

- ☒ Elastic Load Balancing, Amazon EC2, and Auto Scaling
 - ☐ Elastic Load Balancing, Amazon RDS with Multi-AZ, and Amazon S3
 - ☐ Amazon EC2, Amazon Dynamo DB, and Amazon S3
 - ☐ Amazon RDS with Multi-AZ and Auto Scaling
-

Q30) What are the benefits when you implement a Blue Green deployment for your infrastructure or application level changes. Choose 3 answers from the options given below Please select:

- ☐ Ability to deploy with higher risk
 - ☒ Good turnaround time for application deployments
 - ☒ Better rollback capabilities
 - ☒ Near zero-downtime release for new changes
-

Q31)

You have just developed a new mobile application that handles analytics workloads on large scale datasets that are stored on Amazon Red shift. Consequently, the application needs to access Amazon Red shift tables.

Which of the below methods would be the best, both practically and security-wise, to access the tables?

- ☐ Create a RedShift read-only access policy in IAM and embed those credentials in the application.
 - ☐ Create a I-SM client certificate in Red shift and authenticate using this certificate.
 - ☐ Create an AM user and generate encryption keys for that user. Create a policy for RedShift read-only access. Embed the keys in the application.
 - ☒ Use roles that allow a web identity federated user to assume a role that allows access to the RedShift table by providing temporary credentials.
-

Q32) Which of the following can be configured as targets for Cloud watch Events. Choose 3 answers from the options given below

- ☐ Amazon Code Commit
- ☒ Amazon EC2 Instances

- ✓ AWS Lambda Functions
- ✓ Amazon ECS Tasks

Q33)

You recently encountered a major bug in your web application during a deployment cycle. During this failed deployment, it took the team four hours to roll back to a previously working state, which left customers with a poor user experience. During the post-mortem, you team discussed the need to provide a quicker, more robust way to roll back failed deployments.

You currently run your web application on Amazon EC2 and use Elastic Load Balancing for your load balancing needs.

Which technique should you use to solve this problem?

- Create deployable versioned bundles of your application. Store the bundle on Amazon S3. Use an AWS Ops Works stack to redeploy your web application and use AWS Ops Works application versioning to initiate a rollback during failures.
- Create deployable versioned bundles of your application. Store the bundle on Amazon S3. Re-deploy your web application on Elastic Beanstalk and enable the Elastic Beanstalk auto-rollback feature tied to Cloud Watch metrics that define failure.
- ✓ Use an AWS Ops Works stack to re-deploy your web application and use AWS Ops Works Deployment Command to initiate a rollback during failures
- Using Elastic Bean Stalk redeploy your web application and use the Elastic Bean Stalk API to trigger a Failed Deployment API call to initiate a rollback to the previous version.

Q34) Robert is system administrator at TPT Limited and has enabled auto scaling in Multi-availability zone to address high traffic. Robert notes that one of the availability zone is not receiving any traffic. Which of the following options will resolve the issue?

- None of these
- ✓ Availability zone is not added to Elastic load balancer.
- Auto scaling only works for single availability zone.
- Instances need to manually added to availability zone.
- Auto scaling can be enabled for multi AZ only in north Virginia region.

Q35)

You have an I/O and network-intensive application running on multiple Amazon EC2 instances that cannot handle a large ongoing increase in traffic. The Amazon EC2 instances are using two Amazon EBS PIOPS volumes each, and each instance is identical.

Which of the following approaches should be taken in order to reduce load on the instances with the least disruption to the application?

- Create an AMI from an instance, and set up an Auto Scaling group with an instance type that has enhanced networking enabled and is Amazon EBS.
- ✓ Add an Amazon EBS volume for each running Amazon EC2 instance and implement RAID striping to improve I/O performance.
- Stop each instance and change each instance to a larger Amazon EC2 instance type that has enhanced networking enabled and is Amazon EBS-optimized. Ensure that RAID striping is also set up on each instance.
- Create an AMI from each instance, and set up Auto Scaling groups with a larger instance type that has enhanced networking enabled and is Amazon EBS.

Q36)

You are planning on configuring logs for your Elastic Load balancer.

At what intervals does the logs get produced by the Elastic Load balancer service. Choose 2 answers from the options given below ?

- 30 seconds
- ✓ 5 minutes
- ✓ 15 minutes
- 1 minute

Q37) Which of the following are the basic stages of a CI/CD Pipeline. Choose 3 answers from the options below?

- ✓ Build
- ✓ Production
- Run
- ✓ Source Control

Q38)

You work as a Dev ops Engineer for your company. There are currently a number of environments hosted via Elastic beanstalk. There is a requirement to ensure to use the fastest deployment method for changes to the Elastic Beanstalk environment.

Which deployment method is the fastest with Elastic Beanstalk?

- Rolling with additional batch
- Rolling
- ✓ All at Once
- Blue/Green

Q39)

The company you work for has a huge amount of infrastructure built on AWS. However there has been some concerns recently about the security of this infrastructure, and an external auditor has been given the task of running a thorough check of all of your company's AWS assets. The auditor will be in the USA while your company's infrastructure resides in the Asia Pacific (Sydney) region on AWS.

Initially, he needs to check all of your VPC assets, specifically, security groups and NACLs. You have been assigned the task of providing the auditor with a login to be able to do this.

Which of the following would be the best and most secure solution to provide the auditor with so he can begin his Initial Investigations? Choose the correct answer from the options below. Please select:

- ☐ Give him root access to your AWS Infrastructure, because he is an auditor he will need access to every service.
- ☐ Create an IAM user tied to an administrator role. Also provide an additional level of security with MFA.
- ☐ Create an IAM user with full VPC access but set a condition that will not allow him to modify anything if the request is from any IP other than his own.
- ☒ Create an IAM user who will have read-only access to your AWS VPC infrastructure and provide the auditor with those credentials.

Q40)

Your application requires long-term storage for backups and other data that you need to keep readily available but with lower cost.

Which S3 storage option should you use?

- ☐ S3 Standard
- ☐ Reduced Redundancy Storage
- ☐ Glacier
- ☒ Amazon S3 Standard - Infrequent Access

Q41)

You are trying to debug the creation of CloudFormation stack resources.

Which of the following can be used to help in the debugging process? Choose 2 answers from the options below

- ☐ Use AWS Config to debug all the API calls sent by the CloudFormation stack.
- ☐ Use CloudTrail to debug all the API calls sent by the CloudFormation stack.
- ☒ See the logs in the /var/log directory for Linux instances
- ☒ Use the AWS CloudFormation console to view the status of your stack.

Q42) Which of the following CLI commands can be used to describe the stack resources?

- ☒ aws cloudformation list-stack-resources
- ☐ aws cloudformation describe-stack
- ☐ aws cloudformation describe-stack-resources
- ☐ aws cloudformation list-stack

Q43)

You are a DevOps engineer for your company. There is a requirement to host a custom application which has custom dependencies for a development team.

This needs to be

- ☒ Package the application and dependencies with Docker, and deploy the Docker container with Elastic Beanstalk.
- ☐ Package the application and dependencies in an S3 file, and deploy the Docker container with Elastic Beanstalk.
- ☐ Package the application and dependencies with Docker in Elastic Beanstalk, and deploy with Elastic Beanstalk.
- ☐ Package the application and dependencies with Docker, and deploy the Docker container with CloudFormation.

Q44)

Your application has an Auto Scaling group of three EC2 instances behind an Elastic Load Balancer. Your Auto Scaling group was updated with a new launch configuration that refers to an updated AMI. During the deployment, customers complained that they were receiving several errors even though all instances passed the ELB health checks.

How can you prevent this from happening again?

- ☐ Manually terminate the instances with the older launch configuration.
- ☐ Update the launch configuration instead of updating the Auto Scaling Group
- ☒ Create a new launch configuration with the updated AMI and associate it with the Auto Scaling group. Increase the size of the group to six and when instances become healthy revert to three.
- ☐ Create a new ELB and attach the Auto Scaling Group to the ELB

Q45)

Your application is having a very high traffic, so you have enabled auto scaling in multiple availability zones to suffice the needs of your application but you observe that one of the availability zones is not receiving any traffic.

What can be wrong here?

- ☐ Auto scaling can be enabled for multi AZ only in north Virginia region
- ☐ Auto scaling only works for single availability zone
- ☐ Instances need to manually added to availability zone
- ☒ Availability zone is not added to Elastic load balancer

Q46)

You are responsible for an application that leverages the Amazon SDK and Amazon EC2 roles for storing and retrieving data from Amazon S3, accessing multiple Dynamo DB tables, and exchanging message with Amazon SQS queues. Your VP of Compliance is concerned that you are not following security best practices for securing all of this access.

He has asked you to verify that the application's AWS access keys are not older than six months and to provide control evidence that these keys will be rotated a minimum of once every six months.

Which option will provide your VP with the requested information?

- ☐ Create a script to query the IAM list-access keys API to get your application access key creation date and create a batch process to periodically create a compliance report for your VP.
- ☐ Update your application to log changes to its AWS access key credential file and use a periodic Amazon EMR job to create a compliance report for your VP
- ☒ Provide your VP with a link to AM AWS documentation to address the VPCs key rotation concerns.
- ☐ Create a new set of instructions for your configuration management tool that will periodically create and rotate the application's existing access keys and provide a compliance report to your VP.

Q47)

You are using Elastic Beanstalk for your development team. You are responsible for deploying multiple versions of your application.

How can you ensure, in an ideal way, that you don't cross the application version in Elastic beanstalk?

- ☐ Create a lambda function to delete the older versions.
- ☐ Create a script to delete the older versions.
- ☐ Use AWS Config to delete the older versions
- ☒ Use policies 9 Elastic beanstalk

Q48)

You have an Auto scaling Group which is launching a set of t2.small instances. You now need to replace those instances with a larger instance type.

How would you go about making this change in an ideal manner?

- ☐ Change the Instance type of the Underlying EC2 instance directly.
- ☒ Create a new launch configuration with the new instance type and update your Auto scaling Group.
- ☐ Create another Auto scaling Group and attach the new instance type.
- ☐ Change the Instance type in the current launch configuration to the new instance type.

Q49)

You work for an accounting firm and need to store important financial data for clients. Initial frequent access to data is required. but after a period of 2 months, the data can be archived and brought back only in the case of an audit.

What is the most cost-effective way to do this?

- ☐ Store all data in a Glacier
- ☒ Use lifecycle management to move data from S3 to Glacier
- ☐ Store all data in a private S3 bucket
- ☐ Use lifecycle management to store all data in Glacier