Answer Sheet

Q1)

You have an ELB on AWS which has a set of web servers behind them. There is a requirement that the SSL key used to encrypt data is always kept secure. Secondly the logs of ELB should only be decrypted by a subset of users.

Which of these architectures meets all of the requirements?

- Use Elastic Load Balancing to distribute traffic to a set of web servers. To protect the SSL private key. upload the key to the load balancer and configure the load balancer to offload the SSL traffic. Write your web server logs to an ephemeral volume that has been encrypted using a randomly generated AES key.
- Use Elastic Load Balancing to distribute traffic to a set of web servers. Use TCP load balancing on the load balancer and configure your web servers to retrieve the private key from a private Amazon 53 bucket on boot. Write your web server logs to a private Amazon 53 bucket using Amazon 53 server-side encryption.
- Use Elastic Load Balancing to distribute traffic to a set of web servers. Configure the load balancer to perform TCP load balancing, use an AWS Cloud HSM to perform the SSL transactions, and write your web server logs to an ephemeral volume that has been encrypted using a randomly generated AES key.
- ✅ Use Elastic Load Balancing to distribute traffic to a set of web servers, configure the load balancer to perform TCP load balancing, use an AWS Cloud HSM to perform the SSL transactions, and write your web server logs to a private Amazon S3 bucket using Amazon S3 server-side encryption.

Q2) When you implement a lifecycle hook in Auto scaling, by default what is the time limit in which the instance will be a pending state?

- ✅ 60 minutes
- 60 seconds
- 2ominutes
- 5 minutes

Q3) Which of the below is not a lifecycle event in Ops work?

- Shut down
- Setup
- Configure
- ✅ Uninstall

Q4) When deploying applications to Elastic Beanstalk, which of the following statements is false with regards to application deployment ?

- Should not exceed 512 MB in size
- which can be deployed to the application server
- ✅ Can include parent directories
- The application can be bundled in a zip file

Q5)

You have implemented a system to automate deployments of your configuration and application dynamically after an Amazon EC2 instance in an Auto Scaling group is launched.

Your system uses a configuration management tool that works in a standalone configuration, where there is no master node. Due to the volatility of application load, new instances must be brought into service within three minutes of the launch of the instance operating system. The deployment stages take the following times to complete:

1) Installing configuration management agent: 2mins

2) Configuring instance using artifacts: 4mins

3) Installing application framework: 1 5mins

4) Deploying application code: 1 mm

What process should you use to automate the deployment using this type of standalone agent configuration?

- Configure your Auto Scaling launch configuration with an Amazon EC2 User Data script to install the agent.. pull configuration artifacts and application code from an Amazon S3 bucket, and then execute the agent to configure the infrastructure and application.
- ✅ Build a custom Amazon Machine Image that includes all components pre-installed, including an agent. configuration artifacts, application frameworks, and code. Create a startup script that executes the agent to configure the system on startup.
- Create a web service that polls the Amazon EC2 API to check for new instances that are launched in an Auto Scaling group. When it recognizes a new Instance, execute a remote script via SSH to install the agent. SCP the configuration artifacts and application code, and finally execute the agent to configure the system
- Build a custom Amazon Machine Image that includes the configuration management agent and appli framework pre-Installed. Configure your Auto Scaling launch configuration with an Amazon EC2 UserData script tc pull configuration artifacts and application code from an Amazon S3 bucket, and then execute the agent to configure the system.

Q6) Which of the following is a container for metrics in Cloud watch?

- Packages
- Locale
- Metric Collection
- ✅ Namespaces

Q7)

One of your engineers has written a web application in the Go Programming language and has asked your Dev Ops team to deploy It to AWS.

The application code is hosted on a Git repository.

What are your options? (Select Two)

✅ Create a new AWS Elastic Beanstalk application and configure a Go environment to host your application. Using GIt check out the latest version of the code, once the local repository for Elastic Beanstalk Is configured use eb create command to create an environment and then use "eb deploy" command to deploy the application.
⚪ Write a Docker file that installs the Go base image and fetches your application using Git, Create an AWS Cloud Formation template that creates and associates an AWS::EC2::instance resource type with an AWS::EC2::Container resource type.
✅ Write a Docker file that installs the Go base image and fetches your application using Git. Create a new AWS Elastic Beanstalk application and use this Docker file to automate the deployment.
⚪ Write a Docker file that installs the Go base image and uses Git to fetch your application. Create a new AW Ops Works stack that contains a Docker layer that uses the Dockerrun .aws. json file to deploy your container and then use the Dockerfile to automate the deployment.

Q8)

You work for a startup that has developed a new photo-sharing application for mobile devices. Over recent months your application has increased in popularity; this has resulted in a decrease In the performance of the application due to the increased load. Your application has a two-tier architecture that is composed of an Auto Scaling PHP application tier and a My SQL RDS instance initially deployed with AWS Cloud Formation. Your Auto Scaling group has a mm value of 4 and a max value of 8. The desired capacity is now at 8 due to the high CPU utilization of the instances.

After some analysis, you are confident that the performance Issues stem from a constraint in CPU capacity, while memory utilization remains low.

You therefore decide to move from the general-purpose M3 instances to the compute-optimized C3 instances.

How would you deploy thi change while minimizing any interruption to your end users?

⚪ Update the launch configuration specified in the AWS Cloud Formation template with the new C3 instance type. Run a stack update with the new template. Auto Scaling will then update the instances with the new instance type
⚪ Sign into the AWS Management Console, copy the old launch configuration, and create a new launch configuration that specifies the C3 instances. Update the Auto Scaling group with the new launch configuration. Auto Scaling will then update the instance type of all running instances
⚪ Sign into the AWS Management Console and update the existing launch configuration with the new C3 instance type. Add an Update Policy attribute to your Auto Scaling group that specifies an AutoScaling Rolling update.
✅ Update the launch configuration specified in the AWS Cloud Formation template with the new C3 instance type. Also add an Update Policy attribute to your Auto Scaling group that specifies an Auto Scaling Rolling Update. P a stack update with the new template

---

Q9)

You work as a Dev ops Engineer for your company. There are currently a number of environments hosted via Elastic beanstalk. There is a requirement to ensure to ensure that the rollback time for a new version application deployment is kept to a minimal.

Which elastic beanstalk deployment method would fulfill this requirement?

⚪ All at Once
⚪ Rolling with additional batch
✅ Blue/Green
⚪ Rolling

Q10)

You have a web application that is currently running on a three M3 instances in three AZs. You have an Auto Scaling group configured to scale from three to thirty Instances. When reviewing your Cloud Watch metrics, you see that sometimes your Auto Scaling group is hosting fifteen instances. The web application is reading and writing to a Dynamo DBconfigured backend and configured with 800 Write Capacity Units and 800 Read Capacity Units.

Your Dynamo DB Primary Key is the Company ID. You are hosting 25 TB of data in your web application. You have a single customer that is complaining of long load times when their staff arrives at the office at 9:00 AM and loads the website, which consists of content that is pulled from Dynamo DB.

You have other customers who routinely use the web application. Choose the answer that will ensure high availability and reduce the customer?s access times?

⚪ Double the number of Read Capacity Units in your Dynamo DB instance because the instance is probably being throttled when the customer accesses the website and your web application.
⚪ Add a caching layer in front of your web application by choosing Elastic Cache Mem cached instances in one of the AZ5.
✅ Implement an Amazon SQS queue between your Dynamo DB database layer and the web application layer to minimize the large burst In traffic the customer generates when everyone arrives at the office at 9:00AM and begins accessing the website.
⚪ Change your Auto Scaling group configuration to use Amazon C3 instance types. because the web application layer is probably running out of compute capacity.

Q11) Which of the following Cloud formation helper scripts can help install packages on EC2 resources Please select:

⚪ cfn-get-metadata
⚪ cfn-hup

- cfn-signal
- ✅ cfn-init

Q12)

You are a Dev ops Engineer for your company. You are in charge of an application that uses EC2, ELB and Auto scaling. You have been requested to get the ELB access logs. When you try to access the logs .

You can see that nothing has been recorded in S3. Why Is this the case?

- The EC2 instances are not sending the required logs to ELB
- The Auto scaling service is not sending the required logs to ELB
- You do not have the necessary access to the logs generated by ELB.
- ✅ By default ELB access logs are disabled.

Q13)

You are in charge of designing a number of Cloud formation templates for your organization. You need to ensure that no one can update the stack production based resources.

How can this be achieved In the most efficient way?

- Use MFA to protect the resources
- ✅ Use a Stack based policy to protect the production based resources.
- Use 53 bucket policies to protect the resources.
- Create tags for the resources and then create IAM policies to protect the resources.

Q14)

A company has recently started using Docker cloud. This is a SaaS solution for managing docker containers on the AWS cloud and the solution provider is also on the same cloud platform. There is a requirement for the SaaS solution to access AWS resources.

Which of the following would meet the requirement for enabling the SaaS solution to work with AWS resources in the most secured manner? Please select:

- Create an IAM role for EC2 instances, assign it a policy mat allows only the actions required tor the Saas application to work. provide the role ARM to the SaaS provider to use when launching their application Instances. Many SaaS platforms can access aws resources via a Cross account access created in aws. If you go to Roles in your identity management, you will see the ability to add a cross account role.
- ✅ Create an IAM role for cross-account access allows the SaaS providers account to assume the role and assign it a policy that allows only the actions required by the SaaS application.
- Create an IAM user within the enterprise account assign a user policy to the IAM user that allows only the actions required by the SaaS application. Create a new access and secret key for the user and provide these credentials to the SaaS provider.
- From the AWS Management Console. navigate to the Security Credentials page and retrieve the access and secret key for your account.

Q15)

You are writing an AWS Cloud Formation template and you want to assign values to properties that will not be available until runtime. You know that you can use intrinsic functions to do this but are unsure as to which part of the template they can be used in.

Which of the following is correct in describing how you can currently use intrinsic functions in an AWS Cloud Formation template?

- ✅ You can only use intrinsic functions in specific parts of a template. You can use intrinsic functions in resource properties. metadata attributes, and update policy attributes.
- You can use intrinsic functions in any part of a template.
- You can use intrinsic functions only in the resource properties part of a template.
- You can use intrinsic functions in any part of a template. except AWS Template Format Version

---

Q16)

There is a requirement for an application hosted on a VPC to access the On-premise LDAP server. The VPC and the On-premise location are connected via an IPSec VPN.

Which of the below are the right options for the application to authenticate each user, Choose 2 answers from the options below

- ✅ Develop an identity broker that authenticates against LDAP and then calls IAM Security Token Service to get IAM federated user credentials. The application calls the identity broker to get IAM federated user credentials with access to the appropriate AWS service.
- ✅ The application authenticates against LDAP and retrieves the name of an IAM role associated with the use The application then calls the IAM Security Token Service to assume that IAM role. The application can use the temporary credentials to access any AWS resources
- The application authenticates against LDAP the application then calls the AWS Identity and Access Management (IAM) Security service to log in to IAM using the LDAP credentials the application can use the IAM temporary credentials to access the appropriate AWS service. x aniwer Is Ir:tzi
- Develop an identity broker that authenticates against IAM security Token service to assume a IAM role in order to get temporary AWS security credentials The application calls the Identity broker to get AWS temporary security credentials.

Q17)

A user is accessing RDS from an application. The user has enabled the Multi AZ feature with the MS SQL RDS DB.

During a planned outage how will AWS ensure that a switch from DB to a standby replica will not affect access to the

application?

- RDS will have both the DBs running independently and the user has to manually switch over
- The switch over changes Hardware so RDS does not need to worry about access
- ✅ RDS uses DNS to switch over to stand by replica for seamless transition
- RDS will have an internal IP which will redirect all requests to the new DB

Q18)

A custom script needs to be passed to a new Amazon Linux instances created in your Auto Scaling group.

Which feature allows you to accomplish this?

- EC2Config service
- AWS Contig.
- ✅ User data
- IAM roles

Q19)

You are currently planning on using Auto scaling to launch instances which have an application installed.

Which of the following methods will help ensure the instances are up and running in the shortest span of time to take In traffic from the users?

- ✅ Use AMrs which already have the software installed.
- Log into each Instance and Install the software.
- Use User Data to launch scripts to install the software.
- Use Docker containers to launch the software.

Q20)

Your company is planning to develop an application in which the front end is in .Net and the backend is in Dynamo DB. There Is an expectant of a high load on the application.

How could you ensure the scalability of the application to reduce the load on the Dynamo DB database? Choose an answer from the options below.

- Add more Dynamo DB databases to handle the load.
- Increase write capacity of Dynamo DB to meet the peak loads
- Launch Dynamo DB in Multi-AZ configuration with a global index to balance writes
- ✅ Use SQS to assist and let the application pull messages and then perform the relevant operation in Dynamo DB.

Q21)

You work for a company that automatically tags photographs using artificial neural networks (ANNs), which run on GPUs using C++. You receive millions of images at a time, but only 3 times per day on average. These images are loaded into an AWS S3 bucket you control for you in a batch, and then the customer publishes a JSON-formatted manifest into another S3 bucket you control as well.

Each image takes 10 milliseconds to process using a full GPU.

Your neural network software requires 5 minutes to bootstrap. Image tags are JSON objects, and you must publish them to an 53 bucket.

Which of these is the best system architectures for this system?

- Deploy your artificial neural network code to AWS Lambda as a bundled binary for the C++ extension. Make an S3 notification configuration on the manifest which publishes to another AWS Lambda running controller code. This controller code publishes all the images In the manifest to AWS Kinesis. Your ANN code Lambda Function uses the Kinesis as an Event Source. The system automatically scales when the stream contains image events
- ✅ Make an S3 notification configuration which publishes to AWS Lambda on the manifest bucket. Make the Lambda create a Cloud Formation Stack which contains the logic to construct an autoscaling worker tier of EC2 G2 instances with the artificial neural network code on each instance. Create an SQS queue of the images in the manifest. Tear the stack down when the queue is empty.
- Create an Ops Works Stack with two Layers. The first contains lifecycle scripts for launching and bootstrapping an HTTP API on G2 instances for image processing. and the second has an always-on instance which monitors the S3 manifest bucket for new files. When a new file Is detected. request instances to boot on the artificial neural network layer. When the instances are booted and the HTTP APIS are up. submit processing requests to individual instances.
- Create an Auto Scaling, Load Balanced Elastic Beanstalk worker tier Application and Environment. Deploy the artificial neural network code to G2 instances in this tier. Set the desired capacity to 1. Make the code periodically check 53 for new manifests. When a new manifest is detected, push all of the Images in the manifest into the SQS queue associated with the Elastic Beanstalk worker tier.

Q22)

Your firm has uploaded a large amount of aerial image data to 53. In the past, in your onpremises environment, you used a dedicated group of servers to process this data and used Rabbit MQ - An open source messaging system to get job Information to the servers. Once processed the data would go to tape and be shipped offsite.

Your manager told you to stay with the current design, and leverage AWS archival storage and messaging services to minimize cost.

Which is correct?

✅ Setup Auto-Scaled workers triggered by queue depth that use spot instances to process messages in SQS. Once data is processed, change the storage class of the 53 objects to Glacier

⚪ Use SQS for passing job messages. Use Cloud Watch alarms to terminate EC2 worker instances when they become idle. Once data is processed. change the storage class of the 53 objects to Reduced Redundancy Storage.

⚪ Use SNS to pass job messages use Cloud Watch alarms to terminate spot worker Instances when they become idle. Once data is processed. change the storage class of the S3 object to Glacier.

⚪ Change the storage class of the 53 objects to Reduced Redundancy Storage. Setup Auto-Scaled workers triggered by queue depth that use spot Instances to process messages In SQS. Once data Is processed. change the storage class of the S3 objects to Glacier.

Q23)

A user is using Cloud formation to launch an EC2 instance and then configure an application after the instance is launched.

The user wants the stack creation of ELB and Auto Scaling to wait until the EC2 instance is launched and configured properly.

How can the user configure this?

⚪ The user can use the Dependent Condition resource to hold the creation of the other dependent resources

⚪ It is not possible that the stack creation will wait until one service Is created and launched

✅ The user can use the Wait Condition resource to hold the creation of the other dependent resources

⚪ The user can use the Hold Condition resource to wait for the creation of the other dependent resources

---

Q24)

You are in charge of designing a number of Cloud formation templates for your organization. You are required to make changes to the stack resources every now and then based on the requirement.

How can you check the impact of the change to resources In a cloud formation stack before deploying changes to the stack?

⚪ Use Cloud formation Rolling Updates to check for the impact to the changes.

⚪ Use Cloud formation Stack Policies to check for the impact to the changes.

⚪ There is no way to control this. You need to check for the impact beforehand.

✅ Use Cloud formation change sets to check for the impact to the changes.

Q25)

You are working for a company has an on-premise infrastructure. There is now a decision to move to AWS. The plan is to move the development environment first.

There are a lot of custom based applications that need to be deployed for the development community.

Which of the following can help to implement the application for the development team?

⚪ Use Ops Works to deploy the docker containers.

✅ Use Elastic beanstalk to deploy the docker containers.

⚪ Use Cloud formation to deploy the docker containers.

✅ Create docker containers for the custom application components.

Q26) Which of the following run command types are available for ops work stacks? Choose 3 answers from the options given below.

✅ Configure

✅ Update Custom Cookbooks

✅ Execute Recipes

⚪ Un Deploy

Q27)

Of the 6 available sections on a Cloud Formation template (Template Description Declaration, Template Format Version Declaration, Parameters, Resources, Mappings, Outputs).

Which is the only one required for a Cloud Formation template to be accepted? Choose an answer from the options below ?

⚪ Template Declaration

⚪ Mappings

⚪ Parameters

✅ Resources

Q28) Which of the following is not a component of Elastic Beanstalk?

⚪ Environment

⚪ Application Version

⚪ Application

✅ Docker

Q29)

You have created a Dynamo DB table for an application that needs to support thousands of users. You need to ensure that each user can only access their own data in a particular table. Many users already have accounts with a third-party identity provider,

such as Face book, Google, or Login with Amazon.

How would you implement this requirement? Choose 2 answers from the options given below?

- [x] Use Web identity federation and register your application with a third-party identity provider such as Google, Amazon. or Face book.
- [ ] Use a third-party identity provider such as Google, Face book or Amazon so users can become an AWS IAM User with access to the application.
- [x] Create an IAM role which has specific access to the Dynamo DS table.
- [ ] Create an IAM User for all users so that they can access the application.

Q30)

A user is trying to save some cost on the AWS services.

Which of the below mentioned options will not help him save cost? Please select:

- [ ] Release the elastic IP if not required once the instance is terminated
- [ ] Delete the unutilized EBS volumes once the instance is terminated
- [x] Delete the Auto Scaling launch configuration after the instances are terminated
- [ ] Delete the AWS ELB after the instances are terminated

Q31)

An organization is planning to use AWS for their Production Roll Out. The organizations wants to implement automation for deployment, such that it will automatically create a LAMP stack, deploy an RDS My SQL DB instance, download the latest PHP installable from 53 and set up the ELB.

Which of the below mentioned AWS services meets the requirement for making an orderly deployment of the software?

- [ ] AWS Dev Ops
- [ ] AWS Elastic Beanstalk
- [ ] AWS Cloud front
- [x] AWS Cloud formation

Q32) Which of the following are components of the AWS Data Pipeline service. Choose 2 answers from the options given below

- [x] Pipeline definition
- [ ] Workflow Runner
- [ ] Task History
- [x] Task Runner

Q33)

Your company is using an Auto scaling Group to scale out and scale in instances. There is an expectation of a peak in traffic every Monday at 8am, The traffic is then expected to come down before the weekend on Friday 5pm.

How should you configure Auto scaling in this?

- [ ] Manually add Instances to the Auto scaling Group on Monday and remove them on Fr
- [x] Create a scheduled policy to scale up on Monday and scale down on Friday
- [ ] Create dynamic scaling policies to scale up on Monday and scale down on Friday
- [ ] Create a scheduled policy to scale up on Friday and scale down on Monday

---

Q34)

You are using for managing the instances in your AWS environment. You need to deploy a new version of your application. You'd prefer to use all new instances if possible, but you cannot have any downtime, You also don't want to swap any environment urls.

Which of the following deployment methods would you implement?

- [x] Using Rolling Updates' deployment method.
- [ ] Using "All at once" deployment method.
- [ ] Using "Blue Green" with "All at once" deployment method.
- [ ] Using "Blue Green" deployment method.

Q35) Which of the below things can you achieve with the Cloud watch logs service?

- [x] Stream the log data into Amazon Elastic search in near real-time with Cloud Watch Logs subscriptions.
- [ ] Stream the log data to Amazon Kinesis
- [ ] Record API calls for your AWS account and delivers log files containing API calls to your Amazon S3 bucket
- [x] Send the log data to AWS Lambda for custom processing or to load into other systems

---

Q36)

Your company is getting ready to do a major public announcement of a social media site on AWS. The website Is running on EC2 instances deployed across multiple Availability Zones with a Multi-AZ RDS My SQL Extra Large DB Instance.

The site performs a high number of small reads and writes per second and relies on an eventual consistency model. After comprehensive tests you discover that there is read contention on RDS My SQL.

Which are the best approaches to meet these requirements?

- Implement shading to distribute load to multiple RDS My SQL instances
- ✓ Add an RDS My SQL read replica in each availability zone
- Increase the RDS My SQL Instance size and Implement provisioned lops
- ✓ Deploy Elastic Cache in-memory cache running in each availability zone

---

Q37) Which of the following is false when it comes to using the Elastic Load balancer with Op sworks stacks?

- You can attach only one load balancer to a layer.
- Each load balancer can handle only one layer
- ✓ You can use either the Application or Classic Load Balancer with Ops Works stacks.
- You need to create the load balancer before hand and then attach it to the Ops work stack.

Q38)

A vendor needs access to your AWS account. They need to be able to read protected messages in a private S3 bucket. They have a separate AWS account.

Which of the solutions below is the best way to do this?

- Create an IAM User with API Access Keys. Give the vendor the AWS Access Key ID and AWS Secret Access Key for the user.
- Allow the vendor to SSH into your EC2 instance and grant them an IAM role with full access to the bucket.
- ✓ Create a cross-account IAM role with permission to access the bucket, and grant permission to use the role to the vendor AWS account.
- Create an 53 bucket policy that allows the vendor to read from the bucket from their AWS account. .

Q39)

You have a number of Cloud formation stacks in your IT organization.

Which of the following commands will help see all the cloud formation stacks which have a completed status? Please select:

- stacks-complete
- list. templates
- ✓ list-stacks
- describe-stacks

Q40)

You currently have a set of instances running on your Ops work stacks. You need to install security updates on these servers.

What does AWS recommend in terms of how the security updates should be deployed? Choose 2 answers from the options given below

- ✓ Create and start new instances to replace your current online Instances. Then delete the current instances.
- Create a new Ops work stack with the new instances.
- ✓ On Linux-based instances in Chef 11.10 or older stacks, run the Update Dependencies stack command.
- Create a cloud formation template which can be used to replace the instances.

Q41)

Your IT company is currently hosting a production environment in Elastic beanstalk. You understand that the Elastic beanstalk service provides a facility known as Managed updates which are minor and patch version updates which are periodically required for your system.

Your IT supervisor is worried about the impact that these updates would have on the system.

What can you tell about the Elastic beanstalk service with regards to managed updates

- Elastic Beanstalk applies managed updates with no downtime
- ✓ All of the above
- Elastic Beanstalk applies managed updates with no reduction In capacity
- Package updates can be configurable weekly maintenance window

---

Q42) Robert is system administrator at TPT Limited and is planning to configure an AWS Elastic Beanstalk worker tier for easy debugging. But in case there is an issue in completion of queue jobs, what should Robert configure?

- None of these
- Configure Rolling Deployments.
- Configure Blue-Green Deploy me
- ✓ Configure a Dead
- Configure Enhanced Health Reporting.

Q43) Which of the following are ways to secure data at rest and in transit in AWS. Choose 3 answers from the options given below ?

- ✓ Encrypt all EBS volumes attached to EC2 Instances
- Use OPS volumes when working with EBS volumes on EC2 Instances

✅ Use server side encryption for 53
✅ Use SSLIHTTPS when using the Elastic Load Balancer

Q44)

You have a web application hosted on EC2 instances. There are application changes which happen to the web application on a quarterly basis.

Which of the following are example of Blue Green deployments which can be applied to the application? Choose 2 answers from the options given below ?

⚪ Place the EC2 instances behind an ELB. Have a secondary environment with EC2 Instances and ELB in another region. Use Route53 with geo-location to route requests and switch over to the secondary environment.
⚪ Deploy the application to an elastic beanstalk environment. Use the Rolling updates feature to perform a Blue Green deployment.
✅ Deploy the application using Ops work stacks. Have a secondary stack for the new application deployment. Use Route53 to switch over to the new stack for the new application update.
✅ Deploy the application to an elastic beanstalk environment. Have a secondary elastic beanstalk environment in place with the updated application code. Use the swap URL'S feature to switch onto the new environment.

Q45)

Your company is supporting a number of applications that need to be moved to AWS. Initially the thought it moving these applications to the Elastic beanstalk service.

When going to the Elastic beanstalk service . you can see that the underlying platform service Is not an option In the Elastic beanstalk service.

Which of the following options can be used to port your application onto Elastic beanstalk

⚪ Use custom Cloud formation templates to deploy the application into Elastic beanstalk
✅ Create a docker container for the custom application and then deploy It to Elastic beanstalk
⚪ Use the Ops work service to create a stack. In the stack . create a separate custom layer. Deploy the application to this layer and then attach the layer to Elastic beanstalk
⚪ Use custom chef recipes to deploy your application in Elastic beanstalk.

Q46)

You are working for a startup company that is building an application that receives large amounts of data. Unfortunately, current funding has left the start-up short on cash, cannot afford to purchase thousands of dollars of storage hardware, and has opted to use AWS.

Which services would you Implement in order to store a virtually unlimited amount of data without any effort to scale when demand unexpectedly increases? Choose the correct answer from the options below Please select:

✅ Amazon S3, because it provides unlimited amounts of storage data, scales automatically, is highly available and durable
⚪ Amazon Glacier, to keep costs low for storage and scale Infinitely
⚪ Amazon EC2. because EBS volumes can scale to hold any amount of data and, when used with Auto Scaling. can be designed for fault tolerance and high availability
⚪ Amazon Import/Export because Amazon assists in migrating large amounts of data to Amazon S3

Q47) Which of the following commands for the elastic beanstalk CLI can be used to create the current application into the specified environment?

✅ eb create
⚪ en env
⚪ en app
⚪ eb start

Q48)

You are a Dev ops Engineer for your company. Your company is using Ops work stack to rollout a collection of web instances.

When the instances are launched , a configuration file need to be setup prior to the launching of the web application hosted on these Instances.

Which of the following steps would you carry out to ensure this requirement gets fulfilled. Choose 2 answers from the options given below Please select:

✅ Configure a recipe which sets the configuration file and add it to the Configure LifeCycle Event of the speci web layer.
⚪ Configure a recipe which sets the configuration file and add It to the Deploy LifeCycle Event of the specific web layer.
⚪ Ensure that the Ops work stack Is changed to use the AWS specific cookbooks
✅ Ensure that the Ops work stack Is changed to use custom cookbooks

---

Q49)

In reviewing the Auto-Scaling events for your application you notice that your application is scaling up and down multiple times in the same hour.

What design choice could you make to optimize for costs while preserving elasticity?

⚪ Modify the Auto Scaling policy to use scheduled scaling actions

- Modify the Auto Scaling group termination policy to terminate the newest instance first.
- Modify the Amazon Cloud watch alarm period that triggers your Auto Scaling scale down policy.
- Modify the Auto Scaling Group cool down timers

Q50)

Your company has an application hosted in AWS which makes use of Dynamo DB.

There is a requirement from the IT security department to ensure that all source IP addresses which make calls to the Dynamo DB tables are recorded.

Which of the following services can be used to ensure this requirement is fulfilled.

- AWS Cloud Trail
- AWS Code Commit
- AWS Cloud watch
- AWS Code Pipeline
