Answer Sheet

Q1)

You are setting up cloud formation templates for your organization. The cloud formation template consists of creating EC2 Instances for both your development and production environments in the same region.

Each of these instances will have an Elastic P and a security group attached to them which will be done via Cloud formation.

Your cloud formation stack for the development environment gets successfully created, but then the production cloud formation stack fails.

Which of the below could be a reason for this?

- You have chosen the wrong tags when creating the instances in both environments.
- You didn't choose the Production version of the AMI you are using when creating the production stack.
- ✓ You hit the soft limit of S EIP5 per region when creating the development environment. ..?
- You hit the soft limit for security groups when creating the development environment.

Q2)

You have a web application running on six Amazon EC2 instances, consuming about 45% of resources on each instance. You are using auto-scaling to make sure that six instances are running at all times.

The number of requests this application processes is consistent and does not experience spikes.

The application is critical to your business and you want high availability at all times.

You want the load to be distributed evenly between all instances.

You also want to use the same Amazon Machine Image (AMI) for all instances.

Which of the following architectural choices should you make?

- Deploy 3 EC2 Instances in one region and 3 In another region and use Amazon Elastic Load Balancer.
- ✓ Deploy 3 EC2 instances in one availability zone and 3 in another availability zone and use Amazon Elastic Load Balancer
- Deploy 6 EC2 instances in one availability zone and use Amazon Elastic Load Balancer.
- Deploy 2 EC2 Instances in three regions and use Amazon Elastic Load Balancer.

Q3)

You are a Dev ops Engineer for your company. You are responsible for creating Cloud formation templates for your company.

There Is a requirement to ensure that an S3 bucket is created for all resources In development for logging purposes.

How would you achieve this?

- Create separate Cloud formation templates for Development and production.
- Create an S3 bucket from before and then just provide access based on the tag value mentioned in the Cloud formation template
- Use the metadata section in the Cloud formation template to decide on whether to create the S3 bucket or not.
- ✓ Create a parameter in the Cloud formation template and then use the Condition clause in the template to create an S3 bucket if the parameter has a value of development .

Q4) Which of the following files needs to be included along with your source code binaries when deploying code using the AWS Code Deploy service

- app config .json
- appconfig.yml
- ✓ appspec.yml
- apps pec .j son

Q5)

You have the requirement to get a snapshot of the current configuration of the resources in your AWS Account.

Which of the following services can be used for this purpose Please select:

- ✓ AWS Conflg
- AWS Code Deploy
- AWS Trusted Advisor
- AWS lAM

Q6)

Your development team is using an Elastic beanstalk environment. After a week, the environment was torn down and a new one was created.

When the development team tried to access the data on the older environment, it was not available.

Why is this the case?

- This is because before the environment termination. Elastic beanstalk copies the data to Dynamo DB. and hence the data is not present in the EBS volumes

- ✅ This is because the underlying EC2 Instances are created with no persistent local storage
- This is because the underlying EC2 Instances are created with encrypted storage and cannot be accessed once the environment has been terminated.
- This is because the underlying EC2 Instances are created with IOPS volumes and cannot be accessed once1 the environment has been terminated.

Q7)

Your development team use .Net to code their web application. They want to deploy it to AWS for the purpose of continuous Integration and deployment. The application code is hosted in a Git repository.

Which of the following combination of steps can be used to fulfil this requirement. Choose 2 answers from the options given below ?

- Use a chef recipe to deploy the code and attach it to the Elastic beanstalk environment.
- ✅ Create a source bundle for the .Net code and upload It as an application revision.
- ✅ Use the Elastic beanstalk service to provision an IIS platform web environment to host the application.
- Use the Code Pipeline service to provision an uS environment to host the application

---

Q8)

You have a set of EC2 Instances hosting an nix server and a web application that is used by a set of users in your organization. After a recent application version upgrade, the instance runs into technical issues and needs an Immediate restart.

This does not give you enough time to inspect the cause of the issue on the server.

Which of the following options if implemented prior to the incident would have assisted in detecting the underlying cause of the issue?

- ✅ Install Cloud watch logs agent on the instance and send all the logs to Cloud watch logs.
- Create a snapshot of the EBS volume before restart attach it to another instance as a volume and then diagnose the issue.
- Enable detailed monitoring and check the Cloud watch metrics to see the cause of the issue.
- Stream all the data to Amazon Kinesis and then analyze the data in real time.

Q9)

Your company has a set of EC2 Instances that access data objects stored in an S3 bucket. Your IT Security department Is concerned about the security of this architecture and wants you to implement the following

1) Ensure that the EC2 Instance securely accesses the data objects stored in the S3 bucket

2) Ensure that the integrity of the objects stored in S3 is maintained.

Which of the following would help fulfill the requirements of the IT Security department. Choose 2 answers from the options given below Please select:

- ✅ Create an IAM Role and ensure the EC2 Instances uses the IAM Role to access the data in the bucket. ,,
- ✅ Use an 53 bucket policy that ensures that MFA Delete Is set on the objects In the bucket Your answer is partially correct.
- Create an IAM user and ensure the EC2 Instances uses the IAM user credentials to access the data in the bucket.
- Use 53 Cross Region replication to replicate the objects so that the integrity of data is maintained.

Q10)

You have been tasked with building out a duplicate environment in another region for disaster recovery purposes. Part of your environment relies on EC2 instances with preconfigured software.

What steps would you take to configure the Instances in another region? Choose the correct answer from the options below Please select:

- None of the above
- Make the EC2 instance shareable among other regions through IAM permissions
- Create an AMP of the EC2 instance
- ✅ Create an AMP of the EC2 instance and copy the AMI to the desired region

---

Q11)

You have a development team that is planning for continuous release cycles for their application. They want to use the AWS services available to be able to deploy a web application and also ensure they can rollback to previous versions fairly quickly.

Which of the following options can be used to achieve this requirement. Choose 2 answers from the options given below

- Use the Elastic beanstalk service. Create separate environments for each application revision. Revert back an environment incase the new environment does not work.
- Use the Cloud formation service. Create separate templates for each application revision and deploy them accordingly.
- ✅ Use the Ops work service to deploy the web instances. Deploy the app to the Ops work web layer. Rollback using the Deploy app In Ops work.
- ✅ Use the Elastic beanstalk service. Use Application versions and upload the revisions of your application. Deploy the revisions accordingly and rollback to prior versions accordingly.

Q12)

Your company currently has a set of EC2 Instances sitting behind an Elastic Load Balancer. There is a requirement to create an Ops work stack to host the newer version of this application. The idea Is to first get the stack in place, carry out a level of

testing and then deploy it at a later stage. The Ops work stack and layers have been setup.

To complete the testing process, the current ELB is being utilized. But you have now noticed that your current application has stopped responding to requests.

Why is this the case?

- This is because the Ops work web layer is utilizing the current instances after the ELB was attached as an additional layer
- ✓ The ELB would have deregistered the older Instances .-
- This is because the Ops work stack is utilizing the current instances after the ELB was attached as a layer.
- You have configured the Opswork stack to deploy new instances in the same domain the older instances.

Q13) You are having a web and worker role infrastructure defined in AWS using Amazon EC2 resources. You are using SQS to manage the jobs being send by the web role, Which of the following is the right way to ensure the worker processes are adequately setup to handle the number of jobs send by the web role Please select:

- Use ELB to ensure that the load Is evenly distributed to the set of web and worker instances
- ✓ Use Cloud watch monitoring to check the size of the queue and then scale out using Autoscaling to ensure that it can handle the right number of jobs.
- Use Cloud watch monitoring to check the size of the queue and then scale out SQS to ensure that it can handle the right number of jobs
- Use Route53 to ensure that the load is evenly distributed to the set of web and worker instances

---

Q14) Which of the following CLI commands is used to spin up new EC2 Instances? Please select:

- aws ec2 launch-instances
- ✓ aws ec2 run-instances
- aws ec2 new-Instances
- aws ec2 create-instances

Q15) When one creates an encrypted EBS volume and attach it to a supported instance type which of the following data types are encrypted?

- All snapshots created from the volume
- ✓ Data at rest inside the volume
- All data moving between the volume and the instance
- All data copied from the EBS volume to S3

Q16)

You are a Dev ops Engineer in your company. You have been instructed to ensure there is an automated backup solution in place for EBS Volumes.

These snapshots need to be retained only for a period of 20 days.

How can you achieve this requirement in an efficient manner?

- ✓ Use the aws ec2 create-snapshot API to create a snapshot of the EBS Volume. The use the describes nap shot to see those snapshots which are greater than 20 days and then delete them accordingly using the deletes nap shot API call.
- Use Lifecycle policies to push the EBS Volumes to Amazon Glacier. Then use further lifecycle policies to delete the snapshots after 20 days.
- Use Lifecycle policies to push the EBS Volumes to Amazon S3. Then use further lifecycle policies to delete the snapshots after 20 days.
- Use the aws ec2 create-volume API to create a snapshot of the EBS Volume. The use the describe-volume to see those snapshots which are greater than 20 days and then delete them accordingly using the delete-volume API call.

Q17)

You are designing a cloud formation stack which involves the creation of a web server and a database server. You need to ensure that the web server In the stack gets created after the database server is created.

How can you achieve this?

- Ensure that the database server Is defined first and before the web server In the cloud formation template. The stack creation normally goes in order to create the resources.
- Ensure that the web server Is defined as a child of the database server in the cloud formation template.
- Ensure that the database server is defined as a child of the web server in the cloud formation template.
- ✓ Use the Depends On attribute to ensure that the database server is created before the web server.

Q18)

You have an application which consists of EC2 instances in an Auto Scaling group. Between a particular time frame every day, there is an increase In traffic to your website. Hence users are complaining of a poor response time on the application. You have configured your Auto Scaling group to deploy one new EC2 instance when Cpu utilization is greater than 60% for 2 consecutive periods of 5 minutes.

What is the least cost-effective way to resolve this problem?

- Decrease the collection period to ten minutes
- Decrease the threshold cu utilization percentage at which to deploy a new Instance
- Decrease the consecutive number of collection periods
- ✓ Increase the minimum number of instances in the Auto Scaling group

---

Q19)

You have a current Cloud formation template defines in AWS. You need to change the current alarm threshold defined in the Cloud watch alarm.

How can you achieve this?

✅ Update the template and then update the stack with the new template. Only those resources that need to be changed will be changed. All other resources which do not need to be changed will remain as they are.
⚪ Update the template and then update the stack with the new template. Automatically all resources will be changed In the stack.
⚪ Currently there is no option to change what is already defined in Cloud formation templates
⚪ Delete the current cloud formation template. Create a new one which will update the current resources.

Q20)

You are responsible for your company's large multi-tiered Windows-based web application running on Amazon EC2 instances situated behind a load balancer. While reviewing metrics, you've started noticing an upwards trend for slow customer page load time. Your manager has asked you to come up with a solution to ensure that customer load time is not affected by too many requests per second.

Which technique would you use to solve this issue?

⚪ Re-deploy your Infrastructure using an AWS Cloud Formation template. Spin up a second AWS Cloud Formation stack. Configure Elastic Load Balancing Spillover functionality to spill over any slow connections to the second AWS Cloud Formation stack.
⚪ Re-deploy your infrastructure using an AWS Cloud Formation template. Configure Elastic Load Balancing health checks to initiate a new AWS Cloud Formation stack when health checks return failed.
⚪ Re-deploy your application using an Auto Scaling template. Configure the Auto Scaling template to spin up a new Elastic Beanstalk application when the customer load time surpasses your threshold.
✅ Re-deploy your infrastructure using AWS Cloud Formation, Elastic Beanstalk, and Auto Scaling. Set up your Auto Scaling group policies to scale based on the number of requests per second as well as the current customer load time.

Q21)

The project you are working on currently uses a single AWS Cloud Formation template to deploy its AWS infrastructure, which supports a multi-tier web application. You have been tasked with organizing the AWS Cloud Formation resources so that they can be maintained in the future, and so that different departments such as Networking and Security can review the architecture before it goes to Production.

How should you do this in a way that accommodates each department, using their existing workflows?

⚪ Use a custom application and the AWS SDK to replicate the resources defined in the current AWS Cloud Formation template. and use the existing code review system to allow other departments to approve changes before altering the application for future deployments.
⚪ Organize the AWS Cloud Formation template so that related resources are next to each other In the template for each departments use, leverage your existing continuous integration tool to constantly deploy changes from all parties to the Production environment, and then run tests for validation.
⚪ Organize the AWS Cloud Formation template so that related resources are next to each other in the template, such as VPC subnets and routing rules for Networking and security groups and IAM information for Security.
✅ Separate the AWS Cloud Formation template into a nested structure that has individual templates for the resources that are to be governed by different departments, and use the outputs from the networking and security stacks for the application template that you control.

Q22)

You work for a startup that has developed a new photo-sharing application for mobile devices. Over recent months your application has increased in popularity; this has resulted In a decrease in the performance of the application clue to the Increased load. Your application has a two-tier architecture that Is composed of Auto Scaling PHP application tier and a My SQL RDS instance initially deployed with AWS Cloud Formation. Your Auto Scaling group has a mm value of 4 and a max value of 8. The desired capacity is now at because of the high CPU utilization of the instances. After some analysis, you are confident that the performance Issues stem from a constraint In CPU capacity, although memory utilization remains low. You therefore decide to move from the general-purpose M3 instances to the compute-optimized C3 instances.

How would you deploy this change while minimizing any interruption to your end users?

⚪ Update the launch configuration specified in the AWS Cloud Formation template with the new C3 Instance / type. Run a stack update with the new template. Auto Scaling will then update the instances with the new instance type.
✅ Update the launch configuration specified In the AWS Cloud Formation template with the new C3 Instance type. Also add an Update Policy attribute to your Auto Scaling group that specifies Auto Scaling Rolling update.
⚪ Sign into the AWS Management Console, copy the old launch configuration. and create a new launch configuration that specifies the C3 instances. Update the Auto Scaling group with the new launch configuration. Auto Scaling will then update the instance type of all running instances.
⚪ Sign into the AWS Management Console, and update the existing launch configuration with the new C3 instance type. Add an Update Policy attribute to your Auto Scaling group that specifies Auto Scaling Rolling Update.

Q23)

During metric analysis, your team has determined that the company's website is experiencing response times during peak hours that are higher than anticipated. You currently rely on Auto Scaling to make sure that you are scaling your environment during peak windows.

How can you improve your Auto Scaling policy to reduce this high response time? Choose 2 answers.

✅ Increase your Auto Scaling group's number of max servers.
⚪ Push custom metrics to Cloud Watch to monitor your Cpu and network bandwidth from your servers, which will allow your Auto Scaling policy to

have better fine-grain insight.
- Create a script that wns and monitors your servers: when it detects an anomaly in load, it posts to an Amazon SNS topic that triggers Elastic Load Balancing to add more servers to the load balancer.
- ✓ Push custom metrics to Cloud Watch for your application that include more detailed information about your web application, such as how many requests it is handling and how many are waiting to be processed.

Q24)

You have enabled Elastic Load Balancing HTTP health checking. After looking at the AWS Management Console, you see that all Instances are passing health checks, but your customers are reporting that your site is not responding.

What is the cause?

- ✓ The HTTP health checking system is misreporting due to latency in Inter-instance metadata synchronization.
- The health check in place is not sufficiently evaluating the application function.
- The application is returning a positive health check too quickly for the AWS Management Console to respond.
- Latency in DNS resolution Is Interfering with Amazon EC2 metadata retrieval.

Q25)

You have decided that you need to change the instance type of your production instances which are running as part of an Auto Scaling group. The entire architecture is deployed using CloudFormation Template. You currently have 4 instances in Production.

You cannot have any interruption in service and need to ensure 2 instances are always running during the update?

Which of the options below listed can be used for this?

- Auto Scaling scheduled Action
- ✓ Auto Scaling Rolling Update
- Auto Scaling integration Update
- Auto Scaling Replacing update

Q26)

Your application uses Cloud Formation to orchestrate your application's resources. During your testing phase before the application went live, your Amazon RDS instance type was changed and caused the instance to be re-created, resulting In the loss of test data.

How should you prevent this from occurring in the future?

- ✓ In the AWS Cloud Formation template, set the Deletion Policy of the AWS::RDS::DB instances Deletion Policy property to Retain.
- Use an AWS Cloud Formation stack policy to deny updates to the instance. Only allow Update Stack permission to IAM principals that are denied Set Stack Policy.
- Subscribe to the AWS Cloud Formation notification BeforeResourceUpdate7 and call Cancel Stack Update if the resource identified is the Amazon RDS instance.
- Within the AWS Cloud Formation parameter with which users can select the Amazon RDS instance type. set Allowed Values to only contain the current Instance type.

Q27)

Your company has recently extended its datacenter into a VPC on AWS. There is a requirement for on- premise users manage AWS resources from the AWS console. You don't want to create IAM users for them again.

Which of the below options will fit your needs for authentication?

- Use your on-premises SAML2.0-compliat identity provider (IDP) to retrieve temporary security credentials to enable members to sign in to the AWS Management Console.
- Use OAuth 2.0 to retrieve temporary AWS security credentials to enable your members to sign in to the AWS Management Console.
- ✓ Use your on-premises SAML 2 0-compliant identity provider (IDP) to grant the members federated access to the AWS Management Console via the AWS single sign-on (550) endpoint.
- Use web Identity Federation to retrieve AWS temporary security credentials to enable your members to sl> in to the AWS Management Console,

---

Q28) When thinking of AWS Elastic Beanstalk, the ?Swap Environment URLs? feature most directly aids in what?

- Canary Deployments
- Immutable Rolling Deployments
- Mutable Rolling Deployments
- ✓ Blue-Green Deployments

Q29)

You want to use Code Deploy to deploy code that is hosted on your git hub repository.

Which of the following additional services can help fulfill this requirement?

- ✓ Use the Code Batch service
- Use the Code Commit service
- Use the Code Pipeline service
- Use the SQS service

Q30) Which of the following will you need to consider so you can set up a solution that incorporates single sign-on from your corporate AD or LDAP directory and restricts access for each user to a designated user folder in a bucket? Choose 3 Answers from the options below

- Tagging each folder in the bucket
- ✓ Configuring IAM role
- Setting up a matching IAM user for every user in your corporate directory that needs access to a folder in the bucket
- ✓ Using AWS Security Token Service to generate temporary tokens

---

Q31)

Your company has an e-commerce platform which is expanding all over the globe, you have EC2 instances deployed in multiple regions you want to monitor performance of all of these EC2 instances.

How will you setup Cloud Watch to monitor EC2 instances In multiple regions?

- Register instances running on different regions to Cloud Watch
- This Is not possible
- Create separate dash boards in every region.
- ✓ Have one single dashboard to report metrics to Cloud Watch from different region

---

Q32)

Your company owns multiple AWS accounts. There is currently one development and one production account. You need to grant access to the development team to an 53 bucket in the production account.

How can you achieve this?

- When creating the role, define the Development account as a trusted entity and specify a permissions poli that allows trusted users to update the 53 bucket.
- Use web identity federation with a third-party identity provider with AWS STS to grant temporary credentials and membership into the production IAM user.
- Create an IAM user in the Production account that allows users from the Development account (the trusted account) to access the 53 bucket in the Production account.
- ✓ Create an LAM cross account role in the Production account that allows users from the Development account to access the S3 bucket in the Production account.

Q33) Which of the following is not a supported platform on Elastic Beanstalk?

- Node.js
- Packer Builder
- Go
- ✓ Kurber netes

Q34)

You need to deploy a multi-container Ducker environment on to Elastic beanstalk.

Which of the following files can be used to deploy a set of Ducker containers to Elastic beanstalk ?

- ✓ Ducker run.aws.json
- Ducker Multi file
- Ducker run
- Ducker file

Q35)

You have an AWS Ops Works Stack running Chef Version 11.10. Your company hosts its own proprietary cookbook on Amazon 53, and this is specified as a custom cookbook in the stack.

You want to use an open- source cookbook located in an external Git repository.

What tasks should you perform to enable the use of both custom cookbooks?

- . In your cookbook create an S3 sym link object that points to the open source projects cookbook.
- n the Ops Works stack settings add the open source project's cookbook details in addition to your cookt
- Contact the open source project's maintainers and request that they pull your cookbook into theirs, Update the stack to use their cookbook.
- ✓ In the AWS ops works stack settings, enable Berk shelf. Create a new cookbook with a Berksfile that specifies the other two cookbooks. Configure the stack to use this new cookbook.

Q36)

You are designing an application that contains protected health information. Security and compliance requirements for your application mandate that all protected health Information in the application use encryption at rest and in transit.

The application uses a three-tier architecture where data flows through the load balancer and is stored on Amazon EBS volumes for processing and the results are stored in Amazon S3 using the AWS SDK.

Which of the following two options satisfy the security requirements? (Select two) Please select:

✅ Use TCP load balancing on the load balancer. SSL termination on the Amazon EC2 instances. OS-level disk encryption on the Amazon EBS volumes and Amazon 53 with server-side encryption
⚪ use SSL termination on the load balancer. Amazon EBS encryption on Amazon EC2 instances and Amazon S3 with server- side encryption.
✅ Use TCP load balancing on the load balancer. SSL termination on the Amazon EC2 instances and Amazon S3 with server-side encryption.
⚪ Use SSL termination on the load balancer an SSL listener on the Amazon EC2 Instances. Amazon EBS encryption on EBS volumes containing PHI and Amazon S3 with server-side encryption.

Q37)

You are managing the development of an application that uses Dynamo DB to store JSON data. You have already set the Read and Write capacity of the Dynamo DB table.

You are unsure of the amount of the traffic that will be received by the application during the deployment time.

How can you ensure that the Dynamo DB is not highly throttled and does not become a bottleneck for the application? Choose 2 answers from the options below.

✅ Create a Cloud watch alarm which would then send a trigger to AWS Lambda to increase the Read and Write capacity of the Dynamo DB table.
⚪ Create a Cloud watch alarm which would then send a trigger to AWS Lambda to create a new Dynamo DB table.
⚪ Monitor the System Errors metric using Cloud watch
✅ Monitor the Consumed Read Capacity Units and Consumed Write Capacity Jnits metric using Cloud watch.

Q38)

A group of developers in your organization want to migrate their existing application into Elastic Beanstalk and want to use Elastic load Balancing and Amazon SQS.

They are currently using a custom application server.

How would you deploy their system to Elastic Beanstalk?

⚪ Configure an Elastic Beanstalk platform using AWS Ops Works deploy It to Elastic Beanstalk and run a script that creates a load balancer and an Amazon SOS queue.
✅ Use a Ducker container that has the third party application server installed on it and that creates the load balancer and an Amazon SQS queue using the application source bundle feature.
⚪ Create a custom Elastic Beanstalk platform that contains the third party application server and runs a script that creates a load balancer and an Amazon SQS queue.
⚪ Configure an AWS Ops Works stack that installs the third party application server and creates a load balancer and an Amazon SQS queue and then deploys it to Elastic Beanstalk.

Q39)

You are a Dev ops Engineer for your company. You are planning on using Cloud watch for monitoring the resources hosted in AWS.

Which of the following can you do with Cloud watch logs ideally. Choose 3 answers from the options given below ?

⚪ Send the data to SQS for further processing.
✅ Send the log data to AWS Lambda for custom processing
✅ Stream the log data into Amazon Elastic search for any search analysis required.
✅ Stream the log data to Amazon Kinesis for further processing

---

Q40)

An application is currently writing a large number of records to a Dynamo DB table in one region. There is a requirement for a secondary application to just take in the changes to the Dynamo DB table every 2 hours and process the updates accordingly.

Which of the following is an ideal way to ensure the secondary application can get the relevant changes from the Dynamo DB table.

✅ Use Dynamo DB streams to monitor the changes in the Dynamo DB table.
⚪ Create another Dynamo DB table with the records modified in the last 2 hours.
⚪ Transfer the records to 53 which were modified in the last 2 hours
⚪ insert a timestamp for each record and then scan the entire table for the timestamp as per the last 2 hours.

Q41)

Your team is responsible for an AWS Elastic Beanstalk application. The business requires that you move to a continuous deployment model, releasing updates to the application multiple times per day with zero downtime.

What should you do to enable this and still be able to roll back almost Immediately in an emergency to the previous version?

⚪ Enable rolling updates in the Elastic Beanstalk environment, setting an appropriate pause time for application startup.
✅ Create a second Elastic Beanstalk environment running the new application version, and swap the environment CNAME5.
⚪ Develop the application to poll for a new application version In your code repository; download and Install to each running Elastic Beanstalk instance.
⚪ Create a second Elastic Beanstalk environment with the new application version, and configure the old environment to redirect clients, using the HTTP 301 response code, to the new environment

---

Q42) By default in Ops work , how many application versions can you rollback up to? Please select:

- 2
- 1
- 3
- ✓ 4

Q43)

You have an Ops work stack defined with Linux instances. You have executed a recipe, but the execution has failed.

What Is one of the ways that you can use to diagnose what was the reason why the recipe did not execute correctly.

- Use AWS Cloud trail and check the Ops work logs to diagnose the error
- Use AWS Config and check the Ops work logs to diagnose the error
- Deregister the instance and check the EC2 Logs
- ✓ Log into the instance and check if the recipe was properly configured.

Q44)

You currently have an application with an Auto Scaling group with an Elastic Load Balancer configured in AWS. After deployment users are complaining of slow response time for your application.

Which of the following can be used as a start to diagnose the issue ?

- Use Cloud watch to monitor the Memory Utilization
- Use Cloud watch to monitor the CPU Utilization
- ✓ Use Cloud watch to monitor the ELB latency
- Use Cloud watch to monitor the Health Host Count metric

Q45)

You have an a video processing application hosted in AWS. The video?s are uploaded by users onto the site. You have a program that is custom built to process those videos. The program is able to recover Incase there are any failures when processing the videos.

Which of the following mechanisms can be used to deploy the instances for carrying out the video processing activities , ensuring that the cost is kept at a minimum?

- Create a launch configuration with Dedicated Instances. Ensure the User Data section details the installation of the custom software. Create an Auto scaling group with the launch configuration.
- ✓ Create a launch configuration with Spot Instances. Ensure the User Data section details the installation of the custom software. Create an Auto scaling group with the launch configuration.
- Create a launch configuration with On-Demand Instances. Ensure the User Data section details the installation of the custom software. Create an Auto scaling group with the launch configuration.
- Create a launch configuration with Reserved Instances. Ensure the User Data section details the Installation of the custom software. Create an Auto scaling group with the launch configuration.

Q46)

There is a requirement for a vendor to have access to an 53 bucket in your account. The vendor already has an AWS account.

How can you provide access to the vendor on this bucket, Please select:

- Create a new IAM group and grant the relevant access to the vendor on that bucket.
- ✓ Create a cross-account role for the vendor account and grant that role access to the 53 bucket.
- Create an 53 bucket policy that allows the vendor to read from the bucket from their AWS account.
- Create a new IAM user and grant the relevant access to the vendor on that bucket.

Q47)

You need to grant a vendor access to your AWS account. They need to be able to read protected messages in a private S3 bucket at their leisure. They also use AWS.

What is the best way to accomplish this?

- ✓ Create a cross-account IAM Role with permission to access the bucket, and grant permission to use the Ro to the vendor AWS account.
- Create an IAM User with API Access Keys. Grant the User permissions to access the bucket. Give the vendor the AWS Access Key ID and AWS Secret Access Key for the User.
- Generate a signed S3 PUT URL and a signed 53 PUT URL both with wildcard values and 2 year durations, Pass the URL5 to the vendor
- Create an EC2 Instance Profile on your account. Grant the associated IAM role full access to the bucket. Start an EC2 instance with this Profile and give SSH access to the instance to the vendor.

Q48)

You currently have an Auto scaling group that has the following settings Mm capacity - 2 Desired capacity - 2 Maximum capacity - 2 Your launch configuration has AMI?s which are based on the t2.micro instance type. The application running on these instances are now experiencing issues and you have identified that the solution is to change the Instance type of the instances running In the Auto scaling Group.

Which of the below solutions will meet this demand?

Change the desired and maximum size of the Auto scaling Group to 4. Make a copy the Launch configuration. Change the Instance type in the new launch configuration. Attach that to the Auto scaling Group. Change the maximum and Desired size of the Auto scaling Group to 2

Change the Instance type in the current launch configuration. Change the Desired value of the Auto scaling Group to 4. Ensure the new Instances are launched.

Delete the current Launch configuration. Create a new launch configuration with the new instance type and add it to the Auto scaling Group. This will then launch the new instances.

Make a copy the Launch configuration. Change the Instance type In the new launch configuration. Attach that to the Auto scaling Group. Change the maximum and Desired size of the Auto scaling Group to 4. Once the new instances are launched, change the Desired and maximum size back to 2.